
SmallFloats.jl

0.4.0

Jeffrey Sarnoff

2026-08-21T19:56Z

Contents

Contents	ii
I Home	1
1 SmallFloats.jl	2
1.1 The promises	2
1.2 Documentation map	3
1.3 Code point or number?	3
II Start Here	4
2 Installation	5
2.1 Installation	5
3 First Session	6
3.1 First Session	6
4 Core Model	10
4.1 Core Model	10
III User Guide	15
5 Guide Overview	16
5.1 User Guide	16
6 P3109 in One Chapter	18
6.1 P3109 in One Chapter	18
7 Format Anatomy	21
7.1 Format Anatomy	21
8 Values, Code Points, and Conversion	24
8.1 Values, Code Points, and Conversion	24
9 The Exact-Then-Project Contract	29
9.1 The Exact-Then-Project Contract	29
10 Rounding and Saturation	32
10.1 Rounding and Saturation	32
11 Julia's Numeric Model	34
11.1 Julia Numeric Behavior	34
12 Session State and Reproducibility	37
12.1 Session State and Reproducibility	37
13 Defined, Stochastic, and Approximate Results	40
13.1 Defined, Stochastic, and Approximate Results	40
14 Code-Point Algebra	42
14.1 Formal Code-Point Algebra	42

IV	Workflows	46
15	Choose a Format	47
15.1	Choose a Format	47
16	Control Rounding and Overflow	50
16.1	Control Rounding and Overflow	50
17	Quantize and Measure a Tensor	53
17.1	Quantize and Measure a Tensor	53
18	Make Stochastic Work Reproducible	56
18.1	Make Stochastic Work Reproducible	56
19	Use Blocks for Dynamic Range	59
19.1	Use Blocks for Dynamic Range	59
20	Run Operations over Arrays	62
20.1	Run Operations over Arrays	62
21	Pack Values for Storage	65
21.1	Pack Values for Storage	65
22	Interoperate with Float16 and BFloat16	67
22.1	Interoperate with Float16 and BFloat16	67
23	Produce a Conformance Declaration	70
23.1	Read and Export Conformance	70
24	Register and Measure an Approximation	72
24.1	Register an Approximation	72
V	Performance and Examples	74
25	Performance Model	75
25.1	Performance Model	75
26	Example Gallery	77
26.1	Example Gallery	77
27	Applied Sessions	80
27.1	Applied Sessions	80
28	Verification Sessions	87
28.1	Verification Sessions	87
29	Performance Sessions	94
29.1	Performance Evidence	94

VI	Reference	105
30	Formats and Values	106
30.1	Formats and Value Queries	106
31	Projection Specifications	109
31.1	Projection Specifications	109
32	Operation Catalog	112
32.1	Operation Catalog	112
33	Arrays, Blocks, and Packed Storage	116
33.1	Arrays, Blocks, and Packed Storage	116
34	Defaults, Randomness, and Display	119
34.1	Defaults, Randomness, and Display	119
35	Julia Compatibility	122
35.1	Julia Compatibility Register	122
36	Conformance and Approximation	126
36.1	Conformance and Approximation Registry	126
37	Public API Index	128
37.1	Public API Index	128
38	Internal API Index	141
38.1	Internal API Index	141
VII	Technical Guide	151
39	Guide Overview	152
39.1	Technical Guide	152
40	Architecture and Invariants	154
40.1	Architecture and Invariants	154
41	Encoding and Decoding	156
41.1	Encoding and Decoding	156
42	Projection Engine	158
42.1	Projection Engine	158
43	Oracle and Rigor Classes	160
43.1	Oracle and Rigor Classes	160
44	Function Tables and Array Kernels	163
44.1	Function Tables and Array Kernels	163
45	Exact Block Reductions	165
45.1	Exact Block Reductions	165
46	Verification Strategy	169
46.1	Verification Strategy	169
47	Add an Operation	171
47.1	Add an Operation	171
48	Verify Custom Code	179
48.1	Verify Custom Code	179
49	Benchmark Correctly	183
49.1	Benchmark Correctly	183
50	Documentation Claim Ledger	185
50.1	Documentation claim ledger	186

VIII Help	187
51 Troubleshooting	188
51.1 Troubleshooting	188
52 Cheat Sheet	191
52.1 Cheat Sheet	191
53 Glossary	198
53.1 Glossary	198
54 Standard and Design Notes	200
54.1 Standard and Design Notes	200

Part I

Home

Chapter 1

SmallFloats.jl

An exact, policy-visible Julia implementation of the IEEE P3109 small binary formats, from 3- through 16-bit code points.

In an 8-bit format almost every operation rounds: most numbers you construct, and most sums and products you form, are not among the format's few hundred values. The question that decides a small-float experiment is therefore not *whether* rounding happened but *where, and under which policy* — and that is exactly what most implementations hide. SmallFloats.jl keeps it visible: the finite datum set, the rounding/saturation policy as a named argument, and the one projection that turns each exact result back into a code point. All 504 legal formats are supported, byte-sized Code8 representations and wider Code16 representations alike, without the storage difference ever entering the arithmetic semantics.

```
julia> using SmallFloats

julia> x = Binary8p4se(1.6)           # construction projects a value
Binary8p4se(1.625 ≈ 0x45)

julia> Add(Binary8p4se, RNE_SN, x, Binary8p4se(1.75))
Binary8p4se(3.5 ≈ 0x4e)              # exact sum, then one projection

julia> T = format(16, 6, true, true)  # a wider format is an ordinary format
Binary16p6se

julia> codeunit_type(T), SmallFloats.datumcarrier(T)
(UInt16, Float64)
```

The last line is an implementation-facing fact with a user-facing consequence: the code point of a wide format need not fit in a byte, and decode returns the exact carrier selected for the format rather than promising Float64 in every case. The value model itself stays the same.

1.1 The promises

- **Defined results are exact, then projected once.** Construction, conversion, scalar operations, tables, and kernels all agree on RoundToPrecision → Saturate → Encode — one rounding, at the end. No path rounds twice, so no result drifts quietly from the draft's defined answer.
- **Policy is visible.** A ProjSpec names rounding and saturation in the call itself, so a line of code produces the same bits in every session that runs it — repeatability never depends on ambient state.
- **Approximation is declared and measured.** Approximate kernels are opt-in, retrieved by name, and carry a code-point error bound κ measured by exhaustive enumeration at registration time; a declaration that understates the measurement is rejected.

- **Representation does not weaken the datum contract.** Code8 and Code16 select storage, decode, and table policies; both represent the draft's finite datum sets exactly, so widening a format changes cost, never meaning.

1.2 Documentation map

New here? Start with [Installation](#), take the [First Session](#), then read the [Core Model](#). Those three pages establish the vocabulary used everywhere else.

Using the package? Follow the [User Guide](#) through formats, values, projection, arrays, blocks, storage, reproducibility, and Julia interoperability. Use the [Cheat Sheet](#) only when you already know the distinction you are looking up.

Choosing or evaluating a format? Begin with [Choose a Format](#), then [Quantize and Measure a Tensor](#). [Example Gallery](#) routes to runnable basic, AI, machine-learning, and deep-learning sessions.

Maintaining or extending SmallFloats? Read the [Technical Guide](#). It follows a datum through representation, codec, projection, oracle, tables, blocks, verification, and benchmarking before pointing to the extension checklists.

Looking up a contract? The [Formats and Values](#), [Projection Specifications](#), [Operation Catalog](#), and [Julia Compatibility Register](#) are curated references. The public and internal API indices are docstring inventories, not substitutes for the guides.

1.3 Code point or number?

An Unsigned argument names a stored code point directly. Any other Real argument is treated as a number and rounded to the format's nearest datum under the active projection policy.

```
julia> Binary8p4se(0x02), Binary8p4se(2)
(Binary8p4se(0.001953125 ≈ 0x02), Binary8p4se(2.0 ≈ 0x48))
```

The same digit, two questions, answers a factor of 1024 apart — and neither call is an error. `Convert(T, ρ, x)` projects whatever the argument's integer type. [Values, Code Points, and Conversion](#) has the complete rule.

Part II

Start Here

Chapter 2

Installation

2.1 Installation

A working install in about five minutes.

Requirements

- **Julia 1.12** or later.
- **Quadmath** (libquadmath's Float128) is a hard dependency. It is used only inside the oracle and exact-fallback paths — never where it could change a result.

Install

The package is not yet registered, so install by URL:

```
pkg> add https://github.com/JeffreySarnoff/SmallFloats.jl
```

Append `#main` (or a tag/commit) to pin a revision:

```
pkg> add https://github.com/JeffreySarnoff/SmallFloats.jl#main
```

For an editable install, `pkg> dev https://github.com/JeffreySarnoff/SmallFloats.jl` clones into `~/.julia/dev`, or use an existing clone:

```
pkg> dev path/to/SmallFloats.jl
```

Smoke test

Thirty seconds to confirm everything is in place:

```
julia> using SmallFloats

julia> Binary8p4se(1.6) + Binary8p4se(0.25)
Binary8p4se(1.875 ≈ 0x47)

julia> conformance() !== nothing
true
```

The first line proves the projection engine is live (1.6 rounds to 1.625, plus 0.25, lands exactly on 1.875); the second proves the conformance machinery is.

Chapter 3

First Session

3.1 First Session

This session establishes the package's working vocabulary and runs one scalar and one array operation. It assumes [Installation](#) has succeeded.

```
julia> using SmallFloats
```

Construct a value from a number

```
julia> x = Binary8p4se(1.6)
Binary8p4se(1.625 ≈ 0x45)
```

Binary8p4se is an 8-bit, precision-4, signed, extended format: a fixed set of 256 code points. The number 1.6 is not among the values it holds. The nearest data in this format are 1.5, 1.625, and 1.75 — the spacing here is 0.125 — so construction must *project*: it rounds 1.6 to the nearest datum, 1.625, and stores that datum's code point, 0x45. Every construction from a Real answers the same question — *which datum stands in for my number?* — and the constructor never refuses to answer; it rounds.

The display shows both halves of the answer: the datum that now represents 1.6, and the identity under which it is stored. Application output may print more compactly; the documentation uses the typed display so that no projection is invisible.

Construct a value from a code point

```
julia> y = Binary8p4se(0x46)
Binary8p4se(1.75 ≈ 0x46)
```

An Unsigned argument means **code point**, not numeric value. The signedness of the integer type is the entire distinction. Compare:

```
julia> Binary8p4se(0x02), Binary8p4se(2)
(Binary8p4se(0.001953125 ≈ 0x02), Binary8p4se(2.0 ≈ 0x48))
```

The same digit, read two ways: 0x02 asks *give me the datum stored at code point 2* and gets 0.001953125; 2 asks *represent the value 2* and gets 2.0 — results a factor of 1024 apart. Neither call is an error, so an integer of the wrong signedness is not caught; it is a different, valid request. Use an unsigned integer when you mean stored identity — reading a byte from a file, reconstructing a value from a wire format. Use a signed integer or any other Real when you mean a value to project. decode and codepoint ask the two questions in reverse — *what value does this represent?* and *under what identity?*:

```
julia> decode(x), codepoint(x)
(1.625, 0x45)
```

Display

`get_show_style()` returns the process-wide style and `set_show_style!(style)` sets it. `VALID_SHOW_STYLES` contains four values:

Style	Example rendering
<code>:value</code>	1.625
<code>:codepoint</code>	0x45
<code>:datum</code>	(1.625 $\approx$ 0x45)
<code>:typed</code>	Binary8p4se(1.625 $\approx$ 0x45)

An `IIOContext` property, `:binary_show_style`, overrides the process default for one stream. Prefer that form in concurrent code and in libraries that should not change another caller's display.

The four styles answer four different questions about the same value. For the `x = Binary8p4se(1.6)` constructed above:

- `:value` prints 1.625 and answers *what number is this?* Use it when results are read as numbers — comparing against a `Float64` reference, printing residuals, tabulating measurements — where $\approx$ 0x45 on every line is noise.
- `:codepoint` prints 0x45 and answers *which code point is stored?* Use it where identity matters and magnitude does not: diffing a serialized buffer, inspecting a slot in a packed array, checking that a round-trip returned to the same code point. Two values are the same datum exactly when their code points match.
- `:datum` prints (1.625 $\approx$ 0x45) and answers *where did my input land?* Use it when watching a projection: construct from 1.6 and see, on one line, both the datum the input rounded to and the identity it received. Either half alone hides the other.
- `:typed` prints `Binary8p4se(1.625  $\approx$  0x45)` and answers *in which format?* Use it wherever the format cannot be inferred later — logs, saved output, examples meant to be pasted elsewhere, sessions mixing formats — because a code point means nothing without its format: 0x45 names a different value in every format that contains it. This documentation renders `:typed` throughout for exactly that reason.

For a temporary rendering choice, attach the style to the destination stream instead of changing process-global state:

```
io = IOContext(stdout, :binary_show_style => :typed)
show(io, x)
```

`set_show_style!` is appropriate for an interactive session you own. Restore the previous setting before returning control to another caller; library code should use an `IIOContext` supplied by its caller, or create a local one as above. The display style changes presentation only: `decode(x)`, `codepoint(x)`, comparisons, and arithmetic are unaffected.

Compute with an explicit policy

```
julia> Add(Binary8p4se, RNE_SN, x, y)
Binary8p4se(3.5  $\approx$  0x4e)
```

Read the call from left to right:

```
operation(result format, projection specification, operands...)
```

The exact sum, $1.625 + 1.75 = 3.375$, is once again not a datum. It falls *exactly halfway* between the neighbors $3.25 \approx 0x4d$ and $3.5 \approx 0x4e$, so the operation's result is entirely a question of policy: RNE_SN rounds to nearest with ties to even, which selects 3.5, then applies the draft's SatNone boundary rules. A different projection could legitimately land the same exact sum on $0x4d$. That is why the projection sits in the argument list: it is an operand of the operation, not an ambient setting the result quietly depends on.

The ordinary Julia spelling uses the current default projection:

```
julia> x + y
Binary8p4se(3.5 ≈ 0x4e)
```

The two spellings differ in who decides the policy. `+` delegates to the session default, which is right for exploratory same-format arithmetic at the REPL. `Add(T, ρ, ...)` names the policy at the call site, which is right for code whose results must not change because someone, somewhere, changed the session default — the line means the same thing in every session that runs it.

See projection change an overflow

```
julia> w, two = Binary8p4se(200.0), Binary8p4se(2.0);

julia> Multiply(Binary8p4se, RNE_SN, w, two)
Binary8p4se(Inf ≈ 0x7f)

julia> Multiply(Binary8p4se, RNE_SF, w, two)
Binary8p4se(224.0 ≈ 0x7e)
```

A projection has already happened before either `multiply` runs: 200 is not a datum, and it sits exactly halfway between the neighbors 192 and 208, so `w` holds $192.0 \approx 0x7c$ — another tie broken to even. Both calls therefore compute the exact product 384, which exceeds 224, the largest finite datum of the format. The two projections disagree about what exceeding the range should mean. SatNone answers with infinity, preserving the evidence that the range was exceeded; SatFinite clamps to $224.0 \approx 0x7e$, keeping the result finite for pipelines that cannot absorb an infinity. Same exact product, two defined results: the projection is part of the operation's meaning, not a rounding afterthought.

Apply the same contract to an array

```
julia> A = Binary8p4se.([-1.0, 0.5, 2.0]);

julia> Exp(Binary8p4se, RNE_SN, A)
3-element Vector{Binary8p4se}:
 Binary8p4se(0.375 ≈ 0x34)
 Binary8p4se(1.625 ≈ 0x45)
 Binary8p4se(7.5 ≈ 0x57)
```

Each element is the answer the scalar call would give. Look at the middle one: $\exp(0.5) = 1.6487\dots$ projects to $1.625 \approx 0x45$ — the same datum `x` landed on at the top of this page, reached now from a different number by a different operation. The array form adds no semantics of its own: same result format, same projection, same scalar answer per element. Its implementation may serve those answers from a cached table, but a table entry is the scalar result precomputed, so the contract is unchanged whether the table is used or not.

What this leaves open

This session used four things without explaining any of them: a format that did not contain 1.6, a projection named in the call, an overflow that resolved two ways, and an array result that matched the scalar one. [Core Model](#) is those four as one connected idea, and it is the page to read next — everything else in the documentation develops one of its five parts.

If you would rather work first and generalize later, the two workflows that start from real data are [Choose a Format](#) — which format fits numbers you already have — and [Quantize and Measure a Tensor](#), which measures what a choice cost you. Both assume only what this page established.

See also

[Formats and Value Queries](#), [Projection Specifications](#), and [Operation Catalog](#).

Chapter 4

Core Model

4.1 Core Model

Five ideas organize the package, and each answers the question the previous one leaves open. *What is a format?* A finite universe of exact values. *How do you hold one?* A value is a code point, and reading it back is exact. *What happens when you compute?* Exact math, then one projection. *Who decides where the exact result lands?* You do — a projection is a named pair of rounding and saturation choices. *And what if you trade exactness for speed?* Only by name, with the error measured — never by default.

A format is a finite set of datums

A SmallFloats *format* is not a floating-point type in the IEEE-754 sense of a continuum sampled at machine precision. It is a **finite, enumerated set of exact values** — the *datums* — each with a name, its *code point*. The format Binary4p2se (4 bits, precision 2, signed, extended) has 16 code points and they are all of them:

```
julia> decode.(sort(Binary4p2se.(0x00:0x0f)))
16-element Vector{Float64}:
 NaN
 -Inf
 -2.0
 -1.5
 -1.0
 -0.75
 -0.5
 -0.25
 0.0
 0.25
 0.5
 0.75
 1.0
 1.5
 2.0
 Inf
```

That listing *is* the format. There is nothing between the grid points, no hidden state, no approximation in the set itself — the spacing you see (subnormals near zero, binades doubling outward) is the whole structure.

Three consequences fall out immediately:

- **One NaN, no negative zero.** The first entry — NaN sorts below $-\text{Inf}$ in the draft's total order (§4.12.1) — is the format's single NaN; the value 0.0 appears exactly once. IEEE-754 spends thousands of code points on NaN payloads and two on zeros; P3109 spends them on datums.

- **The set has edges.** -Inf, Inf, and NaN occupy three of the sixteen names, which is why the largest *finite* value is 2.0 and not something larger. Finite formats (suffix f instead of e) spend those two infinity code points on more finite datums instead.
- **Larger formats are the same idea with more points.** Binary8p4se is 256 names; Binary16p11se is 65 536. The 504 formats are just the legal choices of how many names (bitwidth K), how much of the budget goes to significand vs exponent (precision P), and whether the set includes negatives (s/u) and infinities (e/f).

Now you can read: [Formats, Aliases & Representations](#) for why the type Binary8p4se is a concrete *representation* of the abstract format Binary{8,4,true,true} — and why `===` is false but `<:` is true.

A value is a wrapper around a code point; decode is exact

Construct a value and you hold the code point; ask what it means and you get the datum back, exactly:

```
julia> x = Binary8p4se(1.6)
Binary8p4se(1.625 ≈ 0x45)

julia> codepoint(x)
0x45

julia> decode(x)
1.625
```

The display `1.625 ≈ 0x45` is the idea in one line: the code point `0x45` *means* the datum `1.625`, and the `≈` is exact, not approximate. `decode` never rounds — it is a table lookup from name to meaning.

The asymmetry that matters: `decode` is exact, but **construction is a projection**. `Binary8p4se(1.6)` did not store `1.6`; it stored the name of the nearest datum, which happens to be `1.625`. The value you get back is the value the format can say, not the value you asked for.

This is why the Unsigned-means-code-point rule (from the First Session) exists: on the code-point side, `codepoint` and the constructor are inverses *by design* —

```
julia> Binary8p4se(codepoint(x)) == x
true
```

— and that round-trip is the format's identity. A value is its code point; everything else is a view.

The inverse runs one way only, and the asymmetry above is the reason. `T(codepoint(x)) == x` holds for every code point, because nothing is discarded going out and back. Start from a number instead and it fails: `decode(T(1.6))` is `1.625`, not `1.6` — construction already chose a datum, and no later step can recover which number asked for it. `decode` and `codepoint` are not the pair here; both merely read a value that has already been decided.

Now you can read: [Values, Code Points, and Conversion](#) for the four ways in and two ways out, and the carriers `decode` returns at wide formats.

Every computation is exact math, then one projection

When you compute, SmallFloats does not simulate low-precision arithmetic step by step. It computes the **mathematically exact result** of the operation on the decoded datums, then applies **one projection** — one rounding plus one saturation decision — to land on a code point:

```
exact result → RoundToPrecision → Saturate → code point
```


$x + y$ from the First Session: exact sum 3.375, projected once to 3.5. The overflow example from the same page: w holds 192 — constructing it from 200 was itself one projection — so the multiply computes the exact product 384, which no finite datum can hold. One projection then decides what an out-of-range 384 becomes, and the next idea says who makes that choice.

Two things follow from "one projection, at the end":

- **There is no accumulated rounding error inside an operation.** A fused $\text{FMA}(x, y, z)$ rounds once, not twice. A `BlockDotProduct` accumulates all 32 lane products exactly and projects once at the end. Intermediate roundings are not "small" — they are *absent*.
- **The default path is bit-exact.** Because the exact math is carried far enough (Float64, Float128, or MPFR enclosures, escalating automatically), the code point you get is *the* defined answer, not a good approximation of it. The package can prove this because the value sets are finite — see the fifth idea.

Now you can read: [The Exact-Then-Project Contract](#) for what "exact" means at each carrier width, and the symbolic-sticky trick that lets directed modes land exactly on asymptotes.

Projection = rounding mode $\times$ saturation mode

A projection specification is the pair of two independent choices:

	what it decides	examples
rounding mode	which neighbor a between-grid value lands on	nearest-ties-even, toward zero, toward $\pm\infty$, stochastic
saturation mode	what an out-of-range result becomes	clamp to finite (SF), propagate infinities (SP), draft rows (SN)

Watch the exact product 384 resolve three ways (Binary8p4se, max finite 224):

```
julia> w, two = Binary8p4se(200.0), Binary8p4se(2.0);

julia> Multiply(Binary8p4se, RNE_SN, w, two)    # nearest + draft rows → Inf
Binary8p4se(Inf ≈ 0x7f)

julia> Multiply(Binary8p4se, RNE_SF, w, two)    # nearest + clamp → 224
Binary8p4se(224.0 ≈ 0x7e)

julia> Multiply(Binary8p4se, RTZ_SN, w, two)    # toward zero + draft rows → 224
Binary8p4se(224.0 ≈ 0x7e)
```

The first two differ in *saturation* (SN vs SF); the first and third differ in *rounding* (RNE vs RTZ) — and the third row is the subtle one: under `SatNone`, a directed mode that points *away* from the overflow clamps to the extremal finite value, so `RTZ_SN` agrees with `RNE_SF` here for different reasons.

The full deterministic grid is 6 rounding modes $\times$ 3 saturation modes = 18 predefined specs (`RNE_SN`, `RNA_SF`, `RTP_SP`, ...), plus three stochastic families parameterized by a random-bit budget. `RNE_SN` is the package-wide default — which is exactly why `+`, `exp`, and `T(1.6)` all agree with `Add(T, RNE_SN, ...)`, `Exp(T, RNE_SN, ...)`, and `Convert(T, RNE_SN, ...)`.

The default is a session default

`+` and friends follow `DefaultProjection()`, which is process-global and stays set. Convenience syntax is for exploratory code; anything that must be reproducible names its spec explicitly. The workflow guides distinguish explicit local policy from `with_default_projection`, which only consumes the current process-global default, before showing `DefaultProjection!`.

Now you can read: [Rounding and Saturation](#) for the experiments, and [Projection Specifications](#) for the complete grid.

A Whole of Five Parts




Figure 4.1: A whole of five parts: a finite set of datums, values as code points, exact computation, one controlled projection, and exact or measured results.

Approximation is opt-in, named, and measured

The package can exploit the first idea's finiteness without weakening it. For example, a pure unary Binary8p4se operation has 256 inputs, so its complete behavior may be cached as a 256-byte result table built through the exact scalar path. Other signatures use the representation-aware table policy or compute directly. In every case, exactness is inherited: a table entry and its scalar call are the same defined function by construction.

The same finiteness makes approximation *honest*. If you register a faster but inexact kernel — say a hardware-style hard-tanh for Tanh — the registry does not take your word for its accuracy. It **measures** the worst-case deviation, exhaustively over every input, as a distance in code points:

```
julia> κ, exhaustive = measure_kappa(
    x -> Clamp(Binary8p4se, RNE_SN, x,
        Negate(one(Binary8p4se)), one(Binary8p4se)),
    :Tanh, Binary8p4se, (Binary8p4se,), RNE_SN)
(4.0, true)
```

$\kappa = 4$ code points, verified over all 256 inputs — that is a measured fact, not a claim. Declare κ smaller than the measurement and registration is *rejected*. Nothing approximate is reachable from the default API; an approximation exists only because you named it, and its κ is part of the conformance declaration the package exports.

This is the package's answer to the problem stated in the Introduction: small formats are easy to implement approximately and hard to implement exactly. SmallFloats implements them exactly, and when you choose otherwise, it makes the choice explicit and quantified.

Now you can read: [Register & Measure an Approximation](#), and [Conformance & Approximation Registry](#).

Visual Synopsis

Read left to right: a value is a name in a finite set (1–2); computing is exact math followed by one projection (3); the projection is your choice of rounding and saturation (4); and the result is exact unless you explicitly register otherwise (5).

Where to go next

These five ideas are the whole public model; the rest of the documentation develops them. Each of the three that raise a further question has a page that answers it. *What does "exact" actually mean, when the exact result is wider than any machine float?* — [The Exact-Then-Project Contract](#), which follows idea 3 down to the carrier. *How do I name a policy, and which one do I want?* — [Rounding and Saturation](#) for the two axes, [Projection Specifications](#) for the complete grid of names. *How is an approximation held to its declared κ ?* — [Conformance and Approximation Registry](#).

For the format-query vocabulary the examples above used in passing, see [Formats and Value Queries](#).

Part III

User Guide

Chapter 5

Guide Overview

5.1 User Guide

Everything needed to use `SmallFloats.jl`, in the order it becomes useful. The public model is deliberately small — a format is a finite datum set; a value names one datum by a code point; an operation evaluates exactly and then projects once under an explicit policy — and each page of this guide answers one question that model raises.

Read [First Session](#) first if you have not constructed a `Binary` value yet. This guide routes; the pages carry the examples and the boundary conditions.

What the standard says, and what a format is

Where do these rules come from? [P3109 in One Chapter](#) is the draft in fifteen minutes: the decode → compute exactly → project once law, the four format parameters, the single NaN and absent negative zero, and the bias formula that makes `Binary16p11se` differ from `Float16` at every code point.

What exactly is `Binary8p4se`? [Format Anatomy](#) answers with the four-parameter space `Binary{K,P,SGN,EXT}`, the line between the abstract format and the concrete representation your arrays must use, and the three routes to all 504 formats when the 120 exported aliases are not enough.

Values and computation

How does a value name a datum? [Values, Code Points, and Conversion](#). The rule worth retaining: an `Unsigned` constructor argument is a code point, every other `Real` is a value to project — and decode is always exact, on the format's own carrier.

What does an operation promise? [The Exact-Then-Project Contract](#): exact math first, one projection at the end, no second write path — plus what "exact" means at each carrier width and why `Rational` input is refused rather than double-rounded.

Who decides where an exact result lands? [Rounding and Saturation](#) develops the two independent policy axes a `ProjSpec` names. Use [Control Rounding and Overflow](#) when you want a short procedure for choosing and testing one.

Living inside Julia

What happens when `Binary` meets ordinary numbers? [Julia's Numeric Model](#): mixed formats fail loudly, ordinary numbers promote to the format's carrier, and no automatic format-widening rule can be built soundly — the page explains why.

What does session state control? [Session State and Reproducibility](#): the five process-global defaults behind `T(x)`, `x + y`, and friends — and why reusable code names its policy instead of relying on them.

Results and their provenance

Why did two runs differ? [Defined, Stochastic, and Approximate Results](#) keeps the three legitimate reasons apart and gives the four-question provenance test that separates all of them from a defect.

What is underneath it all? [Code-Point Algebra](#) states the package as integer identities — where NaN sits, how a code decodes, how ordering is recovered — verified against all 504 formats.

Working with data

The Workflows chapter applies the model. [Choose a Format](#) turns precision, range, domain, and storage requirements into a reproducible choice; [Quantize and Measure a Tensor](#) compares such choices under one fixed projection rather than changing data and policy at once. For bulk work, continue to [Run Operations over Arrays](#), [Use Blocks for Dynamic Range](#), and [Pack Values for Storage](#). One rule holds throughout: acceleration may change how work is scheduled, never which projected code point is returned.

For interoperability and evidence: [Interoperate with Float16 and BFloat16](#), [Make Stochastic Work Reproducible](#), [Produce a Conformance Declaration](#), and [Register and Measure an Approximation](#).

See also

[Cheat Sheet](#) for syntax at a glance, [Example Gallery](#) for runnable end-to-end sessions, and [Technical Guide](#) for the implementation and verification arguments behind these guarantees.

Chapter 6

P3109 in One Chapter

6.1 P3109 in One Chapter

The IEEE P3109 draft — *Arithmetic Formats for Machine Learning* — as this package sees it. Fifteen minutes here buys you the vocabulary every other page assumes. Claims are tagged with the draft clause they come from; the full transliteration is in the repository at docs/other/IEEE_D1.md for when you want the normative text itself.

The one law everything hangs on

Every numeric operation in the draft is three steps (§4.3.2):

decode the operands → compute exactly over the extended reals → project once

There is no arithmetic "in the format." Operands are decoded to exact values, the operation is performed exactly, and a single projection — rounding plus saturation — writes the result back to a code point. The comparisons and predicates stop after decoding (they return Booleans); everything else follows the law. SmallFloats.jl is built as a direct implementation of this sentence.

Formats: four parameters, 504 instances

A format is $\text{Binary}\{K, P, \Sigma, \Delta\}$ (§3.1):

parameter	meaning	range
K	total bitwidth	3 – 16
P	significand precision, implicit bit included	signed: $0 < P < K$; unsigned: $0 < P \leq K$
Σ	signedness: signed or unsigned	
Δ	domain: extended (has $\pm\text{Inf}$) or finite	

The short name reads off the parameters: $\text{Binary8p4se} = K\ 8, p\ 4, \text{signed}, \text{extended}$. The draft's exponent bias is *derived*, not chosen (§3.1, §4.14):

$$B = 2^{(K-P-1)} \quad (\text{signed}) \qquad B = 2^{(K-P)} \quad (\text{unsigned})$$

This symmetric-about-the-midpoint bias is the single most consequential formula in the draft — see "Differences from IEEE 754" below.

Datums and code points

A format's *datum set* is a finite subset of the closed extended reals ($\mathbb{R}$ plus $\pm\text{Inf}$ and NaN), in bijection with the code points $0 \dots 2^K - 1$ (§3, Annex B). Every finite datum has a canonical form

$$(-1)^{\text{sign}} \cdot S \cdot 2^{(1-P)} \cdot 2^E \quad E > -B$$

with significand S classifying it as normal ($2^{(P-1)} \leq S < 2^P$) or subnormal ($0 < S < 2^{(P-1)}$).

The specials are where the draft parts company with IEEE-754 (Annex B, Table 3):

- **Exactly one zero.** There is no -0 ; a signed format has an odd number of finite datums.
- **Exactly one NaN**, sitting at the code point IEEE-754 gives to -0 in signed formats. No payloads, no quiet/signalling distinction — the thousands of code points IEEE spends on NaN encodings are datums here.
- **$\pm\text{Inf}$ only in extended formats**, adjacent to the finite extremes. A finite format spends those code points on more finite values.

Projection: rounding × saturation

A projection specification p is the pair (*rounding mode*, *saturation mode*) (§4.2, §4.7.3–4.7.5).

Six deterministic rounding modes — *NearestTiesToEven*, *NearestTiesToAway*, *TowardPositive*, *TowardNegative*, *TowardZero*, *ToOdd* — plus three stochastic families *StochasticA/B/C*, which consume N random bits as an unsigned draw $R \in 0 \dots 2^N - 1$ (§4.7.4; the draft deliberately leaves the bit source unspecified).

Three saturation modes decide what out-of-range results become (§4.7.5):

mode	finite overflow	$\pm\text{Inf}$ input
<i>SatFinite</i>	clamp to the finite extremes	clamp
<i>SatPropagate</i>	clamp	preserve representable infinities, else clamp
<i>SatNone</i>	the draft's rows: to $\pm\text{Inf}$, finite extreme, or NaN depending on rounding direction, signedness, and domain	preserve, else NaN

One ordering fact the whole clause turns on (§4.7.3 NOTE 1): **rounding precedes saturation**, but the rounding mode is passed *into* saturation so a directed mode can force a finite result. Saturation itself never rounds.

The operation catalog

§4.9–§4.16 define the catalog: *Convert*; sign and basic arithmetic (*Abs* through *Divide*, fused FMA and FAA); the elementary functions (roots, exp/log families, trig, hyperbolic, π -scaled trig, *Softplus*, *Hypot*, *ArcTan2*); the extrema families (*Minimum/Maximum* and their NaN-quieting, *magnitude*, and finite variants, plus *Clamp*); comparisons, predicates, and *TotalOrder*; format-level queries (*BitwidthOf*, *MaxFiniteOf*, ...); and the neighbour operations *NextGreaterThan/NextLessThan*.

Some results are visible only in the pattern tables, and they surprise:

- $\text{Divide}(x, 0) = \text{NaN}$ for every x — including $\pm\text{Inf}$ (§4.10.5 NOTE 2). There is no IEEE-style $\pm\text{Inf}$ quotient.
- $\text{Recip}(\pm\text{Inf}) = 0$ and $\text{Recip}(0) = \text{NaN}$ (§4.10.8).
- *TotalOrder* puts NaN first, below $-\text{Inf}$ — deterministic tie-breaking for search and sort (§4.12.1), and the opposite end from Julia's float-sorting convention.
- $\text{NextGreaterThan}(+\text{Inf}) = \text{NaN}$, unlike IEEE-754 *nextUp* (§4.16 NOTE 2).
- FAA is associative by construction: $\text{Add}(\text{Add}(X, Y), Z) = \text{Add}(X, \text{Add}(Y, Z))$ computed exactly (§4.10.7).

Blocks and scaled operations

Clause 5 adds the MX-style scheme: a *block* is a scale factor in one format plus B elements in another, representing $\text{scale} \times \text{element per lane}$. Block operations decode lanes exactly, apply the operation, and project against a result scale; the reductions (BlockReduceAdd, BlockDotProduct) accumulate exactly and project once. Scaled0p is the block-size-1 form. This is how the draft tracks dynamic range that a bare element format cannot hold.

Conformance and κ

The draft's conformance machinery (§4.4–4.6) has two levels. A *conforming implementation* produces the defined result for every operation, format, and mode. A *κ -approximate implementation* may deviate, but must declare its maximum code-point distance κ from the defined result — measured along the total order, over inputs with finite defined results. Approximation is thus quantified and declared, never silent; SmallFloats implements both levels and measures κ by exhaustive enumeration at registration.

Differences from IEEE 754

Binary16p11se looks like Float16; Binary16p8se looks like BFloat16. They are not (Annex F compares them explicitly):

	Binary16p11se	Float16	Binary16p8se	BFloat16
exponent bias	16	15	128	127
largest finite	65472	65504	3.3762e38	3.3895e38
NaN encodings	1	2046	1	2046
negative zero	no	yes	no	yes

P3109's bias formula makes the exponent range symmetric about the midpoint; IEEE-754's is $2^{(E-1)} - 1$. The two differ by one, so **every code point denotes a value a factor of two away from its IEEE look-alike**. The same bits read as the other type silently halve or double every value. Converting between them is a conversion, never a reinterpretation.

See also

[Core Model, The Standard and this Design, First Session.](#)

Chapter 7

Format Anatomy

7.1 Format Anatomy

Every value in `SmallFloats.jl` belongs to a *format*, and a format is a Julia type. This page explains what that type actually is, why the package draws a hard line between the abstract idea of a format and the concrete type your code uses, and why that line matters for both correctness and speed.

The parametric type

Every format is an instance of the same four-parameter type:

```
Binary{K,P,SGN,EXT}
```

K is the bitwidth, $3 \leq K \leq 16$. P is the precision — significand bits, including the implicit bit. SGN is a `Bool`: signed or unsigned. EXT is a `Bool`: extended (the domain includes $\pm\text{Inf}$) or finite (it does not). Every legal combination of these four parameters is a distinct format, and the package ships all of them: **504 formats** in total.

That is the whole parameter space. There is no fifth axis, and — this is the detail that trips people up when they start comparing formats — K is not an independent budget. It is spent entirely: $K = SGN + E + (P - 1)$, one sign bit if signed, $E = K - P + 1 - SGN$ exponent bits, and $P - 1$ stored significand bits — the leading significand bit is implicit, which is why P enters as $P - 1$. For `Binary8p4se` that is $1 + 4 + 3 = 8$. Raising precision by one bit necessarily takes an exponent bit, because there is nowhere else in the word for the extra bit to come from.

<:, not ==

`Binary{K,P,SGN,EXT}` is an abstract type. It is not what a value's type actually is, and it is not what you should write as an array element type or a constructor target. The concrete type is a *representation* — a draft-named alias like `Binary8p4se` — related to the abstract format by subtyping, not equality:

```
julia> Binary8p4se <: Binary{8,4,true,true}    # `<:`, not `===` — see below
true

julia> formatname(Binary{6,3,true,false})
:Binary6p3sf

julia> SmallFloats.Binary16p6se                # always reachable, unexported
Binary16p6se

julia> format(16, 6, true, true)               # programmatic, runtime params
Binary16p6se

julia> using SmallFloats.Formats                # ...or bring all 504 into scope

julia> Binary16p6se
Binary16p6se
```

`Binary8p4se` is a concrete subtype of `Binary{8,4,true,true}`, but the two type objects are distinct. The distinction is not an implementation accident; it comes from the draft's own vocabulary. A *format* is a datum set and an encoding — an abstract description of which values exist and how they are named. A *representation* is a specific choice of machine word that carries a code point: how many bits, what storage unit, what bit layout. That choice sits below the level of description the standard cares about. `Binary` names the format; `Binary8p4se` is the representation you actually compute with.

The suffix on an alias reads directly: p4 is precision 4, s/u is signed/unsigned, e/f is extended/finite. `Binary8p4se` is 8-bit, precision 4, signed, extended — read the name and you have the four parameters without consulting a table.

Why 504 formats, and why only 120 are exported

Every legal (K, P, SGN, EXT) combination at $K \in 3:16$ gets an alias, because the whole point of the package is that formats are cheap to have — each one is 256 (or fewer) code points and a handful of type-parameter queries, not a hand-written implementation. But 504 identifiers is more than any one script needs in scope at once, so using `SmallFloats` exports only the **120 aliases at $K \leq 8$** — the sizes people actually reach for most: sub-byte and byte-sized formats for quantization, embeddings, and packed storage.

The other 384 (K from 9 to 16) are one opt-in away. You never need to define anything to reach them — they already exist as bindings, just not exported ones:

- `SmallFloats.Binary16p6se` — fully qualified, always available regardless of what's exported.
- `using SmallFloats.Formats` — brings every one of the 504 aliases into scope at once.
- `format(K, P, SGN, EXT)` — the programmatic route, for when the parameters are runtime values rather than something you can spell as a literal type name.

Format introspection

Every format answers a fixed set of queries about its own structure, drawn from the draft's Group M, in both a Julia-flavored spelling and the draft-named form:

```
julia> bitwidth(Binary8p4se), precision(Binary8p4se), expbias(Binary8p4se)
(8, 4, 8)

julia> MaxFiniteOf(Binary8p4se)
Binary8p4se(224.0 ≈ 0x7e)

julia> MinPositiveOf(Binary8p4se)
Binary8p4se(0.0009765625 ≈ 0x01)
```

MinFiniteOf, MinNormalOf, MaxSubnormalOf, expbitwidth, trailingsigbits, issigned, isextended, and their draft-named counterparts (BitwidthOf, PrecisionOf, ExponentBiasOf, ...) round out the set.

The important property of every one of these queries is that it is a pure function of the type parameters — nothing about a particular value changes the answer — so it works identically whether you hand it a type or a value:

```
julia> x = Binary8p4se(1.6);

julia> BitwidthOf(x), PrecisionOf(x), SignednessOf(x), DomainOf(x)
(8, 4, true, true)

julia> MaxFiniteOf(x)
Binary8p4se(224.0 ≈ 0x7e)
```

bitwidth(x) and bitwidth(typeof(x)) are the same query, and both fold to a compile-time literal — there is no runtime cost to asking a value what kind of format it lives in.

The abstract format as an array element type

Because `Binary{K,P,Σ,Δ}` is abstract, declaring an array with it as the element type gives you a non-isbits element type, and every element is boxed:

```
julia> isbitstype(eltype(Vector{Binary{8,4,true,true}}(undef, 3)))
false

julia> isbitstype(eltype(Vector{Binary8p4se}(undef, 3)))
true
```

Nothing about this errors. Every operation still returns the right answer, because the package normalizes the abstract format at every constructor and entry point it can see. What you silently lose is the entire performance story: table gathers, zero-allocation scalar calls, specialized dispatch — all of it depends on the element type being concrete. This is the sharper edge of the type distinction above: comparing the concrete alias with the abstract format makes their different roles visible, while the abstract-array mistake costs you real performance and says nothing at all.

`similar(a, T)` normalizes the request for you, but `Vector{T}(undef, n)` cannot be intercepted at the call site. Always spell the alias, or use `format(K, P, Σ, Δ)` when the parameters are runtime values.

See also

[Values, Code Points, and Conversion](#) explains how a value stores a code point, how decode recovers its exact datum, and how construction differs from raw code-point access.

Chapter 8

Values, Code Points, and Conversion

8.1 Values, Code Points, and Conversion

This page develops the identity of a Binary value: how it is constructed, read, stepped, classified, and converted. Read [First Session](#) first if the distinction between datum and code point is new.

Three ways in, two ways out

A Binary value is an immutable wrapper around a code point. There are three constructors and two accessors:

```
julia> x = Binary8p4se(1.6)           # construct from a Real: projects (rounds)
Binary8p4se(1.625 ≈ 0x45)

julia> Binary8p4se(0x45)             # construct from a UInt8 CODE POINT (validated)
Binary8p4se(1.625 ≈ 0x45)

julia> Convert(Binary8p4se, RNE_SN, 3) # explicit conversion, any mode
Binary8p4se(3.0 ≈ 0x4c)

julia> decode(x)                     # the exact datum, on the format's carrier
1.625

julia> codepoint(x)                  # the code point (extends Base.codepoint)
0x45
```

decode never rounds — it is a lookup from name to meaning. Construction from a Real is the only place rounding happens on the way in.

Important: Unsigned means code point, at every width

The First Session already showed you `Binary8p4se(0x02)` and `Binary8p4se(2)` landing on different values. The rule generalizes to every integer width and every format, and it is worth stating as a standalone hazard because it never raises an error — it just silently gives you the wrong datum if you meant a number:

```
julia> Binary8p4se(0x02), Binary8p4se(2)
(Binary8p4se(0.001953125 ≈ 0x02), Binary8p4se(2.0 ≈ 0x48))
```

‘Unsigned’ means code point — at every width

`T(0x02)` constructs code point `0x02`; `T(2)` projects the numeric value two. Signedness of the integer type is what distinguishes them, so the rule holds for `UInt8`, `UInt16` and any other Unsigned: `Binary16p6se(0x0002)` is code point 2, not the value 2.0. Use `Convert` when intent should be unmistakable.

Out-of-range codes throw for $K < 8$ (`Binary5p3sf(0x20)` is an error). Round-tripping is `T(codepoint(x)) == x` — see [What is, and is not, a round-trip here](#) for the direction that does *not* invert. The unchecked constructor is an internal kernel primitive rather than part of the exported interface.

Convert versus convert

`Convert` (capitalized, from the spec-named register) is the explicit projection you saw above: `Convert(T, p, value)`, numeric for every integer. Base's lowercase `convert` is a different, narrower operation that Julia's promotion machinery calls, and it is value-preserving on Unsigned input rather than code-point reading:

```
julia> Binary8p4se(0x02)
Binary8p4se(0.001953125 ≈ 0x02)
```

Keep the two apart: `T(x::Unsigned)` reads a code point; `Convert(T, p, x)` always projects a numeric value, whatever its type. When you want projection and the argument might be an Unsigned, reach for `Convert`, not the constructor.

What is, and is not, a round-trip here

`decode` and `codepoint` are not inverses of each other — both are read-only accessors on the same `x`, one returning the datum, the other the identity it is stored under. Neither writes a code point, so neither can undo the other.

The actual inverse pair is the code-point constructor and `codepoint` (stated above): `T(codepoint(x)) == x`, exact for every valid code point. The other direction fails on purpose — `codepoint(T(v))` for a `Real v` is not an inverse of anything, because `T(v)` *projects* `v` to the nearest datum first, so you recover that datum's code point, not a code point that would reconstruct `v`. Only the code-point round-trip is exact both ways; the `Real` round-trip loses information by design, every time construction rounds.

`decode` is **always exact**, but the type it returns is a property of the format, not of the package: `Float64` for the 432 formats whose exponent range it holds, `Float128` or an exact dyadic carrier for the other 72. `Float64(x)` always returns a `Float64` and rounds where the datum does not fit.

Enumerating a whole format

Small formats are small enough to look at in full. Binary4p2se has 16 code points, sorted here by the draft's total order — which places the single NaN **first**, below $-\text{Inf}$ (§4.12.1), not last as Float64 sorting does:

```
julia> decode.(sort(Binary4p2se.(0x00:0x0f)))    # broadcast the code-point constructor
16-element Vector{Float64}:
 NaN
 -Inf
 -2.0
 -1.5
 -1.0
 -0.75
 -0.5
 -0.25
  0.0
  0.25
  0.5
  0.75
  1.0
  1.5
  2.0
  Inf
```

That listing *is* the format: one NaN, no negative zero, subnormal spacing near zero, binades doubling outward. This is a good way to build intuition for a format before committing to it.

The one-liner generalizes to any of the 504 formats, in either storage order or total order, as a complete (code, value) table. No projection is involved: decode is exact and policy-free, so the table is a function of the format alone.

```

function codetable(::Type{F}; by::Symbol = :code) where {F<:Binary}
    by in (:code, :value) || throw(ArgumentError("by must be :code or :value"))
    U = codeunit_type(F) # UInt8 at K ≤ 8, UInt16 above
    vals = [F(U(c)) for c in 0:(1 << bitwidth(F)) - 1] # validated code-point constructor
    by == :value && sort!(vals) # the draft's total order: NaN first
    [(code = codepoint(v), value = decode(v)) for v in vals]
end

function printcodetable(io::IO, ::Type{F}; by::Symbol = :code) where {F<:Binary}
    w = 2 * sizeof(codeunit_type(F)) # hex digits per code unit
    for r in codetable(F; by)
        println(io, " 0x", string(r.code; base = 16, pad = w), " ", r.value)
    end
end

printcodetable(stdout, Binary4p2se)

```

```

0x00 0.0
0x01 0.25
0x02 0.5
0x03 0.75
0x04 1.0
0x05 1.5
0x06 2.0
0x07 Inf
0x08 NaN
0x09 -0.25
0x0a -0.5
0x0b -0.75
0x0c -1.0
0x0d -1.5
0x0e -2.0
0x0f -Inf

```

Storage order makes the layout visible — positives ascending from 0x00, the sign bit alone (0x08) naming the single NaN, negatives mirroring above it. `by = :value` gives the sorted listing shown earlier. The value column is the format's datum carrier, so the table stays exact for the 72 wide formats too. The reverse table, value → code, *does* need a projection — it is `codepoint(Convert{F, p, x})` over inputs of your choosing, one rounding per entry.

Stepping and classification

Stepping moves to the adjacent code point; classification names what kind of datum you're holding:

```

julia> NextGreaterThan(Binary8p4se(1.0))
Binary8p4se(1.125 ≈ 0x41)

julia> Class(Binary8p4se(0.01)) # draft classification
ClassPosNormal::FPClass = 0x06

```


`NextGreaterThan/NextLessThan` are the draft's Next operations; `nextfloat` and `prevfloat` alias them, so both spellings work. Standard predicates apply too: `isnan`, `isinf`, `isfinite`, `iszero`, `signbit`, `issubnormal`:

```
julia> Binary8p4se(1e9)           # overflow under the default spec → +Inf
Binary8p4se(Inf ≈ 0x7f)

julia> Binary8p4se(-0.0)          # no -0: projects to the single zero
Binary8p4se(0.0 ≈ 0x00)

julia> isnan(Binary8p4se(NaN)), isfinite(Binary8p4se(2.0))
(true, true)
```

zero, one, eps, typemax

The usual numeric-type vocabulary works on every format too — `zero`, `one`, `eps`, `typemin`, `typemax` — alongside the predicates and stepping operations above; they answer from the format's own grid, not from any IEEE assumption about it.

What this leaves open

Every route in on this page that takes a `Real` rounds, and none of them said by what rule. That rule is a *policy* you choose rather than a property of the format: [The Exact-Then-Project Contract](#) establishes that the rounding happens exactly once and at the end, and [Rounding and Saturation](#) develops the two choices that decide where it lands. The query vocabulary used above — `MaxFiniteOf`, `MinPositiveOf`, and the rest — is catalogued in [Formats and Value Queries](#).

Chapter 9

The Exact-Then-Project Contract

9.1 The Exact-Then-Project Contract

Every number that enters a Binary format goes through the same pipeline, regardless of whether it arrives via a constructor, an arithmetic operator, or an explicit conversion. This page explains what that pipeline promises, what "exact" means at each stage, and the handful of cases where the contract refuses rather than bends.

One write path

There is exactly one way a code point gets produced in SmallFloats.jl:

```
exact result → RoundToPrecision → Saturate → code point
```

The operation is computed *exactly* first — as exact real arithmetic, not as a simulation of some intermediate low-precision format — and only then is the result rounded to the target precision and, if it falls outside the representable range, saturated according to the chosen saturation mode. This holds for a scalar Add, for a fused FMA (one rounding, not two), and for a full BlockDotProduct (every lane product and the accumulation exact, projected once at the end). There is no second path that skips a step or rounds twice; every code point the package ever produces came out of this same pipeline.

The consequence worth sitting with: there is no accumulated rounding error *inside* an operation. Rounding error exists only at the single boundary where the exact result crosses into a code point — never before it.

What "exact" means, per carrier

"Exact math" has to happen somewhere concrete, and the carrier it happens on depends on the format. `decode(x)` always returns the format's exact datum, never a rounded approximation of it — but the *type* it returns is a property of the format, not a package-wide constant.

For 432 of the 504 formats, the datum's exponent range fits inside a Float64, so `decode` returns a Float64 and the exact arithmetic runs there. For the other 72 — formats whose exponent range or precision a Float64 cannot hold without rounding — `decode` returns Float128 or an exact dyadic carrier instead, so no precision is lost representing the datum in the first place. `datumcarrier(T)` names which carrier a given format uses, and it is this carrier, not Float64 universally, that the exact-math half of the pipeline runs on.

This distinction matters because `Float64(x)` always returns a Float64 regardless of the format's own carrier, and rounds wherever the datum does not fit — it is a display/interop convenience, not the internal computation path. At the user level, `promotecarrier(T)` names the *promotion* carrier — Float64 for the 432 formats at the ordinary rung, BigFloat for the other 72 — which is the type an operation between a Binary value and an ordinary number lands in.

Symbolic sticky: projecting "just below" a number

Some values are not just numbers — they are the limit of a sequence approaching a number from one side, and a directed rounding mode needs to know which side to land the projection on. `tanh` approaching 1 as its argument grows is the canonical case: the exact mathematical limit is 1, but $\tanh(x) < 1$ for every finite x , so a correct projection under `TowardNegative` must round *down from* 1, not treat 1 as an ordinary carrier value sitting exactly on a code point.

The engine handles this with a `sticky` argument, `sticky ∈ {-1, 0, +1}`, meaning "the true value is the carrier value plus an infinitesimal of this sign." This lets `project` land exactly on the correct neighbor for enclosure endpoints and asymptotes without inventing a fake carrier value that isn't quite 1:

```
julia> using SmallFloats: project

julia> project(Binary8p4se, ProjSpec(TowardNegative(), SatNone()), 1.0; sticky=-1)
Binary8p4se(0.9375 ≈ 0x3f)      # == NextLessThan(1.0): the engine crossed the binade
```

Read this carefully: the carrier value passed in is 1.0, exactly representable in `Binary8p4se`. Without `sticky`, a directed-toward-negative projection of an exactly-representable value is the identity — it would return 1.0 unchanged. `sticky=-1` tells the engine that the true value sits an infinitesimal *below* 1.0, so `TowardNegative` must move to the next lower code point instead of stopping at 1.0 itself — and it does, crossing the binade boundary down to 0.9375, which is exactly `NextLessThan(1.0)`. This is how the package gets asymptotic behavior like $\tanh \rightarrow 1$ exactly right under every rounding mode, rather than approximately right under most of them.

Why Rational inputs are refused, not double-rounded

`Convert` accepts `Binary` values of any format, `Float16/Float32/Float64`, `Float128`, any `Integer`, and `BigFloat` — every type whose values the package can project *exactly* onto the carrier it needs. `Float128` in particular projects directly, preserving all 113 significant bits, which matters when a value sits just above a rounding midpoint: staging through `Float64` first would round it onto the midpoint and then break the tie the wrong way, a genuine double rounding that changes the answer.

`Rational` is conspicuously missing from that list, and that is deliberate rather than an oversight. A `Rational{Int}` can represent a value that no supported carrier holds exactly (a denominator that is not a power of two, for instance), so converting it would require rounding once to reach a carrier and again to reach the format — silently, with no record of it having happened. That would put a second, invisible rounding step inside what looks like a single conversion, which is exactly the guarantee the one-write-path contract exists to prevent.

So `Convert` refuses `Rational` inputs outright, with an error that tells you what to do about it: convert explicitly to a supported carrier first (`Binary8p4se(Float64(π))` for irrational values, or `Binary8p4se(Float64(r))` for a `Rational r`) and thereby own the double rounding as a choice you made, not a default the package made for you.

Why the contract is a single, named function

All of this — exact math, rounding, saturation, and the sticky adjustment for asymptotic limits — lives behind one function, `project`, rather than being reimplemented at each call site that needs a code point. That centralization is what makes the "one write path" claim checkable rather than aspirational: every operation in the package, from a scalar `Add` to a table build to a block reduction, is required to produce its code point by calling through `project`, and the test suite is able to assert this because there is exactly one place the assertion needs to look. A second hand-rolled rounding step anywhere in the codebase would be a bug by definition, not a matter of style.

This is also why the table-gather performance strategy described elsewhere in the documentation costs nothing in correctness. A cached table entry and a fresh scalar computation are guaranteed identical precisely because both routes go through `project` — the table is not a separate, hopefully-equivalent fast path, it is the *same* path, memoized. Nothing about caching results for speed introduces a second contract to keep in sync with the first.

The same reasoning extends to approximate kernels. When a registered approximation replaces the default path for one operation, it is measured against exactly this contract's output — the exact-math-then-project result — rather than against some looser notion of "close enough." The projection contract is not just how the default path stays correct; it is also the fixed reference every deliberate approximation in the package is honestly scored against.

That is what makes a declared κ a fact about the *difference* between two functions, rather than a vague impression of "close enough" — both sides of the comparison ultimately trace back to the same contract this page describes.

What this leaves open

The contract fixes *when* rounding happens — once, at the end — and says nothing about *which way*. That is the projection's job, and it is two independent choices rather than one: which neighbor a between-grid value lands on, and what a result beyond the format's range becomes. [Rounding and Saturation](#) takes those two axes in turn and shows the one case where they interact. Once you can name a policy, [Control Rounding and Overflow](#) is the procedure for testing that it does what you expect at the values where the modes actually disagree.

Chapter 10

Rounding and Saturation

10.1 Rounding and Saturation

A projection answers two independent questions about one exact result. Rounding: *which neighbor*, when the exact value lies between datums? Saturation: *what may the result become*, when it lies beyond them or hits a special boundary case? This page develops each axis separately, then the one place they interact.

Projection is a product of policies

exact result → rounding decision → saturation decision → code point

SmallFloats represents the pair as ProjSpec(rounding, saturation). A name like RNE_SN is ProjSpec(NearestTiesToEven(), SatNone()) — a convenient value, not a fused algorithm. The independence matters in both directions: changing saturation never changes how an in-range value rounds (RNE_SN and RNE_SF agree on every result inside the finite range), while changing rounding can decide whether the range boundary is even reached — the First Session's exact product 384 becomes Inf under RNE_SN but stops at 224 under RTZ_SN, because a mode directed toward zero never crosses the boundary at all.

Deterministic rounding modes

The deterministic modes answer different questions at the same pair of neighboring datums:

Family	Decision
nearest, ties to even	nearest datum; a tie selects the even code lattice choice
nearest, ties away	nearest datum; a tie increases magnitude
toward positive	least representable datum not below the exact value
toward negative	greatest representable datum not above the exact value
toward zero	neighbor no farther from zero
to odd	select the odd endpoint when discarded information is nonzero

Directed rounding describes order, not simply “up” or “down” in magnitude. For a negative exact value, toward positive moves toward zero while toward negative moves away from zero.

Stochastic rounding modes

Stochastic modes make the landing a function of the exact fractional position and a finite random draw. The distribution is part of the mode's definition. Project $2 + 3/64$ into Binary8p4se under RSA_SN(4): the value sits three sixteenths of the way from 2.0 to the next datum, 2.25, and exactly 3 of the 16 possible draws round upward. The landing probability equals the fractional position. With an explicit draw R the projection is deterministic, so the whole distribution can be tested one draw at a time.

The mode's random-bit budget determines the draw range. An N -bit mode accepts $R \in \mathbb{Z} : (2^N - 1)$; out-of-range draws are contract violations rather than silently wrapped values.

Saturation modes

SatFinite clamps results to the finite range. SatPropagate preserves representable infinities under its contract and handles other overflow by its policy. SatNone applies the draft's full domain-, sign-, direction-, and operation-dependent rows.

The name SatNone is easy to misread. It does not mean “do nothing” and it does not mean “never saturate.” It means neither simplifying policy was selected.

Rounding and range interact

An exact positive result beyond the largest finite datum is not yet an overflow. RoundToPrecision runs first and knows nothing about range — it rounds to the target precision and leaves 384 at 384. Saturation then classifies the out-of-range result *with the rounding mode passed in*: under RNE_SN the draft's rows send 384 to Inf, while under RTZ_SN the toward-zero direction forces the finite extreme, 224. Two specs with the same saturation mode, different boundary outcomes — and saturation itself never rounds; it only consults the direction.

This is why the package projects once through a combined spec rather than rounding in one call and clamping in another. Separating them operationally would introduce an extra representable intermediate and could change the defined result.

Construction, conversion, and operations share projection

$T(x)$, $\text{Convert}(T, p, x)$, and $\text{Op}(T, p, \text{operands} \dots)$ all end at the same projection engine. Their difference is how the exact input to projection is obtained and whether the policy is implicit or explicit.

That shared path supports a useful reasoning rule: once you know the exact mathematical value and the ProjSpec, you can reason about construction, conversion, scalar operations, and array table entries in the same way.

Consequences for users

- Name p in reusable numerical code.
- Test ties and both signs when selecting a rounding policy.
- Test both range boundaries and special values when selecting saturation.
- Record stochastic bit budget and entropy source.
- Do not emulate a spec with a sequence of Base rounding and clamp calls.

What this leaves open

Knowing what the two axes decide is not yet knowing that your choice is right: a single test value cannot distinguish the two nearest modes, and a positive one cannot distinguish toward-zero from toward-negative. [Control Rounding and Overflow](#) is the procedure — which cases to test, and why one value is never enough. If your policy is stochastic, its result varies by design, and pinning that variation so a run is reproducible is [Make Stochastic Work Reproducible](#). For the complete grid of predefined names, [Projection Specifications](#).

Chapter 11

Julia's Numeric Model

11.1 Julia Numeric Behavior

Mixing two Binary formats in an arithmetic expression fails. Mixing a Binary value with an ordinary Julia number succeeds, but the result is not a Binary value. Both behaviors are deliberate, and this page explains the reasoning behind each, along with why the obvious third option — a `promote_rule` that produces a wider Binary format automatically — was considered and rejected.

Mixed Binary formats: deliberate failure, not silent widening

Adding two values from different formats does not silently pick a format wide enough to hold both and round into it. It fails, with an error that names the problem:

```
julia> x = Binary8p4se(1.5);

julia> x + Convert(Binary8p4se, RNE_SN, Binary5p3sf(1.0))
Binary8p4se(2.5 ≈ 0x4a)
```

The line above only works because the `Binary5p3sf` value was converted *explicitly* first. Without that conversion, `x + y` for `x::Binary8p4se` and `y::Binary5p3sf` raises a promotion ... failed to change error rather than silently choosing some third format to compute in. This is not a gap in the promotion machinery — it's the package's stated position that mixing formats is a decision the caller makes explicitly, with `Convert`, never one the package makes for you by picking a resolution.

Binary with ordinary numbers: promotes to the promotion carrier

Combining a Binary value with an ordinary Julia number — an `Int`, a `Float64` — does not fail, but it does not produce a Binary value either. It promotes to the format's **promotion carrier**, an ordinary float type, and the result is that ordinary float:

```
julia> x + 2.0
3.5
```

For 432 of the 504 formats — the ones whose exponent range and precision fit inside a `Float64` without rounding — that carrier is `Float64`, which is what the example above returns. For the other 72, the carrier is `BigFloat`: a `Float64` cannot hold a `Binary16p1uf` datum exactly, and promoting into `Float64` anyway would silently turn a perfectly finite value into $\pm\text{Inf}$ with no P3109 operation anywhere in sight. `promotecarrier(T)` names the target carrier for any format, so the promotion is queryable rather than a fact you have to remember.

If you actually want the result back in the original format, project it back explicitly — the package will not do this step for you, for the same reason it won't guess which format to promote two mixed Binary values into:

```
Convert(Binary8p4se, RNE_SN, decode(x) + 2.0)
```

Why widen is not another format

It's tempting to think the fix for mixed-format arithmetic is obvious: define a `promote_rule` between two Binary types that picks the smallest format containing both operands' datum sets, the way `promote_type(Int, Float64)` picks `Float64`. A companion internal study worked this out in detail, and the answer is that this rule cannot be built soundly. Formats trade precision for range inside a fixed bit budget — raising precision by one bit necessarily removes an exponent bit — so two formats at the same bitwidth are usually *incomparable*: neither contains the other's datum set. Measured over the full 504-format grid, a "smallest format containing both" answer exists for only about 60% of format pairs; the other 40% have no common format at all within the $K \leq 16$ grid, so any Binary-valued rule would have to fall back to something else for those pairs anyway.

Worse, even where the smallest containing format does exist, picking it pairwise is not associative — because Julia's `promote_type(a, b, c)` reduces pairwise, a non-associative rule would make the *result format* depend on the order the operands happened to appear in a `+` chain, which is exactly the kind of instability the package refuses to ship. And even in the cases where the rule is well-defined, it is often surprising on its own terms: two formats that agree on everything except whether they carry infinities do not promote to either input's own bitwidth — they need one bit more, because the wider format has to hold both the finite format's largest value *and* an infinity code point, and at the narrower width there is no code point left for both. A silent promotion landing in a format the user never named, at a width neither operand has, is the wrong kind of convenience for a package whose entire premise is that you choose your result format. The full argument, with the exact percentages and the associativity counterexample, is worked out in the promotion-rules study (`docs/other/promoterules.md` in the repository).

Why this differs from Int/Float64 promotion

It helps to be precise about what makes format promotion a different problem from ordinary numeric promotion. `promote_type(Int, Float64)` works because the target, `Float64`, is fixed in advance and holds every `Int` exactly (up to 2^{53}) — the rule never has to ask "which format contains both operands" because there is only one plausible answer and it never runs out of room. Two Binary formats have no such fixed universal target: the format space is a lattice of finite datum sets, not a single always-sufficient supertype, and which format contains which shifts with every combination of precision, range, signedness, and domain. The question `promote_type` answers for ordinary numbers is trivial; the same question for two arbitrary Binary formats is a genuine combinatorial problem, and the promotion-rules study exists because that problem was worked out in full rather than assumed away.

Explicit Convert as the answer

Given that no automatic rule is both sound and predictable, the package's answer is to make the choice explicit every time: `Convert{T, p, x}` moves any value into the format `T` you name, under the projection `p` you choose. There is no ambiguity about which format the computation happens in, no order-dependence, and no case where the rule silently falls back to a carrier for some inputs and a Binary format for others. The lattice of which format contains which is real and occasionally useful to know explicitly, but it belongs in a query you call when you want the answer, not in an operator that decides for you every time two formats meet.

This is the same philosophy that governs every other write path in the package: one explicit projection, chosen by the caller, rather than a rule that infers intent from context. A `promote_rule` between two Binary formats would be exactly the kind of implicit rounding decision the rest of the design goes out of its way to avoid — it would round somewhere other than inside `project`, which is the one invariant the whole package is built not to break. Treating cross-format arithmetic as a decision you make explicitly, rather than a convenience the library infers, is not a limitation bolted on around an otherwise-automatic feature; it is the same rule that governs rounding modes, saturation modes, and result formats applied consistently to one more place where an implicit choice could have hidden.

What this buys you as a caller

The practical upshot is that a `Binary` expression never surprises you about which format it computed in. If every operand in an expression shares a format, the result is that format, under whatever projection was in play. If you mix in an ordinary number, the result predictably drops to the promotion carrier, and you can query `promotecarrier(T)` in advance to know which carrier that will be. And if you try to mix two different `Binary` formats without converting, the error fires immediately rather than letting a silently-widened result propagate three functions deeper before something notices the format changed. All three outcomes are things you can predict from the types alone, before running anything — which is the actual goal a promotion rule is supposed to serve, achieved here by refusing the one case that could not serve it honestly.

See also

[Values, Code Points, and Conversion](#) gives the constructor and `Convert` rules. [Julia Compatibility Register](#) lists the complete Base surface.

Chapter 12

Session State and Reproducibility

12.1 Session State and Reproducibility

Every convenience call in `SmallFloats.jl` — `T(x::Real)` construction and `x + y` arithmetic, which follow the default projection, and calls that omit a format, which follow the default types — reads one of five process-global defaults. This page names the five, shows how they stay coherent with each other, explains why setting one is a global semantic change rather than a local preference, and gives the safe pattern for consuming them.

The five defaults

Five session-wide defaults are readable as `DefaultX()` and settable as `DefaultX!(v)`:

default	initial value	setter accepts
<code>DefaultType</code>	<code>Binary8p2se</code>	any fully-parameterized Binary type
<code>DefaultReturnType</code>	<code>Binary8p2se</code>	any fully-parameterized Binary type
<code>DefaultRoundingMode</code>	<code>NearestTiesToEven()</code>	mode instance or type
<code>DefaultSaturationMode</code>	<code>SatNone()</code>	mode instance or type
<code>DefaultProjection</code>	<code>RNE_SN</code>	a <code>ProjSpec</code> , or <code>(mode, sat)</code>

Coherence in both directions

`DefaultProjection` and its two component defaults, `DefaultRoundingMode` and `DefaultSaturationMode`, are kept coherent no matter which one you change. Setting `DefaultRoundingMode!` or `DefaultSaturationMode!` rebuilds `DefaultProjection` around the new component; setting `DefaultProjection!` directly decomposes it back into both components. The invariant `DefaultProjection() === ProjSpec(DefaultRoundingMode(), DefaultSaturationMode())` holds always, regardless of which of the three setters you used last:

```
julia> DefaultRoundingMode!(TowardZero())
TowardZero()

julia> DefaultProjection()                # followed the component
(TowardZero, SatNone)

julia> DefaultProjection!(RNA_SF)
(NearestTiesToAway, SatFinite)

julia> DefaultRoundingMode(), DefaultSaturationMode() # followed the projection
(NearestTiesToAway(), SatFinite())

julia> DefaultProjection!(RNE_SN)         # restore: these setters are PROCESS-wide
(NearestTiesToEven, SatNone)
```

The defaults are process-global

The session defaults are process-global and they stay set

DefaultProjection! and its siblings mutate state shared by every module in the process. Nothing scopes them for you: set one in a script, a notebook cell, or a REPL session and every later $T(x::\text{Real})$, $a + b$, and $\text{Exp}(x)$ follows it until something sets it back.

The `with_default_*` combinators consume the current value efficiently; they do not scope or restore mutation. Code needing a local policy should call an explicit form such as `Convert(Binary8p4se, RTZ_SF, 1e9)`.

Every example in the manual assumes the pristine (`NearestTiesToEven`, `SatNone`), which is why each demonstration that changes a default restores it explicitly afterward.

DefaultProjection is not merely advisory — `default_projspec` reads it, so the same-format convenience methods ($a + b$, $\text{Exp}(x)$, ...), the Base-register operators, and $T(x::\text{Real})$ construction all follow it:

```
julia> x, y = Binary8p4se(200.0), Binary8p4se(2.0);

julia> x * y                                     # RNE_SN: overflow → +Inf
Binary8p4se(Inf ≈ 0x7f)

julia> DefaultProjection!(RTZ_SF);

julia> x * y                                     # now clamps to MaxFinite
Binary8p4se(224.0 ≈ 0x7e)

julia> DefaultProjection!(RNE_SN)                # restore before moving on — see below
(NearestTiesToEven, SatNone)

julia> Multiply(Binary8p4se, RNE_SN, x, y)       # explicit p is unaffected
Binary8p4se(Inf ≈ 0x7f)
```

Changing the default is a global semantic change

Every caller of the convenience forms — including code in other packages — sees it. Library code that needs a specific projection should name it explicitly rather than rely on the session default.

The with_default_* combinators

To *consume* a default in your own code without paying dynamic-dispatch costs, go through the `with_default_*` combinators — `with_default_type`, `with_default_returntype`, `with_default_projection` — which call `f(default, args...)`:

```
julia> with_default_type((T, x) -> T(x), 1.5)
Binary8p2se(1.5 ≈ 0x41)
```

The speculation-guard cost model

Following the default costs nothing while it holds its initial value: the convenience forms consume it through a speculation guard, the same mechanism the `with_default_*` combinators use, so they compile against the constant and stay allocation-free with concretely inferred results — this is pinned in the test suite, not an incidental optimization. Once you change the default, they cross a function barrier instead: one dynamic dispatch per call, with everything inside that barrier still fully specialized. The explicit forms, `Add(T, p, x, y)`, are unaffected either way — they never read a default in the first place.

Allocation contract

When `f`'s result type does not depend on the default — `with_default_projection` with the formats fixed by the caller is the normal shape — the call is zero-allocation with a concretely inferred result, pinned in the test suite. When the result's type *is* the default (`with_default_type` used as a constructor, where the returned value's type is whatever `DefaultType` currently names), the value is still computed on the specialized path, but it boxes once where it escapes the function barrier — the irreducible cost of returning a runtime-chosen type.

There is no accumulator default

There is no accumulator default

`DefaultAccumulatorType` existed through v0.1.0 and was removed. The wide-precision accumulator used inside a reduction is not a matter of preference: the span filter measures the operands and picks the cheapest carrier that is *provably exact* for them, so the reduction returns the defined result regardless of which carrier that turns out to be. A knob overriding that choice would make `BlockDotProduct` inexact from the default API — a much larger cost than losing a configuration option. To trade speed for nothing else, set `ENV["SmallFloats_Float128"] = "disable"`, which is bit-identical by construction and only affects build/oracle speed.

See also

[Rounding and Saturation](#) explains the policy stored by `DefaultProjection`. [Make Stochastic Work Reproducible](#) shows when to replace implicit state with an explicit RNG or draw.

Chapter 13

Defined, Stochastic, and Approximate Results

13.1 Defined, Stochastic, and Approximate Results

Two runs disagree — is that a defect? Three reasons are legitimate: the policy changed, the policy itself requests randomness, or a named approximation was selected. Anything else is a defect. The provenance test at the end of this page decides which case you are in.

Defined result

A defined result is the value required by the selected operation, result format, and projection specification. For a deterministic spec, the same inputs always produce the same code point.

```
julia> x, y = Binary8p4se(1.625), Binary8p4se(1.75);  
julia> Add(Binary8p4se, RNE_SN, x, y)  
Binary8p4se(3.5 ≈ 0x4e)
```

The exact sum is 3.375. RNE_SN determines the one permitted landing. The scalar path, array path, and any cached function table must agree on code point 0x4e.

Stochastic defined result

A stochastic projection is still defined. Its additional input is a random draw. Once that draw is fixed, the result is deterministic and testable.

```
σ = RSA_SN(4)  
x = 2.0 + 3/64  
[codepoint(SmallFloats.project(Binary8p4se, σ, x; R)) for R in 0:15]
```

For this value, exactly 3 of the 16 draws round upward. The distribution is part of the stochastic mode's contract. Passing an explicit R tests one point in that contract; passing a seeded RNG reproduces a stream of draws.

The result varies because policy requests variation, not because the implementation is imprecise.

Approximate result

An approximate result comes from an implementation registered as intentionally different from the defined operation. It is reachable only by name through the approximation registry.

The registry records k , the maximum distance in code points between the approximate and defined results over the measured input domain. Registration measures that distance and rejects a claimed bound that is too small.

```
κ, exhaustive = measure_kappa(candidate, :Tanh, Binary8p4se,
                              (Binary8p4se,), RNE_SN)
register_approx!(:fast_tanh, :Tanh, Binary8p4se,
               (Binary8p4se,), RNE_SN, candidate; κ)
```

The pair $(\kappa, \text{exhaustive})$ matters. κ says how far the implementation moved; the Boolean says whether the measurement covered the complete finite input space.

The provenance test

When two runs differ, ask these questions in order:

1. Did the operation, result format, or projection specification change?
2. Is the projection stochastic, and did the draw or RNG stream change?
3. Did the call explicitly select a registered approximation?
4. Did process-global defaults change between calls?

If all answers are no, a changed code point is a defect. This decision rule is more useful than treating every difference as “floating-point noise.”

Consequences for experiments

- Record the result format and projection spec with an experiment.
- Record the RNG algorithm and seed for stochastic work.
- Record approximation name, declared κ , measured κ , and whether measurement was exhaustive.
- Export `conformance_dict()` with artifacts that need an auditable numerical contract.
- Compare code points when verifying conformance; compare decoded application metrics when evaluating usefulness. They answer different questions.

Use it

[Make Stochastic Work Reproducible](#), [Register an Approximation](#), and [Read and Export Conformance](#).

See also

[Projection Specifications](#) and [Conformance and Approximation Registry](#).

Chapter 14

Code-Point Algebra

14.1 Formal Code-Point Algebra

Everything the package does to a $\text{Binary}\{K, P, S, D\}$ value is integer arithmetic on its code point. This page states that arithmetic as identities — where the distinguished code points are, how a code decodes to a value, how the two directions invert each other, and how ordering is recovered.

It is the layer *below* [Formats, Aliases & Representations](#). Read that page for what a format is and how to name one; read this page when you need to compute with code points, prove something about them, or check the package against an independent derivation.

Verified, not transcribed

Every identity below was checked against the package over **all 504 formats** — the distinguished codes for each, and the decoding formula at **every positive finite code point** of every format. The check is the script in [Verification Sessions](#); it is the same kind of exhaustive enumeration the test suite runs, and it reported no disagreement.

Source for the derivations and proofs: *Code-point forms for $\text{Binary}\{K, P, S, D\}$* (IEEE P3109 working material). This page carries the results and the verification; the proofs are there.

Parameters

A format is four numbers. The package writes signedness and domain as `Bool`; the algebra writes them as `0/1`, which is what makes the identities one-liners rather than case analyses.

symbol	meaning	value	package spelling
K	bitwidth — bits in the code	given, $3 \leq K \leq 16$	<code>bitwidth(T)</code>
P	precision — significand digits <i>including</i> the hidden bit	given, $1 \leq P \leq K-S$	<code>precision(T)</code>
S	signedness	0 unsigned, 1 signed	<code>issigned(T)</code>
D	domain	0 finite, 1 extended	<code>isextended(T)</code>

Derived quantities

symbol	meaning	formula	package spelling
Q	total code points	2^K	—
H	sign-bit weight; signed-region offset	$2^{(K-1)}$	signmask(T)
L	trailing-significand patterns per exponent block	$2^{(P-1)}$	—
Ebits	exponent-field bitwidth	$K-P+1-S$	expbitwidth(T)
B	exponent bias	$2^{(K-P-S)} = 2^{(Ebits-1)}$	expbias(T)
η	minimum positive datum; the universal value quantum	$2^{(2-P-B)}$	MinPositiveOf(T)
v	minimum positive normal	$2^{(1-B)} = L\eta$	MinNormalOf(T)
σ	maximum positive subnormal	$(L-1)\eta = v-\eta$ (when $P > 1$)	MaxSubnormalOf(T)
$\delta(j)$	ulp of exponent block j	$2^{(j-B-P+1)} = 2^{(j-1)}\eta$	—

The bias is not the IEEE-754 one

P3109 uses $B = 2^{(Ebits-1)}$, where IEEE-754 uses $2^{(Ebits-1)} - 1$. The difference of one is exactly why Binary16p11se is not Float16 — every code point denotes a different value. See [Why Binary16p11se is not Float16](#).

The NaN anchor

This is the structural keystone. The NaN code is

$$N = Q - 1 - S(H - 1)$$

which specializes to $N = Q-1$ (unsigned: all ones) and $N = H$ (signed: sign bit only, nothing else). Both signed and unsigned formats place the positive exceptional endpoint immediately *below* N, so anchoring at N collapses what would otherwise be a signed/unsigned case split into two one-line identities:

$$\begin{aligned} C(+\infty) &= N - 1 && \text{(when } D = 1) \\ C(\text{MaxFinite}) &= N - 1 - D && \text{always} \end{aligned}$$

Worked for Binary8p4se — $K=8$, $P=4$, $S=1$, $D=1$, so $Q=256$, $H=128$, $L=8$, $B=8$:

quantity	formula	value	confirms
N	$255 - 1 \cdot 127$	$128 = 0x80$	Binary8p4se(NaN) $\approx 0x80$
$C(+\infty)$	$N - 1$	$127 = 0x7f$	Binary8p4se(Inf) $\approx 0x7f$
$M = C(\text{MaxFinite})$	$N - 1 - D$	$126 = 0x7e$	MaxFiniteOf(Binary8p4se) $\approx 0x7e$

Decoding: code to value

Split a positive finite code by Euclidean division on the block size L:

$$\begin{aligned} j &= c \div L && \text{the biased exponent (unbiased: } e = j - B) \\ t &= c \bmod L && \text{the trailing significand} \end{aligned}$$

Then the whole positive finite range decodes with **one** formula in each regime, and the two agree at the boundary:

$$\begin{aligned} V(c) &= c \cdot \eta && 1 \leq c < L && \text{(subnormal)} \\ V(c) &= (L + t) \cdot 2^{(j-1)} \cdot \eta && L \leq c \leq M && \text{(normal)} \end{aligned}$$

The normal form is more familiar written as $V(c) = (1 + t/L) \cdot 2^{(j-B)}$, which is the usual hidden-bit reading; the η -unit form above is the one to compute with, because it is integer arithmetic times a single power of two.

Compute the exponent once

$(L + t) \cdot 2^{(j-1)} \cdot \eta$ and $(L + t) \cdot 2^{(j-1+2-P-B)}$ are the same number, but the first overflows Float64 at the wide formats while the second does not. This is not hypothetical — it is exactly the mistake the verification script for this page made on its first run, and it reported 1405 false "defects" in the package before the arithmetic was corrected. Combine the exponents, then `ldexp` once.

Integer image

For positive finite c , $I(c) = V(c)/\eta$ is an integer — the value counted in quanta. Working in I turns questions about values into questions about integers, which is what makes exhaustive proofs about small formats tractable.

Within a binade

Inside one exponent block, code distance *is* ulp distance. For $c_1 \leq c_2$ with $c_1 \div L = c_2 \div L = j$:

$$V(c_2) - V(c_1) = (c_2 - c_1) \cdot \delta(j)$$

Two consequences the package relies on:

- **Midpoints are exact.** The value halfway between two same-binade neighbours is representable in the next precision up, which is what makes tie-breaking a decidable integer question rather than a floating-point comparison.
- **codedistance is meaningful.** The κ registry measures approximation error in *code points* precisely because within a binade that unit is proportional to the value error, and across binades it is the natural relative unit.

Ordering

Raw code comparison is **not** a numerical order for signed formats: negative values occupy the upper half with $C(x) = H + C(|x|)$, so more-negative values have *larger* codes. The fix is a monotone key. For signed numerical codes, excluding NaN:

$$\text{OrderKey}(c) = \begin{array}{lll} Q - c, & H < c < Q & \text{negative values, folded below } H \\ H, & c = 0 & \text{zero} \\ H + c, & 0 < c < H & \text{positive values, above } H \end{array}$$

with $V(c_1) < V(c_2)$ iff $\text{OrderKey}(c_1) < \text{OrderKey}(c_2)$.

The algebra deliberately says nothing about where NaN goes — "*exceptional codes must be handled according to their specified ordering or unordered semantics*". That is the draft's call, and P3109 §4.12.1 places the single NaN **below** $-\text{Inf}$. The package therefore reserves key 0 for NaN and shifts every datum key up by one, so:

$$\text{order_key}(\text{NaN}) = 0 < \text{order_key}(-\text{Inf}) < \dots < \text{order_key}(+\text{Inf})$$

Which is why sort on a Binary vector puts NaN **first**, not last as Float64 sorting does. See [Julia Compatibility](#) for what that means for generic code.

Sign, magnitude, negation

For signed formats and $c \notin \{0, N\}$:

$$\begin{array}{ll} \text{MagCode}(c) = c \bmod H & \text{code of } |V(c)| \\ \text{NegCode}(c) = c \text{ xor } H & \text{code of } -V(c) \end{array}$$

Negation is an involution with **one exception**: zero. There is a single zero code ($c = 0$), and $0 \text{ xor } H = H$, which is the NaN code N in a signed format — not a negative zero. P3109 has no negative zero, and this is where that shows up in the algebra.

Parameter recovery

The format's parameters are recoverable from its own code lattice, which is the basis of the round-trip validation the test suite performs: precision from the block structure, bias from where the normal range begins, signedness from whether the upper half mirrors the lower, and domain from whether $N-1$ decodes to an infinity or to a finite datum.

```
S = [the upper half mirrors the lower]
D = [V(N-1) is infinite]
M = N - 1 - D
```

Checking your own derivation

The formulas above are worth re-deriving if you are implementing against the draft. To check a derivation against this package:

```
julia> using SmallFloats: nan_code, posinf_code

julia> T = Binary8p4se;

julia> K, P = bitwidth(T), precision(T);

julia> S, D = issigned(T) ? 1 : 0, isextended(T) ? 1 : 0;

julia> Q, H, L, B = 1 << K, 1 << (K - 1), 1 << (P - 1), 1 << (K - P - S);

julia> N = Q - 1 - S * (H - 1);

julia> (N == Int(nan_code(T)),
      N - 1 == Int(posinf_code(T)),
      N - 1 - D == Int(codepoint(MaxFiniteOf(T))),
      B == expbias(T))
(true, true, true, true)
```

The exhaustive form of that check — every distinguished code of every one of the 504 formats, plus the decoding formula at every positive finite code point — is in [Verification Sessions](#).

See also

- [Formats, Aliases & Representations](#) — what a format is, and the type-vs-alias distinction.
- [Encoding and Decoding](#) — how the package walks this algebra at speed.
- [Why Binary16p11se is not Float16](#) — the bias difference, stated in these terms.
- [Format Queries](#) — the package's accessors for every symbol named here.

Part IV

Workflows

Chapter 15

Choose a Format

15.1 Choose a Format

Choosing a format is four decisions: how many bits, how to split them between precision and range, whether the quantity is ever negative, and whether an infinity is a meaningful outcome. This page works through each against real data, because the first three trade off against one another and the honest answer depends on what your numbers actually look like.

Ingredients

- A sample of the data (or a plausible generator for it) so you can measure damage before committing.
- `formatname`, `MaxFiniteOf`, `decode` for inspecting candidates.
- `RNE_SF` (or another `ProjSpec`) to project a sample through each candidate.

Recipe

1. Enumerate the format before committing. Small formats are small enough to look at in full. Sorting `Binary4p2se`'s 16 code points shows the subnormal spacing, the binade structure, and the dynamic range in one line:

```
julia> decode.(sort(Binary4p2se.(0x00:0x0f)))
16-element Vector{Float64}:
 NaN
 -Inf
 -2.0
 -1.5
 -1.0
 -0.75
 -0.5
 -0.25
  0.0
  0.25
  0.5
  0.75
  1.0
  1.5
  2.0
  Inf
```

Read past the leading NaN (the draft's total order puts it below $-\text{Inf}$; see [Core Model](#)) to the format itself: subnormal spacing near zero, binades doubling outward, and $\pm\text{Inf}$ occupying two of the sixteen code points. For a full (code, value) table of any format in either ordering, the `codetable` function in [Values, Code Points, and Conversion](#) does this for all 504.

2. At fixed bitwidth, precision trades against range. For a unit-scale Gaussian tensor, RMSE roughly doubles per lost significant bit while `MaxFinite` grows explosively — spend bits on precision unless the data's magnitude actually needs the range:

```
rng = Xoshiro(7); x = randn(rng, 50_000)
for F in (Binary8p5se, Binary8p4se, Binary8p3se, Binary8p2se)
    back = [decode(Convert(F, RNE_SF, xi)) for xi in x]
    println(rpad(formatname(F), 12), " rmse = ", round(sqrt(mean((x .- back).^2)); sigdigits=3),
           "    maxfinite = ", decode(MaxFiniteOf(F)))
end
```

```
Binary8p5se  rmse = 0.0133    maxfinite = 15.0
Binary8p4se  rmse = 0.0266    maxfinite = 224.0
Binary8p3se  rmse = 0.0528    maxfinite = 49152.0
Binary8p2se  rmse = 0.103     maxfinite = 2.147483648e9
```

If your data's dynamic range genuinely exceeds one binade's reach at your precision budget, don't buy that range with a wider element — get it from a block scale instead.

3. Bounded-in-[0,1] quantities want unsigned. A membership degree, a probability, a normalized score — anything that is never negative — is a job for an unsigned finite format, not a signed one spending a code on a sign it never uses. `Binary8p6uf` gives a `[0, 15.5]` range at 1/32 resolution near 1:

```
julia> F = Binary8p6uf;

julia> membership(x, a, b, c, d) = clamp(min((x - a) / (b - a), 1.0), (d - x) / (d - c), 0.0, 1.0);

julia> warm = F(membership(26.0, 15, 20, 24, 28)); # membership of 26 °C in "warm"

julia> hot = F(membership(26.0, 24, 30, 100, 101)); # ... and in "hot"

julia> (decode(warm), decode(hot),
        decode(Minimum(F, RNE_SN, warm, hot)), # warm AND hot
        decode(Maximum(F, RNE_SN, warm, hot)), # warm OR hot
        decode(Multiply(F, RNE_SF, warm, hot))) # product t-norm
(0.5, 0.3359375, 0.3359375, 0.5, 0.16796875)
```

4. Pick finite (f) when Inf would be meaningless. If nothing in your domain should ever legitimately produce or consume an infinity — a scale factor, a bounded score, a probability-like quantity — a finite-domain format spends the code points IEEE-style formats give to $\pm\text{Inf}$ on two more finite datums instead. Reach for e (extended) only when overflow-to-infinity is a real outcome you want to detect downstream.

An unsigned format is not free precision

Going unsigned doubles your positive-side resolution only if the quantity truly never needs a sign. Do not use an unsigned format for anything that can legitimately go negative (gradients, residuals, log-odds) — projecting a negative Real into an unsigned format is not a sign flip, it is a domain violation and produces a different failure than you likely intend. Reach for the fuzzy-inference example above only when the quantity is already bounded in `[0, ∞)` by construction.

What can go wrong

Enumeration intuition does not scale past small K

`decode.(sort(F.(0x00:...)))` is a genuine format-choice tool for K small enough to print — beyond roughly $K = 8$ the code-point space is too large to eyeball, so switch to the RMSE/MaxFinite sweep instead.

RMSE alone hides overflow

A format can have a low RMSE on the bulk of your data and still overflow its tails to Inf. Always check `count(isinf, back) / length(back)` alongside RMSE and max error — see [Quantize a Tensor or Model](#) for the full three-number check.

Next steps

[Quantize a Tensor or Model](#), [Use Blocks to Track Dynamic Range](#).

Chapter 16

Control Rounding and Overflow

16.1 Control Rounding and Overflow

Two questions have to be answered before a computation is reproducible: where a between-grid value lands, and what happens when a result leaves the format's range. This workflow names both in the call rather than inheriting them from session state, and then tests that the choice does what you expect at the cases where the modes actually disagree.

Select the result format

The result format is the first argument to a spec-named operation:

```
F = Binary8p4se
x, y = F(1.625), F(1.75)
```

Do not rely on cross-format promotion to choose a destination. Convert operands explicitly before combining formats.

Select a deterministic projection

Projection specs combine one rounding mode with one saturation mode. The predefined names use ROUND_SATURATION:

```
julia> x, y = Binary8p4se(1.625), Binary8p4se(1.75);

julia> Add(Binary8p4se, RNE_SN, x, y)
Binary8p4se(3.5 ≈ 0x4e)
```

RNE is nearest, ties to even. SN is SatNone, the draft's domain- and direction-dependent boundary policy. Use SF when the application requires a finite clamp. Use SP when representable infinities should propagate according to SatPropagate.

Test a between-grid case

Choose a value that cannot be represented and compare the policies you intend to support:

```
z = 1.6
[(p, decode(Convert(Binary8p4se, p, z))) for p in
 (RNE_SN, RNA_SN, RTP_SN, RTN_SN, RTZ_SN, RT0_SN)]
```

1.6 sits between 1.5 and 1.625, closer to 1.625. Four modes land on 1.625 (both nearest modes, toward-positive, and to-odd); toward-negative and toward-zero land on 1.5 — they agree here only because the value is positive.

That is why one value is not a policy test. This value is not a tie, so the two nearest modes cannot be told apart by it, and it is positive, so toward-zero and toward-negative cannot be told apart either. Include a tie, a negative value, and a value near a binade boundary before concluding anything about a mode.

Test positive and negative overflow

```
julia> w, two = Binary8p4se(200.0), Binary8p4se(2.0);

julia> Multiply(Binary8p4se, RNE_SN, w, two)
Binary8p4se(Inf ≈ 0x7f)

julia> Multiply(Binary8p4se, RNE_SF, w, two)
Binary8p4se(224.0 ≈ 0x7e)
```

Note that `w` holds 192.0, not 200 — 200 ties between 192 and 208 and construction already rounded it. The exact product is therefore 384, past the largest finite datum, 224.

Repeat the check for a negative result: -192×2 gives $-\text{Inf}$ under `RNE_SN` and -224.0 under `RNE_SF`, mirroring the positive case. Repeat it again for an unsigned result format if you use one. Saturation behavior depends on domain, signedness, direction, and the rounded classification; "overflow clamps" is not a sufficient policy description.

Keep policy local

Prefer explicit forms in libraries and experiments:

```
Multiply(F, RNE_SF, x, y)
```

The `with_default_*` combinators efficiently pass the current process-global default to a function; they do not create a scope or restore anything. Use explicit `T` and `p` arguments whenever numerical policy must be local. Process-global setters are convenient for controlled REPL setup, but their effects are persistent and shared by all tasks in the process.

Validate the complete boundary

For small formats, enumerate all code points and test the operation under each supported projection. At minimum assert:

- exact-grid inputs remain exact;
- ties land according to the named rule;
- positive and negative out-of-range results follow the selected saturation;
- NaN and infinity cases follow the operation's P3109 rows;
- scalar and array forms return the same code points.

What can go wrong

SatNone does not mean finite clamp

`SatNone` means that neither of the two simplifying saturation policies was selected. It applies the draft's full boundary rows and can produce an infinity, an extremal finite value, or NaN depending on the case.

A default is state, not a declaration

A default setter changes process-global state. Name the projection in code that must be reproducible or safe under concurrency.

Understand it

[Rounding and Saturation](#) and [The Exact-Then-Project Contract](#).

See also

[Projection Specifications](#).

Chapter 17

Quantize and Measure a Tensor

17.1 Quantize and Measure a Tensor

A quantization choice needs evidence, and one number is not enough evidence. This workflow produces three that answer different questions: root-mean-square error (how much damage on average), maximum absolute error (how bad the worst case is), and overflow or clamp rate (whether the format's range was adequate at all). A format can look good on the first and fail on either of the others.

Use a block workflow instead when the data contains groups with very different scales and one bare element format cannot cover them economically.

Prepare representative data

A quantizer can look excellent on synthetic unit-scale data and fail on the tails of a real tensor. Prefer a stable sample from the actual model or signal. The example uses a fixed seed only to make the documentation reproducible.

```
using SmallFloats, Random, Statistics

rng = Xoshiro(42)
w = randn(rng, 50_000) .* 0.25
```

Project explicitly

Name the projection used for evaluation. RNE_SF makes the choice clear: round to nearest, ties to even, and clamp outside the finite range.

```
F = Binary8p4se
q = [Convert(F, RNE_SF, x) for x in w]
back = decode.(q)
```

Using `F.(w)` is shorter and follows the session default. The explicit form is better in an evaluation because the report remains meaningful if a default changes elsewhere in the process.

Measure error and boundary use

```
function quantization_report(::Type{F}, x; projection = RNE_SF) where {F}
    q = [Convert(F, projection, xi) for xi in x]
    y = decode.(q)
    finite = isfinite.(y)
    (
        format = formatname(F),
        rmse = sqrt(mean((x[finite] .- y[finite]).^2)),
        max_error = maximum(abs.(x[finite] .- y[finite])),
        nonfinite_rate = count(!isfinite, y) / length(y),
        boundary_rate = count(v -> abs(v) == decode(MaxFiniteOf(F)), y) / length(y),
    )
end

quantization_report(Binary8p4se, w)
```

```
(format = :Binary8p4se, rmse = 0.0066, max_error = 0.0622,
 nonfinite_rate = 0.0, boundary_rate = 0.0)
```

Read the four together: average damage is under 1% of the data's scale, the worst single value is off by ten times that, and nothing overflowed — this sample never approached 224, so the range was never the binding constraint.

RMSE describes aggregate damage. Maximum error catches a bad local miss. Nonfinite and boundary rates reveal a range problem that an average can hide. Under SatFinite, overflow becomes an extremal finite value rather than Inf, so the boundary rate is the relevant companion metric.

Compare candidates under the same policy

```
reports = [quantization_report(F, w) for F in
    (Binary8p5se, Binary8p4se, Binary8p3se, Binary8p2se)]
```

Binary8p5se	rmse=0.00356	max=0.0256	nonfinite=0.0	boundary=0.0
Binary8p4se	rmse=0.00662	max=0.0622	nonfinite=0.0	boundary=0.0
Binary8p3se	rmse=0.0132	max=0.0881	nonfinite=0.0	boundary=0.0
Binary8p2se	rmse=0.0261	max=0.162	nonfinite=0.0	boundary=0.0

RMSE roughly doubles per significand bit given up, exactly as the spacing does. For *this* sample the extra range those bits bought is worthless — every format holds the data without overflow, so Binary8p5se wins outright. On a tensor with a wider spread the last two columns would stop reading zero, and the ranking would change. That is the whole reason to report four numbers rather than one.

At fixed bitwidth, increasing P improves local resolution and reduces exponent range. Compare formats on the same sample and projection. Do not mix a format decision with a saturation-policy change in the same experiment.

Choose the narrowest format that satisfies an application threshold for all reported metrics. Record the selected format and projection with the model or dataset artifact.

Validate the choice on held-out data

Repeat the report on data that did not participate in the choice. Also inspect the worst few errors:

```
r = quantization_report(Binary8p4se, w)
q = [Convert(Binary8p4se, RNE_SF, x) for x in w]
err = abs.(w .- decode.(q))
worst = partialsortperm(err, 1:10; rev = true)
[(w[i], decode(q[i]), err[i]) for i in worst]
```

```
(-1.06279, -1.125, 0.06221)
(-1.06092, -1.0, 0.06092)
(-1.06967, -1.125, 0.05533)
```

The worst misses cluster just past 1.0, where Binary8p4se's spacing widens to 0.125 at the binade boundary — a structural property of the format, not an unlucky sample. Knowing *where* the errors live tells you whether a different format would help or whether the data needs rescaling.

A satisfactory average does not excuse systematic failure on rare but important values. If failures cluster by row or channel scale, move to [Use Blocks for Dynamic Range](#).

What can go wrong

MSE can hide the tail

Always report maximum error and boundary/nonfinite rate. A few saturated values can matter greatly while barely moving a global mean.

The constructor hides policy

`F.(x)` is valid, but it uses the current default projection. Use `Convert(F, ρ, x)` in experiments whose policy must be recorded.

Packing is a later decision

First choose and validate the numerical format. Pack the resulting code points only after that choice; packing changes storage, not quantization.

Understand it

[Format Anatomy](#) and [Rounding and Saturation](#).

Chapter 18

Make Stochastic Work Reproducible

18.1 Make Stochastic Work Reproducible

Stochastic rounding varies by design, which makes "did my change break something?" hard to answer unless the randomness is pinned. Two things can be pinned, and they suit different jobs: an RNG seed reproduces a whole stream, while an explicit draw R reproduces one decision without depending on any RNG algorithm at all. Tests usually want the second.

Ingredients

- An explicit `AbstractRNG` (e.g. `Xoshiro(seed)`) passed as `rng`.
- An explicit `R` (the raw stochastic draw) for tests that must be bit-exact.
- `Random.seed!` when intentionally using Julia's task-local default RNG.
- The `projection` keyword on scalar `rand`/`randn` forms.

Recipe

1. Prefer an explicit rng stream over the task-local default. Passing any `AbstractRNG` first gives a stream independent of global state:

```
julia> rand(Xoshiro(1), Binary8p4se)
Binary8p4se(0.0703125 ≈ 0x21)

julia> rand(Xoshiro(8), Binary8p4se, 4)
4-element Vector{Binary8p4se}:
 Binary8p4se(0.40625 ≈ 0x35)
 Binary8p4se(0.021484375 ≈ 0x13)
 Binary8p4se(0.75 ≈ 0x3c)
 Binary8p4se(0.28125 ≈ 0x31)
```

If you rely on the implicit task-local generator instead, `Random.seed!` controls it exactly as for any other type:

```
julia> Random.seed!(1234); rand(Binary8p4se)
Binary8p4se(0.3125 ≈ 0x32)

julia> Random.seed!(1234); rand(Binary8p4se)      # same seed, same draw
Binary8p4se(0.3125 ≈ 0x32)
```

2. A stochastic projection draws from the same rng as the value draw. One underlying Float64 draw, landed on the grid four different ways — the seed fixes the draw, the projection decides the landing:

```
julia> rand(Xoshiro(2), Float64)           # the underlying uniform draw
0.0022505868897625403

julia> rand(Xoshiro(2), Binary8p4se)       # default: floor (stays in [0,1))
Binary8p4se(0.001953125 ≈ 0x02)

julia> rand(Xoshiro(2), Binary8p4se; projection = RTP_SN) # ceiling
Binary8p4se(0.0029296875 ≈ 0x03)

julia> rand(Xoshiro(2), Binary8p4se; projection = RNE_SN) # nearest
Binary8p4se(0.001953125 ≈ 0x02)

julia> rand(Xoshiro(2), Binary8p4se; projection = RSA_SN(8)) # stochastic
Binary8p4se(0.001953125 ≈ 0x02)
```

Because the stochastic projection draws its random bits from the *same* rng as the uniform draw, seeded streams stay reproducible end to end — you do not need a second seed for the rounding decision.

3. For a test, fix the raw draw R instead of an rng. This makes a single projection exactly reproducible without depending on an RNG's internal algorithm at all — the right tool when you want to test both edges of a stochastic rounding decision deterministically:

```
julia> σ = RSA_SN();                       # ≡ ProjSpec(StochasticA{8}(), SatNone())

julia> Add(Binary8p4se, σ, Binary8p4se(2.0), Binary8p4se(0.03125); rng = Xoshiro(1))
Binary8p4se(2.0 ≈ 0x48)

julia> Add(Binary8p4se, σ, Binary8p4se(2.0), Binary8p4se(0.03125); R = 0)
Binary8p4se(2.0 ≈ 0x48)      # smallest draw: rounds down

julia> Add(Binary8p4se, σ, Binary8p4se(2.0), Binary8p4se(0.03125); R = 255)
Binary8p4se(2.25 ≈ 0x49)    # largest draw: rounds up
```

4. Omitted RNGs follow Julia. SmallFloats does not keep a second package-wide RNG. An operation with no rng argument uses `Random.default_rng()`, so `Random.seed!` controls that task-local stream. Prefer an explicit RNG in library code and tests so ownership of the stream remains visible.

What can go wrong

Array and 'i' forms always use the defaults

`rand(T, dims)`, `randn!(A)`, and friends always use the default projection (floor for `rand`, nearest + `SatFinite` for `randn`) — there is no projection keyword on the array forms. For arrays under another projection, draw scalars explicitly: `[rand(rng, T; projection = ρ) for _ in 1:n]`.

R must lie in '0:2^N-1'

For an N-bit stochastic mode, an out-of-range explicit R is a contract violation, not a silently wrapped value — pick R from the mode's actual bit budget (`RSA_SN(8)` takes `R ∈ 0:255`).

Nearest/ceiling can return exactly 1.0 from 'rand'

The default floor projection is the contract-keeper for rand: it guarantees results stay in $[0, 1)$. Opting into RTP_SN or RNE_SN can return exactly 1.0 (mass near the top rounds up) — only opt out of the default if your code can tolerate that.

Next steps

[Choose a Format](#), [Quantize a Tensor or Model](#).

Chapter 19

Use Blocks for Dynamic Range

19.1 Use Blocks for Dynamic Range

A narrow element format has good resolution over a small range. When your data spans more range than that, the usual fix — a wider element format — spends bits on exponent that you wanted for precision. A block spends them once per group instead: one shared scale, B elements, each representing $\text{scale} \times \text{element}$. This page builds one, measures what it buys on data that actually needs it, and shows the one staging mistake that silently throws the benefit away.

Scale $\times$ element

Small formats have small ranges. A `Block{B, FS, FE}` fights that by sharing one scale, in format `FS`, across B elements in format `FE`; the value each element *represents* is $\text{scale} \times \text{element}$, not the element alone. Construct one directly from a scale and its elements:

```
julia> b = Block(Binary8pluf(4.0), Binary8p4se(1.5), Binary8p4se(-0.75),  
                Binary8p4se(2.0), Binary8p4se(0.5))  
Block{4, Binary8pluf, Binary8p4se}(Binary8pluf(4.0  $\approx$  0x82), (...))
```

Read the type parameters: block size 4, scale format `Binary8pluf` (unsigned — scales are non-negative — with the huge range that precision 1 buys), element format `Binary8p4se`. The represented values are $4.0 \times 1.5 = 6.0$, $4.0 \times -0.75 = -3.0$, and so on.

ConvertToBlockMaxAbsFinite

Building a block by hand from a scale you already chose is one path; the more common one is to quantize a tuple of values and let the draft's algorithm pick the scale for you, from the maximum finite `|element|`, then project each element against it:

```
julia> ConvertToBlockMaxAbsFinite(Binary8pluf, Binary8p4se, RNE_SN, RNE_SN,  
                                (Binary8p4se(100.0), Binary8p4se(-12.0), Binary8p4se(0.5), Binary8p4se(3.0)))  
Block{4, Binary8pluf, Binary8p4se}(Binary8pluf(64.0  $\approx$  0x86), (...))
```

Read the arguments in order: scale format, element format, a projection spec for the scale, a projection spec for the elements, then the values. The algorithm picked scale 64 from the largest-magnitude element (100) and projected each of the four elements against that scale.

ConvertFromBlock

Read a block back out as plain values in a format of your choosing: each $\text{scale} \times \text{element}$ is decoded exactly and projected once.


```
julia> ConvertFromBlock(Binary8p3se, RNE_SN, b) # decode scale×elem, project
(Binary8p3se(6.0 ≈ 0x4a), Binary8p3se(-3.0 ≈ 0xc6), Binary8p3se(8.0 ≈ 0x4c), Binary8p3se(2.0 ≈ 0x44))
```

The stage-wide pitfall

ConvertToBlockMaxAbsFinite takes already-Binary elements as input. If you stage your raw Float64 data through the narrow *element* format before quantizing into a block, SatFinite clamps the large values **before** the block scale ever gets a chance to absorb them — and the block buys you nothing. This is the single most consequential mistake with this API, so watch it happen on real data before you write code that makes it.

Each row of *W* here lives at a different scale, spanning roughly 2^{16} across the matrix — exactly the shape MX-style blocking exists for:

```
using SmallFloats, Random, Statistics

rng = Xoshiro(3)
W = randn(rng, 64, 64) .* 2 .* (2.0 .^ (collect(0:63) ./ 4)) # per-row scales

function mx_rows(W, ::Type{FST}, ::Type{FS}, ::Type{FE}, ::Val{B}) where {FST,FS,FE,B}
    n, m = size(W)
    [begin
        # stage through a WIDE-RANGE format so nothing clamps before scaling
        seg = ntuple(k -> Convert(FST, RNE_SF, W[i, (j-1)*B + k]), Val(B))
        ConvertToBlockMaxAbsFinite(FS, FE, RNE_SN, RNE_SN, seg)
    end for i in 1:n, j in 1:m ÷ B]
end

blocks = mx_rows(W, Binary8p2se, Binary8p1uf, Binary8p4se, Val(32))
recon = [let b = blocks[i, (j-1) ÷ 32 + 1]
    decode(b.s) * decode(b.x[(j-1) % 32 + 1])
end for i in 1:64, j in 1:64]
plain = [decode(Convert(Binary8p4se, RNE_SF, w)) for w in W]

relerr(A) = sqrt(mean(((W .- A) ./ max.(abs.(W), 1e-9)).^2))
relerr(plain), relerr(recon)
```

```
(0.616, 0.109) # plain 8p4se clamps the large rows; MX tracks them
```

The plain, unblocked Binary8p4se conversion clamps every row whose scale exceeds what Binary8p4se alone can hold — relative error 0.616. Staging through the wide-range Binary8p2se first, then letting the per-row block scale absorb the range, brings that down to 0.109 — the residual is now just the staging format's own precision.

Stage wide, then block-quantize

ConvertToBlockMaxAbsFinite takes already-Binary elements. If you stage your Float64 data through the *element* format first, SatFinite clamps the large values **before** the block scale can absorb them, and blocks buy you nothing — we measured exactly that (0.617 vs 0.616) with Binary8p4se staging. Stage through a wide-range format (Binary8p2se above); the residual 0.109 here is the staging format's own precision, which bounds what any downstream scheme can keep.

At $B = 32$ with an 8-bit scale the storage cost is 8.25 bits/value.

BlockDotProduct: one rounding, at the end

BlockDotProduct computes every lane product and the accumulation *exactly*, in a wide carrier, and projects only once, at the very end — there is no hidden intermediate rounding lane by lane:

```
rng = Xoshiro(9)
a64, b64 = randn(rng, 32), randn(rng, 32)
qb(v) = ConvertToBlockMaxAbsFinite(Binary8pluf, Binary8p4se, RNE_SN, RNE_SN,
    ntuple(i -> Convert(Binary8p4se, RNE_SF, v[i]), Val(32)))
dq = BlockDotProduct(Binary8p4se, RNE_SN, qb(a64), qb(b64))
(a64'b64, decode(dq))
```

```
(5.0634, 4.5) # difference is input quantization only, never accumulation error
```

The gap between 5.0634 and 4.5 here is entirely the up-front quantization of a64 and b64 into 4-bit-precision elements — accumulation itself contributes nothing, because the lane products and their sum are carried exactly until the single final projection. Reductions besides the dot product follow the same shape: BlockReduceAdd, BlockReduceMultiply.

BlockVector

For many blocks of the same shape, BlockVector stores them in a structure-of-arrays layout rather than as an array of individually boxed Block values — a memory-bound storage form for blocks, the same way packed code-point storage exists for bare values.

Try it

You have raw Float64 data with widely varying per-row magnitude, and you want to quantize it into Block{32, Binary8pluf, Binary8p4se}. Name the staging mistake to avoid, and the one-line fix.

<details> <summary>Answer</summary>

The mistake is converting straight to Binary8p4se, the narrow element format.

Passing those narrowed elements to ConvertToBlockMaxAbsFinite lets SatFinite clamp large rows before the block scale can absorb them. Stage through a wide-range format first, such as Convert(Binary8p2se, RNE_SF, w), and pass those staged elements onward—the `seg = ntuple(k -> Convert(FST, RNE_SF, ...), ...)` line in the demo above.

</details>

Next steps

[Stochastic Rounding, Arrays, Blocks & Packed Storage.](#)

Chapter 20

Run Operations over Arrays

20.1 Run Operations over Arrays

Array operations return exactly what the scalar call would return, element by element — that is the whole semantic story, and it does not change no matter which of the four spellings you use. What differs between them is ergonomics and, on a first call, timing. This page covers the spellings, then the one measurement mistake the caching makes easy.

The two registers

Every scalar operation comes from one of two registers, and the same split carries over to arrays:

- **The spec-named register** exposes every draft operation under its draft name with an explicit result format and spec: `Op(fr, p, operands...)`. This is `Add`, `Multiply`, `Exp`, and the rest of the operation catalog, where you name `p` yourself.
- **The Base register** makes each format an ordinary Julia number under its *default* spec (`RNE_SN` unless you changed it): `+`, `-`, `*`, `/`, `exp`, `sqrt`, `min`, `max`, comparisons — same-format operands only, no silent cross-format promotion.

Both registers exist because they serve different code: library code that must be insensitive to the session default names `p` explicitly through the spec-named register; exploratory code reads naturally through the Base register's operators. Both extend to arrays without any new vocabulary.

Array calls: `Add(T, p, A, B)`, `Exp(T, p, A)`

Every registered operation has array methods with exactly the same argument shape as the scalar form, operands replaced by arrays:

```
julia> A = Binary8p4se.(randn(Xoshiro(7), 4) .* 2)
4-element Vector{Binary8p4se}:
 Binary8p4se(-0.875 ≈ 0xbe)
 Binary8p4se(4.0 ≈ 0x50)
 Binary8p4se(-3.25 ≈ 0xcd)
 Binary8p4se(2.25 ≈ 0x49)

julia> Exp(Binary8p4se, RNE_SN, A)
4-element Vector{Binary8p4se}:
 Binary8p4se(0.40625 ≈ 0x35)
 Binary8p4se(56.0 ≈ 0x6e)
 Binary8p4se(0.0390625 ≈ 0x1a)
 Binary8p4se(9.0 ≈ 0x59)
```

Add(*T*, *ρ*, *A*, *B*) reads the same way, elementwise, for a two-array call.

Understand cold and warm calls

The first deterministic array call for a new (operation, formats, projection) specialization may build a result table. Later calls reuse it. Warm the exact specialization before measuring steady-state performance, and use `table_bytes()` or `empty_tables!()` only when you need to inspect or reset the cache.

This does not create another numerical path. Every table entry is produced by the scalar operation, so scalar and array forms have the same code-point contract. Stochastic operations compute per element because each projection consumes its own draw.

The cache design and ternary tier policy belong to [Function Tables and Array Kernels](#); they are not required to use the array API correctly.

vmap and vmap!

`vmap/vmap!` are the generic, table-aware entry points behind the named array calls above. `vmap!` is the in-place form and is what you want when you already have a destination array to fill rather than a fresh one to allocate:

```
C = vmap(:Add, T, ρ, A, B)
vmap!(C, Val(:Add), T, ρ, A, B)
```

Reach for `vmap!` in a loop that reuses a buffer; reach for the named form (`Add(T, ρ, A, B)`) everywhere else — they run the same kernel.

The Op callable, with map and broadcast

`Op(T, ρ, A)` above already reads like `map`. When you want Julia's own array verbs — `map`, `broadcast` — rather than the named array call, bind the operation, result format, and projection into an `Op` value and use it as a function:

```
julia> A = Binary8p4se.([1.5, 0.25]); B = Binary8p4se.([0.25, 1.5]);

julia> map(Op(:Add, Binary8p4se, RNE_SN), A, B)
2-element Vector{Binary8p4se}:
 Binary8p4se(1.75 ≈ 0x46)
 Binary8p4se(1.75 ≈ 0x46)

julia> Op(:Exp, Binary8p4se, RNE_SN).(A)
2-element Vector{Binary8p4se}:
 Binary8p4se(4.5 ≈ 0x51)
 Binary8p4se(1.25 ≈ 0x42)
```

`Op` is a singleton type — the operation, format, and projection all live in its type parameters — so a call specializes exactly as `Add(Binary8p4se, RNE_SN, x, y)` does and allocates nothing beyond the array `map` or `broadcast` itself builds.

Sorting

Sorting is special-cased for every Binary format: values compare through integer order keys, and vectors of Binary sort with an $O(n)$ **counting sort** installed as the default algorithm. `sort(A)` follows P3109's total order, with the single NaN first (last under `rev=true`):

```
julia> decode.(sort(Binary4p2se.(0x00:0x0f)))    # broadcast the code-point constructor
16-element Vector{Float64}:
 NaN
 -Inf
 -2.0
 -1.5
 -1.0
 -0.75
 -0.5
 -0.25
 0.0
 0.25
 0.5
 0.75
 1.0
 1.5
 2.0
 Inf
```

`TotalOrder(x, y)` exposes the draft's total order directly if you need the comparison itself rather than a sorted array. This NaN placement differs from `Float64` sorting and is intentional.

Try it

Given `A = Binary8p4se.([1.5, 0.25])` and `B = Binary8p4se.([0.25, 1.5])` from above, write the 0p-callable version of `Add(Binary8p4se, RNE_SN, A, B)` using broadcast rather than `map`.

<details> <summary>Answer</summary>

```
Op(:Add, Binary8p4se, RNE_SN).(A, B)
```

This broadcasts the same bound callable used with `map` above and returns the same two-element vector, `[1.75, 1.75]` in `Binary8p4se`.

</details>

Next steps

Use [Blocks for Dynamic Range](#), [Operation Catalog](#), [Arrays](#), [Blocks & Packed Storage](#).

Chapter 21

Pack Values for Storage

21.1 Pack Values for Storage

A `Vector{Binary5p2se}` spends a whole byte on each 5-bit code point — three bits per element wasted, 37% of the array. `PackedVector` stores them at their true width instead. The trade is explicit: storage shrinks by exactly $\text{bitwidth}(F)/8$, and every read pays a small unpack cost, so this is a win when memory or bandwidth binds and a loss when only compute does.

Ingredients

- `PackedVector(v)` built from an existing `Vector{F}`.
- Standard indexing, collect, iteration — `PackedVector` is a full `AbstractVector{F}`.
- `vmap` for computing directly on packed data (it unpacks tiles internally).
- The sizing rule: $\text{bitwidth}(F)$ bits per element, rounded up to whole storage words.

Recipe

1. Build a packed vector from ordinary values. `PackedVector{F}` stores each code point in exactly $\text{bitwidth}(F)$ bits:

```
julia> v = Binary5p2se.(rand(Xoshiro(3), 6) .* 4);

julia> pv = PackedVector(v); sizeof(pv.data)
8           # 6 × 5 bits = 30 bits → one 64-bit word

julia> pv[3] == v[3]
true
```

Six values at 5 bits each is 30 bits, which rounds up to one 64-bit backing word — `sizeof(pv.data)` reports 8 bytes.

2. Index, collect, and iterate like any other vector. `PackedVector` is a full `AbstractVector{F}`:

```
pv[2]
pv[2] = Binary5p2se(0.75)
collect(pv)
```

3. Compute with `vmap` directly on the packed vector. `vmap` accepts a `PackedVector` argument directly, unpacking cache-friendly tiles internally — you do not need to collect first:

```
out = vmap(:Exp, Binary5p2se, p, pv)
```

4. Size larger vectors the same way. The packing ratio is exactly $\text{bitwidth}(F) / 8$ against a plain byte-per-element vector (for $K \leq 8$ formats). At a million elements of a 5-bit format:

```
julia> n = 1_000_000;

julia> model = [Binary5p2se(UInt8(rand(0:31))) for _ in 1:n];

julia> pv = PackedVector(model);

julia> (sizeof(model), sizeof(pv.data))
(1000000, 625000)
```

1,000,000 bytes of plain storage becomes 625,000 bytes packed — $n \times K / 8$ bytes, in general, rounded to whole words.

What can go wrong

There is deliberately no in-place packed arithmetic

The rule is *store packed, compute unpacked*. Don't reach for an in-place packed $+=$ -style operation — it doesn't exist. Compute with `vmap` (which unpacks internally and returns a fresh result) or unpack with `collect` first, compute, then re-pack with `PackedVector` if you want to keep the result compressed at rest.

Packing helps storage bandwidth, not per-element compute speed

Reach for `PackedVector` when memory footprint or bandwidth is the bottleneck (e.g. a large quantized model kept resident). If you are only computing, not storing, an ordinary `Vector{F}` is simpler and does not pay unpacking cost. Its code-unit storage is one byte for `Code8` and two bytes for `Code16`; compare that representation-specific cost with the packed bitwidth before choosing a layout.

Next steps

[Quantize a Tensor or Model, Choose a Format.](#)

Chapter 22

Interoperate with Float16 and BFloat16

22.1 Interoperate with Float16 and BFloat16

Move data between SmallFloats formats and IEEE Float16/BFloat16 without silently changing its value by a factor of two.

Ingredients

- `Convert(T, p, x)` — the only safe path between formats with different exponent bias.
- `datumsexact` to check whether the two datum sets permit an exact round trip.
- Never reinterpret across formats with different bias — see the warning below.

Recipe

1. Understand the mismatch before you touch either format. The $K \leq 16$ grid contains two formats that look like Float16 and BFloat16 but are not — the difference is the kind that produces correct-looking numbers rather than errors:

	Binary16p11se	Float16	Binary16p8se	BFloat16
precision P	11	11	8	8
exponent bias	16	15	128	127
largest finite	65472	65504	3.3762e38	3.3895e38
NaN encodings	1	2046	1	2046
negative zero	no	yes	no	yes
domain is a parameter	yes (se/sf)	no	yes	no

P3109 defines the bias so the exponent range is symmetric about the format's midpoint; IEEE-754 defines it as $2^{(E-1)} - 1$. The two differ by one, so **every datum in the format is a factor of two away from its IEEE counterpart** — Binary16p11se and Float16 share a significand layout and disagree on the value of every code point.

2. Always convert — never reinterpret. Convert is a real conversion: Float16(1.5) becomes Binary16p11se(1.5), unchanged in value:

```
julia> using SmallFloats.Formats
```

```
julia> Convert(Binary16p11se, RNE_SF, Float16(1.5))    # a conversion: 1.5 ↦ 1.5
Binary16p11se(1.5 ≈ 0x4200)
```

```
julia> reinterpret(Binary16p11se, Float16(1.5))        # the SAME BITS: 1.5 ↦ 0.75
Binary16p11se(0.75 ≈ 0x3e00)
```

Float16(1.5) has the bit pattern 0x3e00; read as a Binary16p11se, those same bits denote **0.75**, because the exponent bias differs by one. Nothing errors, nothing warns, and the answer is off by a factor of two.

3. Check datumsexact when you need to know if the round trip is lossless. The package converts to and from both Float16 and BFloat16 exactly, wherever the datum sets permit — datumsexact tells you where that holds so you don't have to assume.

4. Remember the other differences are draft decisions, not bugs. A single NaN encoding means there is no payload to propagate and no quiet/signalling distinction — the 2045 code points IEEE spends on NaN payloads are ordinary datums here. No negative zero means -0.0 and 0.0 are the same code point. The domain parameter (e for extended with $\pm\text{Inf}$, f for finite) has no IEEE analogue: Binary16p11sf spends the two infinity code points on finite values instead. None of these need "fixing" before converting — Convert already accounts for all of them.

What can go wrong

'reinterpret' silently returns the wrong value, off by 2×

reinterpret(Binary16p11se, Float16(1.5)) gives 0.75, not 1.5 — the same bit pattern, read under a different exponent bias. This is the single most dangerous operation in cross-format interop: it never errors, so a pipeline that accidentally reinterprets instead of converting produces plausible-looking numbers that are systematically wrong by a factor of two. Always use Convert, never reinterpret, when moving between a SmallFloats format and Float16/BFloat16.

If you want IEEE semantics, use Float16/BFloat16 directly

Binary16p11se and Binary16p8se are not drop-in replacements for Float16/BFloat16 — they exist for the P3109 draft's own reasons (symmetric exponent range, single NaN, no negative zero, parametric domain). If your downstream consumer specifically needs IEEE bit patterns, keep the data in Float16/BFloat16 and convert only at the SmallFloats boundary.

Next steps

[Choose a Format, Read & Export the Conformance Declaration.](#)

Chapter 23

Produce a Conformance Declaration

23.1 Read and Export Conformance

A result is only as interpretable as the record of what produced it. The conformance declaration is that record in machine-readable form: which formats and operations exist, which draft revision they implement, which table specializations this session actually built, and which approximations — if any — were in play. Export it beside an artifact and a later reader can tell exactly what your numbers came from without rerunning anything.

Ingredients

- `conformance()` — the live declaration as a Julia value.
- `conformance_report()` — a printed, human-readable form.
- `conformance_dict()` — a plain nested `Dict{String,Any}` for JSON/TOML serialization.
- `draft_revision()` — the P3109 draft revision the package implements.
- `draft_identity()` — the retained draft source and its reproducibility digest.
- Any JSON/TOML writer your project already depends on.

Recipe

1. Inspect the live declaration. `conformance()` returns the formats, operations, mode vocabulary, the table specializations actually instantiated this session, and every registered approximation. `conformance_report()` prints the same information for a human to read at the REPL.

2. Export it as a plain Dict for serialization.

```
d = conformance_dict()           # plain nested Dict{String,Any}
(d["package"], length(d["formats"]), length(d["operations"]),
 [s["op"] for s in d["cached_specializations"]])
```

```
("SmallFloats.jl 0.4.0", 504, 52, Any[])
```

All 504 formats and all 52 operations, always — those are properties of the package, not of your session. The fourth entry is the one that varies: it lists the table specializations built so far, and it is empty above only because that session had not yet made an array call. Run one first and the operation appears there. That is what the "live, not static" note below means in practice — snapshot the declaration after the work, not before it.

Serialize `d` with any JSON/TOML writer to attach a machine-readable conformance statement to experiment artifacts — for example, alongside a saved model checkpoint or a benchmark result, so a later reader can confirm exactly which formats and approximations produced it without re-running the session.

3. Record the draft revision alongside it. `draft_revision()` reports which revision of the P3109 draft this package implements — include it next to the conformance dict so an artifact remains interpretable even if the draft itself changes in a later package version.

`conformance_dict()` now includes this structured identity directly. Use `draft_identity()` when an experiment manifest needs the same facts without the session-dependent cache and approximation entries.

4. Confirm registered approximations show up. Any approximation registered with `register_approx!` appears in the conformance declaration automatically — `conformance_report()` lists it alongside the exact operation catalog, so the declaration always reflects the true provenance of a session's results, exact and approximate alike.

What can go wrong

The declaration is live, not static

`conformance()` and `conformance_dict()` reflect the *current* session — including which table specializations have actually been built so far. Call them at the point you want to snapshot (typically right before saving results), not at the top of a script, or the declaration may under-report what the run actually exercised.

A conformance dict does not travel with unregistered approximations

If you registered an approximation in one session and load results from it in a fresh session, `conformance_dict()` in the new session will not know about that approximation unless you re-register it. Save the exported dict itself with the artifact — do not rely on regenerating it later from a differently configured session.

JSON/TOML serialization is your responsibility, not the package's

`conformance_dict()` returns a plain nested Dict; the package does not ship a JSON or TOML writer. Use whatever serialization library your project already depends on — the dict is deliberately unopinionated about the output format.

Next steps

[Register & Measure an Approximation \(\$\kappa\$ \)](#), [Choose a Format](#).

Chapter 24

Register and Measure an Approximation

24.1 Register an Approximation

Register a faster, inexact kernel in place of the bit-exact default, with an honest, measured error bound.

Ingredients

- `measure_kappa(fn, opname, resultformat, operandformats, p)` to measure the worst-case code-point deviation exhaustively.
- `register_approx!(name, opname, resultformat, operandformats, p, fn; κ)` to register the implementation under a declared bound.
- `approx`, `list_approx`, `kappa`, `kappa_measured` to retrieve and inspect registered approximations.
- `unregister_approx!` to remove one.

Recipe

1. Measure κ before declaring anything. `measure_kappa` runs your approximate function against the defined operation over every input and returns the worst code-point distance, exhaustively verified:

```
julia> step2(x) = (r = Exp(Binary8p4se, RNE_SN, x);  
                 isfinite(decode(r)) ? NextGreaterThan(NextGreaterThan(r)) : r);  
  
julia> measure_kappa(step2, :Exp, Binary8p4se, (Binary8p4se,), RNE_SN)  
(2.0, true)           # κ = 2 code points, verified over all 256 inputs
```

2. Register with the measured (or a larger) κ . `register_approx!` takes the same signature plus the implementation and a declared κ :

```
julia> register_approx!(:cheater, :Exp, Binary8p4se, (Binary8p4se,), RNE_SN, step2; κ=1)  
ERROR: ArgumentError: declared κ = 1.0 understates measured κ = 2.0 — registration rejected
```

Declaring $\kappa = 1$ when the measured deviation is 2 is rejected outright — the registry measures your honesty at registration time; you cannot understate the bound.

3. Register with a truthful κ and retrieve it. Using $\kappa=1$ failed above; declaring the true value (or higher) succeeds. The same shape applied to a hardware-style approximation (`hardtanh` standing in for `Tanh`) shows the full round trip:

```
julia> one4 = Binary8p4se(1.0);

julia> hardtanh(x) = Clamp(Binary8p4se, RNE_SN, x, Negate(one4), one4);

julia>  $\kappa$ , exhaustive = measure_kappa(hardtanh, :Tanh, Binary8p4se, (Binary8p4se,), RNE_SN)
(4.0, true) # worst deviation: 4 code points, verified on all 256 inputs

julia> impl = register_approx! (:hardtanh_act, :Tanh, Binary8p4se, (Binary8p4se,),
                             RNE_SN, hardtanh;  $\kappa=4$ );

julia> (kappa(:hardtanh_act), kappa_measured(impl), :hardtanh_act in list_approx())
(4.0, 4.0, true) # declared = measured; conformance_report() now lists it
```

4. Remove a registration when you're done with it.

```
unregister_approx! (:hardtanh_act)
```

What can go wrong

Understating κ is not silently clamped — it throws

You cannot register `:cheater` with $\kappa=1$ "just to see" and get a warning; `register_approx!` throws `ArgumentError` and refuses the registration entirely. There is no partial-registration state to clean up.

NaN-mismatching implementations need $\kappa = \text{NaN}$, not a number

If your approximate kernel disagrees with the defined operation on whether a NaN is produced, you must acknowledge this explicitly by declaring $\kappa = \text{NaN}$ — a numeric κ asserts NaN behavior matches.

Approximate implementations never enter the default API

Registering `:my_fast_exp` or `:hardtanh_act` does not change what `Exp` or `Tanh` return by default — the bit-exact path is always the one `Exp`, `+`, and friends use. Retrieve an approximation explicitly with `approx(:name)` and call it yourself; `list_approx()` enumerates what's registered, and `conformance_report()` lists them alongside the exact catalog.

Next steps

[Read & Export the Conformance Declaration.](#)

Part V

Performance and Examples

Chapter 25

Performance Model

25.1 Performance Model

Nothing in this package makes a result faster by making it different. Tables, carrier selection, packing, and threading all schedule the same exact-then-project computation; the code point never depends on which one ran. That narrows performance work to two questions, in order: can Julia see the format and projection at compile time, and does your signature qualify for a table? Get the first wrong and you measure dynamic dispatch — roughly a microsecond per call, swamping the operation entirely. The second decides whether an array call gathers from memory or computes per element.

Pass formats and policies statically

Give scalar code a concrete format type and, in a hot path, a concrete projection specification. A function barrier is enough when a format arrives at run time:

```
function add_one(::Type{T}, ρ, x, y) where {T<:Binary}
    Add(T, ρ, x, y)
end
```

This does not require every application to hard-code its format. It gives Julia a point at which the type has become known, so calls inside the barrier can specialize. A benchmark that reads a format from a non-const global measures dynamic dispatch as well as the operation and cannot be used to compare codec, projection, or carrier paths.

Let the signature choose the array shape

For a pure projection, an array operation may be a Shape-A table gather or a Shape-B scalar compute loop. Ask the package which one applies:

```
julia> table_policy(:Add, Binary8p4se, Binary8p4se, Binary8p4se, RNE_SN)
(shape = :A, entries = 65536, bytes = 65536, reason = "within the 2^16-entry build band")
```

Read the answer as a budget decision. A same-format binary Add needs one entry per operand pair — $256 \times 256 = 65\,536$ entries at one byte each — which fits the build band, so it gathers. A unary Exp on the same format needs 256 entries and fits easily. Change one format to Code16 and the entry count grows by $256\times$, which is what the budget exists to refuse.

The first Shape-A call may build a table; later calls can reuse it. Benchmark cold construction and warm use separately. Shape B is not an error or a less correct mode: it runs the same defined scalar semantics once per element and may thread for long, pure-spec arrays.

Representation alone does not determine the shape. A Code16 unary or mixed-width signature can fit the current budget, while a larger or stochastic signature may compute. Conversely, an all-Code8 ternary signature can become too large for an eager table. Use `table_policy`, not a K-only rule of thumb.

Stochastic work preserves order

Stochastic projection consumes an RNG stream in index order. The corresponding array paths stay sequential so a seeded run does not depend on thread scheduling. Pass an explicit RNG when reproducibility is part of the experiment; use a pure projection when independent per-element computation is the desired performance case.

Store packed; compute unpacked

`PackedVector{F}` stores code points at `bitwidth(F)` bits per element when that saves space. It is useful for bandwidth and persistence, not an invitation to perform arithmetic directly on packed bits. Kernels unpack the working data they need and preserve the ordinary scalar/array semantics. `BlockVector` provides the analogous layout for collections of same-shaped blocks.

Sorting and measurement discipline

`SmallFloats` has a representation-aware counting-sort path for sufficiently large homogeneous vectors; short inputs deliberately use the stock algorithm because key-space setup can cost more than comparison sorting. NaN is first in the forward P3109 total order and last under reverse order.

Treat timings, allocation counts, and thread scaling as measurements with a scope, not universal package properties. Record the format, operation, projection, input shape, Shape A/B decision, Julia version, thread count, hardware, and cold/warm state. [Performance Evidence](#) holds recorded results; [Benchmark Correctly](#) explains how to obtain new ones without measuring dispatch or cache setup by accident.

Do not use `@fastmath`

`@fastmath` is not a `SmallFloats` optimization strategy. The package's value is that each defined result follows its stated semantic path, including NaN, single-zero, and projection rules. If a workload is slow, make types static, use an appropriate array call, inspect `table_policy`, or reduce storage pressure with `PackedVector`; do not relax the semantics you are trying to study.

See also

[Function Tables and Array Kernels](#), [Performance Evidence](#), and [Benchmark Correctly](#).

Chapter 26

Example Gallery

26.1 Example Gallery

Every worked example in the documentation, in one list. Each title links to the example itself. All are runnable as shown; outputs were captured from real sessions and seeds are fixed, so you can reproduce them exactly.

The examples live on two pages, split by the level they work at rather than by subject:

- **Applied** — the public API only. Format choice, quantization, blocks, stochastic rounding, packing.
 - **Internals** — unexported machinery. Pipeline introspection, exhaustive self-verification, κ measurement, benchmarking. Stable enough to learn from, not covered by semver.
-

Basic

example	what it shows
One value under several rounding modes	deterministic policies compared on one between-grid value
Enumerating a whole format	a 16-point format printed in full, in total order
Saturation in one line each	SatNone / SatFinite / SatPropagate on the same overflow
Random values under a chosen projection	rand/randn land on code points through the engine
Watching RoundToPrecision work	the exact (sign, S, Q) decomposition before saturation sees it
Symbolic sticky	how directed modes land exactly on asymptotes
Tables are the scalar path, memoized	table entry $\equiv$ scalar result, by construction
Order keys	the sign-magnitude fold that comparisons and the counting sort run on

General AI

example	what it shows
Log-odds belief updating	the running belief drifts far from exact while the threshold decision stays correct — order survives quantization when magnitude does not
Fuzzy inference in an unsigned format	Gödel t-norm/s-norm and product t-norm as bit-exact registry ops, so a rule base reproduces to the code point across machines
Search heuristics keep the ordering	200 scores collapse onto 78 code points; argmax and top-k survive
Exhaustive monotonicity audit	proving by enumeration that a format conversion never inverts the total order
κ-safe decision margins	two results more than 2κ code points apart cannot have their comparison inverted — and the exhaustive check of that rule

Machine Learning

example	what it shows
Quantizing a weight tensor	MSE, max error, overflow fraction — the three-number check for any quantization
Picking a format: precision vs range	RMSE and MaxFinite across four $K = 8$ formats on the same data
Stochastic rounding is unbiased	where nearest accumulates a systematic drift and stochastic does not
Verifying a quantizer exhaustively	your own code checked against an independent reference over every input
The κ workflow, including the refusal	registration rejecting an understated κ — the registry measures, it does not trust
Exporting the conformance declaration	what you hand a reviewer

Deep Learning

example	what it shows
MX-style block quantization	shared-scale blocks, and the staging mistake that silently costs accuracy
Quantized dot products, one rounding	all lane products accumulated exactly, projected once at the end
The swamping demo	why a training loop stalls under nearest and does not under stochastic
Activations are 256-byte tables	the Binary8p4se unary table is one exact, reusable finite function; the example records the measurement scope
Packing a quantized model	bit-packed storage at K not a multiple of 8
Proving a fused kernel exact	BlockDotProduct against 512-bit truth
Stochastic rounding, audited	the full-R sweep: every random draw, not a sample
Hard-tanh κ measurement	an accelerator-style activation, measured rather than assumed
Benchmarking without measuring the dispatcher	the measurement post-mortems this package learned from

Verification

example	what it shows
Checking the code-point algebra	every distinguished code of all 504 formats, and the decoding formula at every positive finite code point

Use it

Task-shaped instructions, as opposed to worked demonstrations, begin with [Choose a Format](#). A workflow tells you what to do; an example shows a complete session and interprets what its output means.

Chapter 27

Applied Sessions

27.1 Applied Sessions

Runnable examples spanning basic use, general AI, machine learning, and deep learning. Every output shown was captured from a real session; seeds are fixed so you can reproduce them exactly.

Basic

Two foundational examples live with their explanations

One value, every rounding mode is in [Rounding and Saturation](#), and Enumerating a whole format is in [Values, Code Points, and Conversion](#). Those pages explain the ideas the examples expose, so this collection links to them rather than repeating them.

Saturation in one line each

```
julia> w, two = Binary8p4se(200.0), Binary8p4se(2.0); # MaxFinite is 224
```

```
julia> Multiply(Binary8p4se, RNE_SN, w, two)
Binary8p4se(Inf ≈ 0x7f)
```

```
julia> Multiply(Binary8p4se, RNE_SF, w, two)
Binary8p4se(224.0 ≈ 0x7e)
```

Random values under a chosen projection

`rand` and `randn` take a projection keyword on their scalar `::Type` forms. One underlying `Float64` draw, landed on the format's grid four different ways — the seed fixes the draw, the projection decides the landing:

```
julia> rand(Xoshiro(2), Float64)           # the underlying uniform draw
0.0022505868897625403

julia> rand(Xoshiro(2), Binary8p4se)       # default: floor (stays in [0,1))
Binary8p4se(0.001953125 ≈ 0x02)

julia> rand(Xoshiro(2), Binary8p4se; projection = RTP_SN) # ceiling
Binary8p4se(0.0029296875 ≈ 0x03)

julia> rand(Xoshiro(2), Binary8p4se; projection = RNE_SN) # nearest
Binary8p4se(0.001953125 ≈ 0x02)

julia> rand(Xoshiro(2), Binary8p4se; projection = RSA_SN(8)) # stochastic
Binary8p4se(0.001953125 ≈ 0x02)
```

A stochastic projection draws its random bits from the *same* rng as the uniform draw, so seeded streams stay reproducible. Note the default is the contract-keeper: a nearest or upward projection can return exactly 1.0 (mass near the top rounds up), which is why floor is the default.

The same keyword on `randn` — here the draw is -1.9757...:

```
julia> randn(Xoshiro(6), Binary8p4se)      # default: nearest + SatFinite
Binary8p4se(-2.0 ≈ 0xc8)

julia> randn(Xoshiro(6), Binary8p4se; projection = RTZ_SF) # toward zero
Binary8p4se(-1.875 ≈ 0xc7)
```

Where the saturation half of the default earns its keep: `Binary3p1se` has `MaxFinite` = 1.0, so normal tail draws overflow constantly. Seed 360 draws `z` = 3.326...:

```
julia> randn(Xoshiro(360), Binary3p1se)    # SatFinite clamps the tail
Binary3p1se(1.0 ≈ 0x02)

julia> randn(Xoshiro(360), Binary3p1se; projection = RNE_SN)
Binary3p1se(Inf ≈ 0x03)           # opt out, get the overflow
```

The array and `!` forms (`rand(T, dims)`, `randn!(A, ...)`) always use the defaults; for arrays under another projection, draw scalars: `[rand(rng, T; projection = p) for _ in 1:n]`.

General AI

Log-odds belief updating — and where tiny formats bite reasoning

Classic probabilistic reasoning: accumulate independent evidence as log-likelihood ratios. A sensor with hit rate 0.85 and false-alarm rate 0.30 contributes $\log(0.85/0.30) \approx 1.0415$ per alarm. In Binary8p3se that datum is 1.0 — and the running belief stalls exactly the way the gradient accumulator does in the Deep Learning section:

```
using SmallFloats

function logodds_demo(nalarms)
    llr = Binary8p3se(log(0.85 / 0.30))    # quantized evidence weight
    acc = Binary8p3se(0.0)
    for _ in 1:nalarms
        acc = Add(Binary8p3se, RNE_SN, acc, llr)
    end
    post(l) = 1 / (1 + exp(-l))
    exact = nalarms * log(0.85 / 0.30)
    (decode(llr), round(exact; digits=2), decode(acc),
     round(post(exact); digits=5), round(post(decode(acc)); digits=5))
end
logodds_demo(12)
```

```
(1.0, 12.5, 8.0, 1.0, 0.99966)
# llr datum, exact log-odds, quantized log-odds, exact posterior, quantized posterior
```

The accumulator is wrong by a third (8.0 vs 12.5 — near 8 the ulp is 2.0, so 8 + 1 rounds straight back to 8), yet the *decision* is untouched: both posteriors are ≈ 1 . Threshold decisions need order, not magnitude — which is why coarse formats work in reasoning pipelines far past the point where their arithmetic looks broken. When magnitude does matter, the cures are the same as for gradients: a wider accumulator format, or stochastic rounding.

Fuzzy inference in an unsigned format

Membership degrees live in $[0, 1]$ and are never negative — a job for an unsigned finite format. Binary8p6uf gives a $[0, 15.5]$ range with 1/32 resolution at 1; the classic fuzzy connectives are registry operations (Minimum = Gödel t-norm, Maximum = s-norm, Multiply = product t-norm):

```
julia> F = Binary8p6uf;

julia> trap(x, a, b, c, d) = clamp(min((x - a) / (b - a), 1.0, (d - x) / (d - c)), 0.0, 1.0);

julia> warm = F(trap(26.0, 15, 20, 24, 28));    # membership of 26 °C in "warm"

julia> hot = F(trap(26.0, 24, 30, 100, 101));    # ... and in "hot"

julia> (decode(warm), decode(hot),
        decode(Minimum(F, RNE_SN, warm, hot)), # warm AND hot
        decode(Maximum(F, RNE_SN, warm, hot)), # warm OR hot
        decode(Multiply(F, RNE_SF, warm, hot))) # product t-norm
(0.5, 0.3359375, 0.3359375, 0.5, 0.16796875)
```

Because every connective is a bit-exact registry op, a fuzzy rule base evaluated in Binary8p6uf is *reproducible across machines to the code point* — a property an implementation that substitutes unmeasured approximations cannot promise.

Search heuristics: what quantized scores keep is the ordering

Game-tree and beam search consume evaluation scores only through comparisons. Quantizing 200 heuristic scores to 8 bits collapses them onto 78 distinct code points — yet the argmax survives, and sorting is exact (integer order keys, $O(n)$ counting sort, the single NaN ordered first, deterministically):

```
julia> rng = Xoshiro(21);

julia> scores = randn(rng, 200) .* 3;

julia> q = Binary8p4se.(scores);

julia> (argmax(scores), argmax(decode.(q)), argmax(scores) == argmax(decode.(q)),
        decode(sort(q; rev=true)[1:5]), length(unique(codepoint.(q))))
(140, 140, true, [8.0, 7.5, 6.5, 6.0, 6.0], 78)
```

The two 6.0s are the caveat: quantization creates *ties* the exact scores did not have. TotalOrder's deterministic tie behavior means the search expands the same node on every run — ties change which answer you get, not whether you can reproduce it.

Machine Learning

Quantizing a weight tensor and measuring the damage

```
using SmallFloats, Random, Statistics

rng = Xoshiro(42)
w = randn(rng, 10_000) .* 0.25      # typical trained-weight scale
q = Binary8p4se.(w)                 # project every weight
back = decode.(q)

mean((w .- back).^2), maximum(abs.(w .- back)), count(isinf, back) / length(back)

(4.41e-5, 0.0312, 0.0)                # MSE, max |error|, overflow fraction
```

The max error is exactly half an ulp of the largest binade the data reached — the worst case nearest rounding permits.

Picking a format: precision vs range

For the same Gaussian tensor, $K = 8$ formats trade significand bits against exponent range. RMSE doubles per lost significand bit; MaxFinite grows explosively:

```
rng = Xoshiro(7); x = randn(rng, 50_000)
for F in (Binary8p5se, Binary8p4se, Binary8p3se, Binary8p2se)
    back = [decode(Convert(F, RNE_SF, xi)) for xi in x]
    println(rpad(formatname(F), 12), " rmse = ", round(sqrt(mean((x .- back).^2)); sigdigits=3),
            " maxfinite = ", decode(MaxFiniteOf(F)))
end
```

```
Binary8p5se  rmse = 0.0133  maxfinite = 15.0
Binary8p4se  rmse = 0.0266  maxfinite = 224.0
Binary8p3se  rmse = 0.0528  maxfinite = 49152.0
Binary8p2se  rmse = 0.103   maxfinite = 2.147483648e9
```

For unit-scale data, spend bits on precision; buy range only when your data needs it (or get it from a block scale — see below).

Stochastic rounding is unbiased where nearest is not

Nearest rounding maps 0.30078125 to 0.3125 every single time — a systematic +0.0117 bias. Stochastic rounding is right *on average*:

```
target = 0.30078125                                # v = 3/16 of an ulp above 0.296875
σ = RSA_SN(16)                                     # StochasticA{16} with SatNone
rng = Xoshiro(1)
m = mean(decode(Convert(Binary8p4se, σ, target; rng)) for _ in 1:200_000)
(decode(Binary8p4se(target)), m)
```

```
(0.3125, 0.300765)                                # RNE result vs stochastic mean ≈ target
```

This is the property that makes stochastic rounding matter for accumulating many small contributions (gradients, activations statistics) in a low-precision format.

Deep Learning

MX-style block quantization — and the staging pitfall

Block formats pair a shared scale with narrow elements, tracking dynamic range that a bare element format cannot. Here each row of W lives at a different scale, spanning $\sim 2^{16}$ across the matrix:

```
using SmallFloats, Random, Statistics

rng = Xoshiro(3)
W = randn(rng, 64, 64) .* 2 .* (2.0 .^ (collect(0:63) ./ 4)) # per-row scales

function mx_rows(W, ::Type{FST}, ::Type{FS}, ::Type{FE}, ::Val{B}) where {FST,FS,FE,B}
    n, m = size(W)
    [begin
        # stage through a WIDE-RANGE format so nothing clamps before scaling
        seg = ntuple(k -> Convert(FST, RNE_SF, W[i, (j-1)*B + k]), Val{B})
        ConvertToBlockMaxAbsFinite(FS, FE, RNE_SN, RNE_SN, seg)
    end for i in 1:n, j in 1:m ÷ B]
end

blocks = mx_rows(W, Binary8p2se, Binary8pluf, Binary8p4se, Val{32})
recon = [let b = blocks[i, (j-1) ÷ 32 + 1]
    decode(b.s) * decode(b.x[(j-1) % 32 + 1])
end for i in 1:64, j in 1:64]
plain = [decode(Convert(Binary8p4se, RNE_SF, w)) for w in W]

relerr(A) = sqrt(mean(((W .- A) ./ max.(abs.(W), 1e-9)).^2))
relerr(plain), relerr(recon)
```

```
(0.616, 0.109) # plain 8p4se clamps the large rows; MX tracks them
```

Stage wide, then block-quantize

ConvertToBlockMaxAbsFinite takes already-Binary elements. If you stage your Float64 data through the *element* format first, SatFinite clamps the large values **before** the block scale can absorb them, and blocks buy you nothing — we measured exactly that (0.617 vs 0.616) with Binary8p4se staging. Stage through a wide-range format (Binary8p2se above); the residual 0.109 here is the staging format's own precision, which bounds what any downstream scheme can keep.

At $B = 32$ with an 8-bit scale the storage cost is 8.25 bits/value.

Quantized dot products with one final rounding

BlockDotProduct computes every lane product and the accumulation *exactly*, then projects once:

```
rng = Xoshiro(9)
a64, b64 = randn(rng, 32), randn(rng, 32)
qb(v) = ConvertToBlockMaxAbsFinite(Binary8p4se, Binary8p4se, RNE_SN, RNE_SN,
    ntuple(i -> Convert(Binary8p4se, RNE_SF, v[i]), Val(32)))
dq = BlockDotProduct(Binary8p4se, RNE_SN, qb(a64), qb(b64))
(a64' b64, decode(dq))
```

```
(5.0634, 4.5)      # difference is input quantization only, never accumulation error
```

Why training loops like stochastic rounding: the swamping demo

Accumulate 400 gradient steps of 0.011 into a Binary8p3se accumulator. Under nearest rounding, the moment the accumulator dwarfs the increment, every add rounds back to the accumulator — it **stalls**. Stochastic rounding keeps absorbing the increments in expectation:

```
function accumulate_demo(nsteps, g)
    σ = RSA_SN(16)
    rng = Xoshiro(11)
    acc_rne = Binary8p3se(0.0); acc_sto = Binary8p3se(0.0)
    for _ in 1:nsteps
        acc_rne = Add(Binary8p3se, RNE_SN, acc_rne, Convert(Binary8p3se, RNE_SN, g))
        acc_sto = Add(Binary8p3se, σ, acc_sto, Convert(Binary8p3se, σ, g; rng); rng)
    end
    (nsteps * g, decode(acc_rne), decode(acc_sto))
end
accumulate_demo(400, 0.011)
```

```
(4.4, 0.125, 5.0)  # exact sum, RNE accumulator (stalled!), stochastic accumulator
```

The stochastic result is noisy (5.0 vs 4.4 on this seed) but unbiased; the RNE result is *wrong by 35×* and no amount of steps will fix it. In practice you keep a wider accumulator when you can — and use stochastic rounding when you can't.

Activation functions are 256-byte lookup tables

For pure specs, a unary activation over an 8-bit format is a *finite function* — the whole nonlinearity is one 256-byte table, built bit-exactly through the scalar path and then gathered per element. This is precisely the activation LUT you would burn into an accelerator:

```
using SmallFloats, Random
empty_tables!()
rng = Xoshiro(4)
pre = Binary8p4se.(randn(rng, 4096) .* 2.5)      # pre-activations,  $\pm 2.5\sigma$ 
act = Tanh(Binary8p4se, RNE_SN, pre)              # table-gather kernel

(table_bytes(),
 round(count(v -> abs(decode(v)) == 1.0, act) / 4096; digits=3),
 length(unique(codepoint.(act))),
 decode(NextLessThan(one(Binary8p4se))))
```

```
(256, 0.384, 100, 0.9375)
# LUT size; fraction saturated to  $\pm 1$ ; distinct output codes; last value below 1
```

Two things worth staring at. First, **saturation is severe at this input scale**: 38% of pre-activations land exactly on ± 1 — under RNE, tanh reaches 1 as soon as the true value rounds past the midpoint of the last gap, and the entire saturated tail becomes indistinguishable to the next layer. (Under TowardNegative the asymptote is honored instead: tanh of any finite positive input projects to 0.9375, never 1 — the projection mode is part of the activation design space.) Second, the choice of nonlinearity changes *information retention* in code points, not just shape:

```
actS = Softplus(Binary8p4se, RNE_SN, pre)
(length(unique(codepoint.(actS))), table_bytes())
```

```
(61, 512)      # softplus keeps 61 distinct codes here; two LUTs now cached
```

Auditing an activation for a format is this cheap: enumerate all 256 inputs, look at the output histogram, count what survives.

Packing a quantized model

```
n = 1_000_000
model = [Binary5p2se(UInt8(rand(0:31))) for _ in 1:n]
pv = PackedVector(model)
(sizeof(model), sizeof(pv.data))
```

```
(1000000, 625000) # 8 bits → 5 bits per value; indexing and vmap work directly on pv
```

Chapter 28

Verification Sessions

28.1 Verification Sessions

Internals-level recipes: pipeline introspection, exhaustive verification of your own code, κ measurement, and doctrine-compliant benchmarking. These use unexported names (`SmallFloats.round_to_precision`, `SmallFloats.project`, `SmallFloats.get_table`, `SmallFloats.order_key`, ...), which are stable enough to learn from but are not covered by semantic-versioning guarantees.

Outputs are captured from real sessions.

Basic

Watching RoundToPrecision work

`round_to_precision` returns the draft's exact (sign, S, Q) decomposition before saturation ever sees it:

```
julia> using SmallFloats: round_to_precision

julia> r = round_to_precision(4, 8, NearestTiesToEven(), 2.30078125, 0, 0); # P=4, bias=8

julia> (r.sign, r.S, r.Q, r.S * 2.0^r.Q)
(1, 9, -2, 2.25)          # S·2^Q = 9/4: the nearest P=4 value
```

Symbolic sticky: projecting just below a number

The engine accepts sticky $\in \{-1, 0, +1\}$ meaning "the true value is the carrier plus an infinitesimal of this sign". This is how enclosure endpoints and asymptotes ($\tanh \rightarrow 1^-$) are projected exactly:

```
julia> using SmallFloats: project

julia> project(Binary8p4se, ProjSpec(TowardNegative(), SatNone()), 1.0; sticky=-1)
Binary8p4se(0.9375 ≈ 0x3f)    # == NextLessThan(1.0): the engine crossed the binade
```

Tables are the scalar path, memoized

```
julia> using SmallFloats: get_table

julia> empty_tables!();

julia> tbl = get_table(:Exp, Binary8p4se, Binary8p4se, RNE_SN);

julia> (tbl[Int(0x45) + 1],
        codepoint(Exp(Binary8p4se, RNE_SN, SmallFloats.rawvalue(Binary8p4se, 0x45))),
        table_bytes())
(0x52, 0x52, 256)          # table entry ≡ scalar result; one 256-byte table cached
```

Order keys

Ordering is integer arithmetic: a sign-magnitude fold, monotone with the total order, NaN at the bottom (key 0). This is what comparisons and the $O(n)$ counting sort run on:

```
using SmallFloats: order_key
[(v, order_key(v)) for v in (Binary8p4se(-1.0), Binary8p4se(0.0),
                             Binary8p4se(1.0), Binary8p4se(NaN))]
```

```
-1.0 → 64      0.0 → 129      1.0 → 193      NaN → 0
```

General AI

Ranking safety: an exhaustive monotonicity audit of score conversion

Decision and ranking systems consume scores through *order*. If scores are produced in one format and compared in another, the conversion must be monotone — and formats are small enough to prove it by enumeration rather than trust it. Walk every Binary8p3se datum in total order and check that its Binary8p4se image never goes backward on the integer order keys:

```
using SmallFloats
using SmallFloats: order_key

function monotone_conversion(::Type{From}, ::Type{To}) where {From, To}
    prev = nothing
    U = codeunit_type(From)
    values = [SmallFloats.rawvalue(From, U(c))
               for c in UInt64(0):(UInt64(1) << bitwidth(From)) - 1]
    for v in sort(values)
        isnan(decode(v)) && continue
        g = Convert(To, RNE_SN, v)
        prev != nothing && order_key(g) < order_key(prev) && return false
        prev = g
    end
    true
end
monotone_conversion(Binary8p3se, Binary8p4se)
```

```
true
```

Run this for every (From, To, ρ) triple your system actually uses; it is a few hundred integer comparisons per triple. A ranking pipeline whose conversions all pass this audit cannot invert a preference by changing formats — a guarantee no amount of spot-testing provides.

κ -safe decision margins for approximate evaluators

Search under a compute budget often wants a cheap, approximate evaluation function. The κ registry turns "how approximate?" into a *measured* code-point bound — and code-point bounds compose into a decision rule: **if two defined evaluations differ by more than 2κ code points along the total order, the approximate evaluator cannot invert their comparison.** Verify the rule exhaustively for a $\kappa = 2$ evaluator:

```
using SmallFloats: order_key

fast(x) = (r = Exp(Binary8p4se, RNE_SN, x);
           isfinite(decode(r)) ? NextGreaterThan(NextGreaterThan(r)) : r)
κ, exhaustive = measure_kappa(fast, :Exp, Binary8p4se, (Binary8p4se,), RNE_SN)

function margin_audit(fast, κ)
    codes = [SmallFloats.rawvalue(Binary8p4se, UInt8(c)) for c in 0:255]
    safe = violations = 0
    for a in codes, b in codes
        da, db = Exp(Binary8p4se, RNE_SN, a), Exp(Binary8p4se, RNE_SN, b)
        (isnan(decode(da)) || isnan(decode(db))) && continue
        codedistance(da, db) > 2κ || continue
        safe += 1
        (order_key(fast(a)) < order_key(fast(b))) ==
            (order_key(da) < order_key(db)) || (violations += 1)
    end
    (safe, violations)
end
(κ, exhaustive, margin_audit(fast, κ)...)

```

```
(2.0, true, 49522, 0)
```

49,522 operand pairs clear the 2κ margin, and the approximate evaluator agrees with the defined ordering on every one of them — zero violations, exhaustively. Inside the margin, comparisons are genuinely undecidable at this κ ; a search can treat sub-margin comparisons as ties to expand, or escalate those few nodes to the exact evaluator. Either way the pruning is *provably* sound, with κ measured at registration rather than promised.

Machine Learning

Verifying a quantizer exhaustively against an independent reference

Formats are small enough that "test on a few points" is never necessary. Here every in-range datum of `Binary8p3se` is projected into `Binary8p4se` and checked against a brute-force nearest search under 256-bit arithmetic (distance agreement; the draft's tie rule is the projection engine's job and is pinned elsewhere in the suite):

```
using SmallFloats

function ref_nearest_distance(::Type{T}, x) where {T}
    U = codeunit_type(T)
    last = (UInt64(1) << bitwidth(T)) - 1
    fins = [v for v in (SmallFloats.rawvalue(T, U(c)) for c in UInt64(0):last)
            if isfinite(decode(v))]
    minimum(abs(setprecision(() -> BigFloat(decode(v)) - BigFloat(x), BigFloat, 256))
            for v in fins)
end

function verify_quantizer()
    ok = true
    for c in 0x00:0xff
        x = decode(SmallFloats.rawvalue(Binary8p3se, c))
        (isfinite(x) && abs(x) <= decode(MaxFiniteOf(Binary8p4se))) || continue
        got = Convert(Binary8p4se, RNE_SN, x)
        ok &= abs(decode(got) - x) == Float64(ref_nearest_distance(Binary8p4se, x))
    end
    ok
end

verify_quantizer()
```

```
true
```

This pattern — enumerate the inputs, compare against an independently written reference — is how the entire package is tested, and it is available to *your* quantization code at trivial cost.

The κ workflow, including the part where it says no

Register an approximate kernel and the registry measures it; understate the bound and it refuses:

```
step2(x) = (r = Exp(Binary8p4se, RNE_SN, x);
            isfinite(decode(r)) ? NextGreaterThan(NextGreaterThan(r)) : r)

measure_kappa(step2, :Exp, Binary8p4se, (Binary8p4se,), RNE_SN)
```

```
(2.0, true)           #  $\kappa$  = 2 code points, verified over all 256 inputs
```

```
register_approx!(:cheater, :Exp, Binary8p4se, (Binary8p4se,), RNE_SN, step2;  $\kappa$ =1)
```

```
ERRORR: ArgumentError: declared  $\kappa$  = 1.0 understates measured  $\kappa$  = 2.0 – registration rejected
```

Exporting the conformance declaration

```
d = conformance_dict()           # plain nested Dict{String,Any}
(d["package"], length(d["formats"]), length(d["operations"]),
 [s["op"] for s in d["cached_specializations"]])
```

Serialized with any JSON/TOML writer to attach a machine-readable conformance statement to experiment artifacts.

Deep Learning

Proving your fused kernel exact: BlockDotProduct vs 512-bit truth

The block layer promises *one* projection with exact lane arithmetic. Trust, then verify — against big-float truth, over random blocks with mixed scales and an honest special-value mix:

```
using SmallFloats, Random

rng = Xoshiro(5)
mk() = Block(Binary8p4uf(2.0^rand(rng, -3:3)),
             ntuple(_ -> SmallFloats.rawvalue(Binary8p4se, UInt8(rand(rng, 0:255))), 32))

function verify_dots(trials)
  for _ in 1:trials
    bx, by = mk(), mk()
    lx = [decode(bx.s) * decode(v) for v in bx.x]
    ly = [decode(by.s) * decode(v) for v in by.x]
    (any(!isfinite, lx) || any(!isfinite, ly)) && continue
    truth = setprecision(() -> sum(BigFloat(lx[i]) * BigFloat(ly[i]) for i in 1:32),
                                   BigFloat, 512)
    got = BlockDotProduct(Binary8p4se, RNE_SN, bx, by)
    ref = setprecision(() -> SmallFloats.project(Binary8p4se, RNE_SN, truth), BigFloat, 512)
    codepoint(got) == codepoint(ref) || return false
  end
  true
end
verify_dots(200)
```

```
true
```

The same shape verifies any custom fused kernel you write: compose the truth in BigFloat, project once, compare code points.

Stochastic rounding, audited: the full-R sweep

Every stochastic projection is a deterministic function of the draw R , so its *distribution* is checkable exactly. For $x = 2 + 3/64$ in Binary8p4se the fraction is $v = 3/16$ of an ulp, so StochasticA{4} must round up for exactly 3 of the 16 draws:

```
σ4 = RSA_SN(4) # ≡ ProjSpec(StochasticA{4}(), SatNone())
x = 2.0 + 3/64
count(decode(SmallFloats.project(Binary8p4se, σ4, x; R)) == 2.25 for R in 0:15)
```

```
3
```

Sweeping R like this turns "is my stochastic pipeline unbiased?" from a statistical question into an exhaustive one — the pattern the shipped test suite uses.

An accelerator-style activation, κ -measured: hard-tanh vs Tanh

Hardware activation units often ship piecewise approximations. The κ machinery makes the substitution honest: measure the approximation's worst code-point deviation from the defined activation, exhaustively, and register it under that measured bound. Hard-tanh — `clamp(x, -1, 1)` — as a stand-in for Tanh:

```
one4 = Binary8p4se(1.0)
hardtanh(x) = Clamp(Binary8p4se, RNE_SN, x, Negate(one4), one4)

κ, exhaustive = measure_kappa(hardtanh, :Tanh, Binary8p4se, (Binary8p4se,), RNE_SN)
(κ, exhaustive)
```

```
(4.0, true)          # worst deviation: 4 code points, verified on all 256 inputs
```

Where is it worst? At $x = 1.0$: hard-tanh returns 1.0 while the defined `tanh(1.0)` is 0.75 — four code points along the total order. The registration round-trip, including the conformance surface:

```
impl = register_approx! (:hardtanh_act, :Tanh, Binary8p4se, (Binary8p4se,),
                        RNE_SN, hardtanh; κ=4)
(kappa(:hardtanh_act), kappa_measured(impl), :hardtanh_act in list_approx())
```

```
(4.0, 4.0, true)      # declared = measured; conformance_report() now lists it
```

(`unregister_approx! (:hardtanh_act)` removes it.) Declaring $\kappa=3$ would be *rejected* — understatement is impossible by construction. The decision of whether a $\kappa = 4$ activation is acceptable now belongs to your accuracy budget, not to hope: combined with the margin rule from the General AI section, any two pre-activations whose defined Tanh images sit more than 8 code points apart keep their order under hard-tanh, provably.

Benchmarking without measuring the dispatcher

The package's benchmark doctrine, in one snippet (needs the benchmarking/ environment for Chairmarks). Format types enter as **type parameters**; operands come from Chairmarks' *untimed* setup; and if you retrieve functions reflectively (`getfield(SmallFloats, op)`), pass them through an argument barrier so they specialize:

```
using Chairmarks, SmallFloats, Random
using Statistics: median

function bench_add(::Type{T}) where {T<:Binary}          # T: type parameter, not a global
    U = codeunit_type(T)
    last = (UInt64(1) << bitwidth(T)) - 1
    pool = [SmallFloats.rawvalue(T, U(rand(UInt64(0):last))) for _ in 1:4096]
    @be (rand(pool), rand(pool)) (t -> Add(T, RNE_SN, t[1], t[2]))( )
end

b = bench_add(Binary8p4se)
(round(median(b).time * 1e9; digits=1), median(b).allocs)
```

On the host recorded in the [benchmark report](#), the corresponding in-domain Add row has a 9.0 ns median and zero allocations. Treat that as a check on the methodology, not a portable expected value: the tuple printed by this snippet depends on your CPU, Julia build, and benchmark environment.

The 60× misapplication

The same call with `T` read from a non-const global measures $\sim 1 \mu\text{s}$ — Julia's dynamic keyword dispatch, not this package. Two project post-mortems trace to exactly this mistake; the shipped `benchmarking/benchmarking.jl` asserts specialization (zero warm-path allocation) before it believes any number, and so should yours.

Checking the code-point algebra against the package

The identities on [Formal Foundations](#) are claims about every format, so they are checked over every format — the distinguished codes for all 504, and the decoding formula at every positive finite code point of each.

Two things this exercise teaches beyond the identities themselves. First, the combined-exponent rule: computing $(L+t) \cdot 2^{(j-1)} \cdot \eta$ as two multiplications overflows `Float64` at the wide formats, and the first run of this script reported 1405 "defects" that were all in the checking arithmetic. Second, the oracle has to have the range of the thing it is checking — `BigFloat` with one `ldexp` does; `Float64` does not.

```
using SmallFloats, SmallFloats.Formats
using SmallFloats: _NAMED, nan_code, posinf_code, MaxFiniteOf, rawvalue,
                  codeunit_type, codepoint

bad = 0
for (nm, T) in sort(collect(_NAMED); by = first)
    K, P = bitwidth(T), precision(T)
    S, D = issigned(T) ? 1 : 0, isextended(T) ? 1 : 0
    Q, H, L = 1 << K, 1 << (K - 1), 1 << (P - 1)
    B = 1 << (K - P - S)
    N = Q - 1 - S * (H - 1)           # the NaN anchor
    M = N - 1 - D                     # max positive finite code

    B == expbias(T)                   || (bad += 1)
    N == Int(nan_code(T))              || (bad += 1)
    M == Int(codepoint(MaxFiniteOf(T))) || (bad += 1)
    D == 1 && (N - 1 == Int(posinf_code(T)) || (bad += 1))

    setprecision(BigFloat, 256) do
        for c in 1:M                  # every positive finite code
            v = decode(SmallFloats.rawvalue(T, codeunit_type(T)(c)))
            want = c < L ? ldexp(BigFloat(c), 2 - P - B) :
                    ldexp(BigFloat(L + (c % L)), (c ÷ L) - 1 + 2 - P - B)
            BigFloat(v) == want || (bad += 1)
        end
    end
end
bad
```

```
0          # 504 formats, every distinguished code, every positive finite code point
```

Chapter 29

Performance Sessions

29.1 Performance Evidence

Generated: 2026-08-04 00:56 UTC · Julia 1.12.6 · alderlake (32 logical CPUs, 8 Julia threads) · Float128 paths: enabled · Chairmarks 1.3.1

Source: SmallFloats.jl 0.4.0 · commit ed6a0b7ba6cfdcc0452ac4c8fdab91522ba82fc0 · tree dirty · seed 2026

Reference format for per-operation tables: Binary8p4se under (NearestTiesToEven, SatNone). Every table names its operand class: the scalar-operation tables appear in four variants — all code points (NaN and $\pm\text{Inf}$ sampled), NaN excluded, finite-only, and per-operation in-domain — and every other sampled table uses the all-code-points pool, identified in its note. Times are per call; minima first, medians alongside. Methodology per the recorded benchmark doctrine: type-parameterized barriers, untimed setup, specialization preflight.

Core primitives

The decode/compare/step/classify layer plus the projection engine. Operands: all code points — NaN and $\pm\text{Inf}$ sampled.

operation	min	median	allocs
decode	1.4 ns	1.4 ns	0
order_key	1.4 ns	1.6 ns	0
$x < y$	1.6 ns	1.6 ns	0
TotalOrder	1.6 ns	2.1 ns	0
Class	1.2 ns	2.1 ns	0
NextGreaterThan	1.2 ns	1.8 ns	0
project (RNE·SatNone)	1.6 ns	6.7 ns	0
project (StochasticA[8], R drawn)	1.8 ns	7.3 ns	0

Scalar operations

Each arity is measured under four operand classes — safe args, no NaN/Inf, no NaN, and all code points — so operand-class effects (NaN fast rows, $\pm\text{Inf}$ special rows) are separable.

Unary (30)

Safe args

Finite operands within each operation's safe domain — the fully unmasked per-operation scalar cost. The argument-restricted ops (Sqrt, RSqrt, Log, Log2, LogOnePlus, Recip, Divide, ArcSin, ArcCos, ArcCosh, ArcTanh) draw from explicit per-argument safe-domain predicates; every other op uses finite operand tuples whose defined result is not NaN (oracle-derived). Sorted by median.

operation	min	median	allocs
Negate	4.8 ns	6.9 ns	0
Abs	4.4 ns	7.1 ns	0
Recip	7.9 ns	18.2 ns	0
Sqrt	4.9 ns	18.7 ns	0
Cosh	7.3 ns	20.6 ns	0
RSqrt	9.6 ns	20.7 ns	0
Sin	5.7 ns	21.7 ns	0
Cos	7.3 ns	22.2 ns	0
Exp2	7.3 ns	22.3 ns	0
Exp	7.2 ns	22.5 ns	0
ArcSin	6.0 ns	23.1 ns	0
Tanh	5.6 ns	23.3 ns	0
ExpMinusOne	5.7 ns	23.4 ns	0
ArcCos	5.6 ns	24.4 ns	0
ArcSinPi	5.3 ns	25.6 ns	0
Sinh	5.5 ns	25.6 ns	0
Log	5.7 ns	25.8 ns	0
Log2	6.0 ns	26.6 ns	0
LogOnePlus	5.5 ns	26.7 ns	0
Tan	5.7 ns	26.8 ns	0
ArcCosPi	5.4 ns	26.8 ns	0
CosPi	6.6 ns	28.0 ns	0
SinPi	6.0 ns	28.1 ns	0
ArcTan	4.9 ns	28.4 ns	0
ArcTanh	6.1 ns	30.8 ns	0
TanPi	6.3 ns	31.1 ns	0
ArcTanPi	5.3 ns	31.1 ns	0
ArcCosh	5.6 ns	33.3 ns	0
ArcSinh	5.3 ns	33.3 ns	0
Softplus	14.9 ns	33.4 ns	0

No NaN, Inf args

Operands exclude NaN and $\pm\text{Inf}$; finite datums only (zeros and subnormals kept). Domain-restricted ops still take NaN fast rows on out-of-domain finite operands. Sorted by median.

No NaN args

Operands exclude the NaN code point; $\pm\text{Inf}$ and every finite datum are sampled. Sorted by median.

All code points

Operands drawn uniformly over ALL code points — NaN and $\pm\text{Inf}$ are sampled; medians of domain-restricted ops are diluted by instant NaN rows. Sorted by median. Transcendental rows mix special-row fast returns with enclosure-path evaluations, so these are *scalar-path* costs; bulk unary work routes through 256-byte tables (see Array kernels).

operation	min	median	allocs
ArcCosh	1.6 ns	1.6 ns	0
Log2	1.8 ns	1.8 ns	0
Sqrt	1.8 ns	1.9 ns	0
ArcSinPi	1.9 ns	2.0 ns	0
RSqrt	1.8 ns	2.2 ns	0
Log	1.9 ns	5.5 ns	0
Abs	4.7 ns	6.7 ns	0
Negate	4.6 ns	6.8 ns	0
ArcCosPi	1.8 ns	7.5 ns	0
ArcCos	1.9 ns	16.3 ns	0
ArcSin	1.8 ns	17.2 ns	0
Recip	1.8 ns	17.9 ns	0
Sin	4.7 ns	21.6 ns	0
Cos	7.1 ns	21.6 ns	0
Cosh	7.0 ns	21.6 ns	0
Exp2	7.0 ns	22.2 ns	0
Exp	7.0 ns	22.3 ns	0
Tanh	5.0 ns	23.1 ns	0
ExpMinusOne	5.3 ns	23.3 ns	0
Sinh	5.6 ns	25.7 ns	0
LogOnePlus	1.6 ns	26.3 ns	0
Tan	4.9 ns	26.7 ns	0
CosPi	6.5 ns	27.8 ns	0
SinPi	5.0 ns	27.8 ns	0
ArcTan	5.1 ns	28.3 ns	0
ArcTanh	1.9 ns	28.8 ns	0
TanPi	5.7 ns	30.8 ns	0
ArcTanPi	4.5 ns	30.9 ns	0
Softplus	14.5 ns	33.5 ns	0
ArcSinh	5.0 ns	33.6 ns	0

operation	min	median	allocs
ArcCosh	1.6 ns	1.6 ns	0
Sqrt	1.8 ns	1.8 ns	0
ArcTanh	1.9 ns	1.9 ns	0
ArcSin	1.8 ns	2.0 ns	0
Log	1.9 ns	5.5 ns	0
Abs	4.7 ns	6.7 ns	0
Negate	4.6 ns	7.0 ns	0
ArcSinPi	1.9 ns	7.2 ns	0
ArcCosPi	1.8 ns	7.3 ns	0
Log2	1.8 ns	8.6 ns	0
RSqrt	1.8 ns	9.4 ns	0
TanPi	1.8 ns	12.3 ns	0
Recip	1.8 ns	17.9 ns	0
Cos	1.5 ns	21.7 ns	0
Sin	2.0 ns	21.7 ns	0
Exp2	4.9 ns	22.3 ns	0
Exp	4.5 ns	22.3 ns	0
ExpMinusOne	5.2 ns	22.9 ns	0
Tanh	5.0 ns	23.1 ns	0
ArcCos	1.9 ns	23.8 ns	0
Cosh	5.2 ns	25.0 ns	0
Sinh	5.3 ns	25.7 ns	0
LogOnePlus	1.6 ns	26.2 ns	0
Tan	2.0 ns	26.9 ns	0
SinPi	1.6 ns	27.8 ns	0
CosPi	1.6 ns	27.8 ns	0
ArcTan	5.2 ns	28.2 ns	0
ArcTanPi	4.5 ns	30.9 ns	0
Softplus	5.0 ns	33.5 ns	0
ArcSinh	4.9 ns	33.5 ns	0

operation	min	median	allocs
ArcCosh	1.6 ns	1.6 ns	0
Log2	1.8 ns	5.9 ns	0
Sqrt	1.8 ns	6.2 ns	0
Abs	1.9 ns	6.7 ns	0
Negate	1.8 ns	6.9 ns	0
RSqrt	1.8 ns	9.4 ns	0
ArcSin	1.6 ns	17.1 ns	0
Recip	1.8 ns	18.1 ns	0
Cosh	1.8 ns	21.2 ns	0
Cos	1.5 ns	21.5 ns	0
Sin	1.8 ns	21.9 ns	0
Exp2	1.6 ns	22.2 ns	0
Exp	1.8 ns	22.3 ns	0
ExpMinusOne	1.9 ns	23.1 ns	0
Tanh	1.9 ns	23.4 ns	0
ArcCos	1.9 ns	23.9 ns	0
ArcSinPi	1.8 ns	24.3 ns	0
Log	1.9 ns	25.3 ns	0
Sinh	1.8 ns	25.7 ns	0
LogOnePlus	1.6 ns	26.4 ns	0
Tan	1.8 ns	26.4 ns	0
ArcCosPi	1.8 ns	26.6 ns	0
CosPi	1.6 ns	27.8 ns	0
SinPi	1.4 ns	27.8 ns	0
ArcTan	1.8 ns	28.2 ns	0
ArcTanh	1.8 ns	28.5 ns	0
ArcTanPi	2.0 ns	30.8 ns	0
TanPi	1.7 ns	30.8 ns	0
Softplus	1.8 ns	33.5 ns	0
ArcSinh	1.8 ns	33.6 ns	0

Binary (18)**Safe args**

Finite operands within each operation's safe domain — the fully unmasked per-operation scalar cost. The argument-restricted ops (Sqrt, RSqrt, Log, Log2, LogOnePlus, Recip, Divide, ArcSin, ArcCos, ArcCosh, ArcTanh) draw from explicit per-argument safe-domain predicates; every other op uses finite operand tuples whose defined result is not NaN (oracle-derived). Sorted by median.

operation	min	median	allocs
MinimumMagnitude	4.6 ns	7.1 ns	0
MaximumMagnitude	6.9 ns	7.1 ns	0
Maximum	4.6 ns	7.2 ns	0
Minimum	4.6 ns	7.2 ns	0
MaximumFinite	4.5 ns	7.2 ns	0
MinimumFinite	5.3 ns	7.2 ns	0
MaximumNumber	4.8 ns	7.3 ns	0
MinimumNumber	5.1 ns	7.3 ns	0
MaximumMagnitudeNumber	7.1 ns	7.3 ns	0
CopySign	4.6 ns	7.3 ns	0
MinimumMagnitudeNumber	5.3 ns	7.3 ns	0
Multiply	5.1 ns	7.5 ns	0
Add	6.8 ns	9.0 ns	0
Subtract	7.4 ns	9.2 ns	0
Divide	6.7 ns	18.7 ns	0
Hypot	7.7 ns	28.0 ns	0
ArcTan2	6.5 ns	33.1 ns	0
ArcTan2Pi	9.1 ns	35.6 ns	0

No NaN, Inf args

Operands exclude NaN and $\pm\text{Inf}$; finite datums only (zeros and subnormals kept). Domain-restricted ops still take NaN fast rows on out-of-domain finite operands. Sorted by median.

operation	min	median	allocs
CopySign	4.8 ns	7.2 ns	0
Maximum	4.7 ns	7.2 ns	0
Minimum	4.4 ns	7.2 ns	0
MinimumMagnitude	5.0 ns	7.2 ns	0
MaximumFinite	5.3 ns	7.2 ns	0
MaximumMagnitude	6.9 ns	7.2 ns	0
MinimumFinite	4.9 ns	7.3 ns	0
MaximumMagnitudeNumber	7.1 ns	7.4 ns	0
Multiply	5.2 ns	7.4 ns	0
MaximumNumber	5.1 ns	7.4 ns	0
MinimumNumber	4.9 ns	7.5 ns	0
MinimumMagnitudeNumber	4.7 ns	7.6 ns	0
Add	6.6 ns	9.0 ns	0
Subtract	7.5 ns	9.3 ns	0
Divide	2.5 ns	18.4 ns	0
Hypot	7.7 ns	28.0 ns	0
ArcTan2	6.8 ns	33.1 ns	0
ArcTan2Pi	4.2 ns	35.9 ns	0

No NaN args

Operands exclude the NaN code point; $\pm\text{Inf}$ and every finite datum are sampled. Sorted by median.

operation	min	median	allocs
MaximumMagnitude	5.4 ns	7.2 ns	0
CopySign	4.8 ns	7.2 ns	0
Maximum	4.7 ns	7.2 ns	0
Minimum	4.8 ns	7.2 ns	0
MinimumNumber	4.8 ns	7.3 ns	0
MaximumNumber	4.9 ns	7.3 ns	0
MaximumFinite	5.6 ns	7.4 ns	0
Multiply	5.1 ns	7.4 ns	0
MinimumMagnitude	5.2 ns	7.4 ns	0
MinimumFinite	5.1 ns	7.5 ns	0
MaximumMagnitudeNumber	5.5 ns	7.5 ns	0
MinimumMagnitudeNumber	4.6 ns	7.5 ns	0
Add	6.5 ns	9.1 ns	0
Subtract	6.7 ns	9.2 ns	0
Divide	2.5 ns	18.3 ns	0
Hypot	5.2 ns	28.0 ns	0
ArcTan2	6.7 ns	33.0 ns	0
ArcTan2Pi	7.3 ns	35.4 ns	0

All code points

Operands drawn uniformly over ALL code points — NaN and $\pm\text{Inf}$ are sampled; medians of domain-restricted ops are diluted by instant NaN rows. Sorted by median.

operation	min	median	allocs
MinimumMagnitude	1.8 ns	7.1 ns	0
MaximumMagnitude	1.9 ns	7.2 ns	0
Maximum	1.6 ns	7.2 ns	0
MaximumFinite	5.5 ns	7.2 ns	0
MinimumFinite	4.6 ns	7.2 ns	0
CopySign	1.8 ns	7.3 ns	0
Minimum	1.6 ns	7.3 ns	0
MinimumMagnitudeNumber	4.6 ns	7.3 ns	0
MinimumNumber	4.8 ns	7.3 ns	0
MaximumMagnitudeNumber	5.6 ns	7.3 ns	0
MaximumNumber	4.9 ns	7.4 ns	0
Multiply	1.6 ns	7.5 ns	0
Add	3.4 ns	9.2 ns	0
Subtract	4.0 ns	9.4 ns	0
Divide	2.2 ns	17.9 ns	0
Hypot	2.2 ns	27.5 ns	0
ArcTan2	3.0 ns	33.0 ns	0
ArcTan2Pi	3.4 ns	35.5 ns	0

Ternary (3)

Safe args

Finite operands within each operation's safe domain — the fully unmasked per-operation scalar cost. The argument-restricted ops (Sqrt, RSqrt, Log, Log2, LogOnePlus, Recip, Divide, ArcSin, ArcCos, ArcCosh, ArcTanh) draw from explicit per-argument safe-domain predicates; every other op uses finite operand tuples whose defined result is not NaN (oracle-derived). Sorted by median.

operation	min	median	allocs
Clamp	5.1 ns	8.1 ns	0
FMA	7.5 ns	9.6 ns	0
FAA	7.4 ns	9.9 ns	0

No NaN, Inf args

Operands exclude NaN and $\pm\text{Inf}$; finite datums only (zeros and subnormals kept). Domain-restricted ops still take NaN fast rows on out-of-domain finite operands. Sorted by median.

operation	min	median	allocs
Clamp	5.1 ns	8.1 ns	0
FMA	8.2 ns	9.6 ns	0
FAA	7.5 ns	9.7 ns	0

No NaN args

Operands exclude the NaN code point; $\pm\text{Inf}$ and every finite datum are sampled. Sorted by median.

operation	min	median	allocs
Clamp	5.0 ns	8.1 ns	0
FMA	5.0 ns	9.7 ns	0
FAA	4.3 ns	9.8 ns	0

All code points

Operands drawn uniformly over ALL code points — NaN and $\pm\text{Inf}$ are sampled; medians of domain-restricted ops are diluted by instant NaN rows. Sorted by median.

operation	min	median	allocs
Clamp	1.6 ns	8.0 ns	0
FMA	3.9 ns	9.6 ns	0
FAA	3.7 ns	9.7 ns	0

Sensitivity studies

How the scalar costs move with the format and with the projection specification.

Format sensitivity

Same three binary ops across formats; Binary8pluf exercises the wide-exponent-spread escalations, small-K formats the tiny-table regime. Operands: all code points — NaN and $\pm\text{Inf}$ sampled.

operation	min	median	allocs
Add(Binary8p4se)	3.4 ns	9.0 ns	0
Divide(Binary8p4se)	2.3 ns	18.3 ns	0
Multiply(Binary8p4se)	1.8 ns	7.4 ns	0
Add(Binary16p8se)	14.6 ns	77.8 ns	0
Divide(Binary16p8se)	11.0 ns	24.5 ns	0
Multiply(Binary16p8se)	9.5 ns	13.2 ns	0
Add(Binary16p5se)	136.6 ns	698.3 ns	34
Divide(Binary16p5se)	428.8 ns	1.81 μs	63
Multiply(Binary16p5se)	94.1 ns	96.2 ns	0
Add(Binary16pluf)	17.1 ns	21.8 ns	0
Divide(Binary16pluf)	1.5 μs	1.61 μs	81
Multiply(Binary16pluf)	9.8 ns	13.2 ns	0
Add(Binary3plse)	3.7 ns	7.2 ns	0
Divide(Binary3plse)	2.5 ns	6.7 ns	0
Multiply(Binary3plse)	1.8 ns	6.2 ns	0

Projection by rounding/saturation mode

`project(Binary8p4se, p, x)` over the mode vocabulary (stochastic budgets $N = 8$). Operands: all code points — NaN and $\pm\text{Inf}$ sampled.

operation	min	median	allocs
NearestTiesToEven	1.6 ns	6.6 ns	0
NearestTiesToAway	2.6 ns	7.6 ns	0
TowardPositive	2.5 ns	6.0 ns	0
TowardNegative	2.5 ns	6.0 ns	0
TowardZero	2.3 ns	5.3 ns	0
ToOdd	1.8 ns	6.6 ns	0
StochasticA[8]	1.8 ns	7.2 ns	0
StochasticB[8]	1.8 ns	7.2 ns	0
StochasticC[8]	1.6 ns	7.0 ns	0
RNE · SatFinite	1.8 ns	6.3 ns	0
RNE · SatPropagate	1.8 ns	6.4 ns	0

Kernels and storage

Array-shaped work: the table-gather and compute kernels, the ternary bitwidth policy, sorting, table builds, blocks, and packed/converted storage.

Array kernels (vmap)

Warm caches: table specializations prebuilt, so table rows measure the gather; scalar-loop rows measure the full compute pipeline per element. The ternary row here is Binary8p4se (K=8, always the compute path); see the next section for how the ternary bitwidth policy behaves across K. Operands: all code points — NaN and $\pm\text{Inf}$ sampled.

operation	min	median	allocs	per element
vmap unary (table gather), n=65536	8.32 μs	8.68 μs	0	0.13 ns/elem — 7.55 Gelem/s
vmap binary (table gather), n=65536	16.35 μs	16.67 μs	0	0.25 ns/elem — 3.93 Gelem/s
vmap ternary (scalar loop), n=65536	135.26 μs	207.61 μs	42	3.17 ns/elem — 0.32 Gelem/s
vmap binary stochastic (scalar loop), n=65536	642.39 μs	680.83 μs	0	10.39 ns/elem — 0.1 Gelem/s
vmap unary through PackedVector, n=65536	80.72 μs	85.78 μs	773	1.31 ns/elem — 0.76 Gelem/s

Ternary bitwidth tiers (FMA/FAA)

FMA/FAA/Clamp are total functions on $2^{(K1+K2+K3)}$ code points, but that count spans 512 B (K=3) to 16 MiB (K=8), so the array kernel tables small operand formats eagerly, tables mid-size ones adaptively (after enough elements amortize the build; not shown here — see the adaptive-cache gate in `test/ternary_opt.jl`), and always runs the scalar compute kernel at K=8, threaded above a size cutoff when `Threads.nthreads() > 1`. Each tier's optimized row is paired with a scalar-loop baseline (policy Refs forced off around the measurement, restored after) so the win is visible per tier; this process has 8 Julia threads. Same reference format, p, and operand pool discipline as Array kernels above.

operation	min	median	allocs	per element
FMA K=4 (eager table), n=65536	19.47 μs	20.72 μs	0	0.32 ns/elem — 3.16 Gelem/s
FMA K=4 (eager table), scalar-loop baseline, n=65536	920.36 μs	972.5 μs	0	14.84 ns/elem — 0.07 Gelem/s
FMA K=6 (eager table), n=65536	24.07 μs	25.15 μs	0	0.38 ns/elem — 2.61 Gelem/s
FMA K=6 (eager table), scalar-loop baseline, n=65536	999.05 μs	1.01 ms	0	15.41 ns/elem — 0.06 Gelem/s
FMA K=8 (compute), n=65536	131.64 μs	209.46 μs	42	3.2 ns/elem — 0.31 Gelem/s
FMA K=8 (compute), scalar-loop baseline, n=65536	996.15 μs	1.01 ms	0	15.48 ns/elem — 0.06 Gelem/s
FMA K=8 (compute), threaded [8t], n=65536	133.02 μs	211.11 μs	42	3.22 ns/elem — 0.31 Gelem/s
FMA K=8 (compute), sequential [1t], n=65536	987.69 μs	1.0 ms	0	15.31 ns/elem — 0.07 Gelem/s

Sorting (64 K values)

Counting sort is installed as the default algorithm for Binary vectors. Operands: all code points — NaN and $\pm\text{Inf}$ sampled.

operation	min	median	allocs
sort! (counting sort via defalg), n=65536	159.58 μs	161.56 μs	2
sort! (stock comparison sort), n=65536	1.34 ms	1.35 ms	3
sort! rev=true (counting sort), n=65536	160.9 μs	165.64 μs	2

Table builds (oracle + projection, Float128-first)

Cold cache per sample (empty_tables! in untimed setup); JIT pre-warmed. The warm-hit column is the steady-state cost of get_table when the specialization is already cached (min / median). Table entries enumerate every code point by construction (NaN and $\pm\text{Inf}$ included).

operation	min	median	allocs	warm hit
Exp(8p4se) (256 entries)	5.62 μs	6.01 μs	1	65.9 ns / 68.2 ns
Tanh(8p4se) (256 entries)	6.05 μs	6.15 μs	1	66.1 ns / 67.6 ns
Add(8p4se $\times$ 8p4se) (64 K entries)	79.47 μs	158.88 μs	44	62.3 ns / 66.9 ns
Divide(8p4se $\times$ 8p4se) (64 K)	131.26 μs	143.89 μs	44	65.4 ns / 67.3 ns
Add(8p1uf $\times$ 8p1uf) (64 K, wide-spread)	2.57 ms	3.83 ms	536594	64.5 ns / 66.2 ns

Block and scaled operations

Elements Binary8p4se, scales Binary8p1uf, B = 32. Operands: all code points — NaN and $\pm\text{Inf}$ sampled.

operation	min	median	allocs	per lane
BlockAdd (B=32)	964.9 ns	1.03 μs	35	32.28 ns/lane
BlockDotProduct $\rightarrow$ 8p4se (B=32)	320.0 ns	399.0 ns	3	12.47 ns/lane
BlockReduceAdd $\rightarrow$ 8p4se (B=32)	45.0 ns	493.0 ns	0	15.41 ns/lane
ConvertToBlockMaxAbsFinite (B=32)	3.75 μs	3.98 μs	111	124.34 ns/lane

Conversions and packed storage

Operands: all code points — NaN and $\pm\text{Inf}$ sampled.

operation	min	median	allocs	per element
T:: <uint8) code-point="" constructor<="" td=""><td>1.2 ns</td><td>1.2 ns</td><td>0</td><td></td></uint8)>	1.2 ns	1.2 ns	0	
rawvalue (unchecked kernel route)	1.2 ns	1.2 ns	0	
T:: <float64) (projects)<="" constructor="" numeric="" td=""><td>2.5 ns</td><td>7.9 ns</td><td>0</td><td></td></float64)>	2.5 ns	7.9 ns	0	
Convert 8p4se → 8p3se (scalar)	1.8 ns	6.7 ns	0	
Float64 → 8p4se (project)	1.8 ns	6.6 ns	0	
PackedVector pack, n=65536	32.45 μs	34.75 μs	3	0.53 ns/elem
PackedVector unpack (collect), n=65536	28.86 μs	29.04 μs	3	0.44 ns/elem

Parallelism and allocation

The two properties the doctrine states in prose and this report previously left unmeasured.

Threading

Both variants are warmed before either is measured, and the switch is a Ref read inside an already-compiled branch — not a second compilation, which is the harness failure the benchmark doctrine records. Binary16p8se's binary table would be 2^{32} entries, so table_for declines and the Shape-B compute kernel is the *only* array path that format has; that is what threading it buys. The build rows separate two different wins: Multiply is cheap per entry and was dominated by a per-entry dynamic dispatch (now fixed by the _fill_*! function barrier), while Exp is dominated by MPFR enclosure work and is where threading itself pays. This process has 8 Julia threads.

operation	min	median	allocs
Shape-B compute (Binary16p8se) Add, n=65536 – 8 threads	2.69 ms	5.14 ms	597767
Shape-B compute (Binary16p8se) Add, n=65536 – sequential	16.19 ms	18.31 ms	597725
build Exp(Binary16p8se) (64 K entries) – 8 threads	14.57 ms	16.63 ms	1654196
build Multiply(Binary8p4se ²) (64 K entries) – 8 threads	64.46 μ s	132.86 μ s	44
build Exp(Binary16p8se) (64 K entries) – sequential	74.51 ms	77.37 ms	1764746
build Multiply(Binary8p4se ²) (64 K entries) – sequential	445.35 μ s	451.85 μ s	2

Allocation by rung

Warm-path bytes for one call of each operation class, per carrier rung. The **exact selections are zero at every rung, unconditionally** — a selection returns one of its operands, so there is nothing to escalate into, and a nonzero reading there is plumbing rather than arithmetic (this is the condition preflight aborts on). Arithmetic is *recorded, not gated*: it allocates exactly when the operand spread exceeds the carrier's exact range and evaluation escalates to MPFR, which happens at rung 1 too. The enclosure ladder allocates at rungs 2 and 3 by construction. Times are not meaningful for these calls and are omitted.

format	rung	carrier	selection	add	fma	ladder
Binary8p4se	1	Float64	0 B	0 B	0 B	0 B
Binary16p8se	1	Float64	0 B	0 B	0 B	0 B
Binary16p5se	2	Float128	0 B	0 B	0 B	2544 B
Binary16p1uf	3	Dyadic	0 B	0 B	0 B	2968 B
Binary3p1se	1	Float64	0 B	0 B	0 B	0 B

All numbers from this machine/run; absolute values vary by host. Regenerate with `julia --project=benchmarking -t auto benchmarking/benchmarking.jl`; the thread count matters for the Threading table and is recorded in the header above.

Part VI

Reference

Chapter 30

Formats and Values

30.1 Formats and Value Queries

This page is the lookup contract for format names, representation types, and format queries. For the conceptual model, read [Format Anatomy](#); for a worked introduction, use [First Session](#).

Naming grammar

Binary K p P (s|u) (e|f)

```

  Binary K p P (s|u) (e|f)
      |  |  |  |  |
      |  |  |  |  └─ extended (Inf) or finite domain
      |  |  |  └─── signed or unsigned
      |  |  └────── significand precision, including the implicit bit
      |  └────────── total bitwidth, 3 through 16
      └──────────
  
```

Field	Values	Meaning
K	3-16	total bitwidth
P	1-K	significand precision, including the implicit bit
s u	signed / unsigned	sign field present or absent
e f	extended / finite	domain includes $\pm\text{Inf}$, or spends those code points on finite datums

There are 504 legal (K,P,SGN,EXT) combinations, each with a draft-named alias. 120 aliases at $K \leq 8$ are exported by using `SmallFloats`.

Example alias	K	P	Signed	Domain
Binary8p4se	8	4	signed	extended
Binary8p4sf	8	4	signed	finite
Binary6p3ue	6	3	unsigned	extended
Binary5p5uf	5	5	unsigned	finite
Binary16p6se	16	6	signed	extended (opt-in export)
Binary16p1uf	16	1	unsigned	finite

<: versus ===

- Binary is abstract; Binary8p4se is its concrete representation.
- Binary8p4se <: Binary{8,4,true,true} is true.
- Binary8p4se === Binary{8,4,true,true} is false. Use the alias, or `format(K, P,  $\Sigma$ ,  $\Delta$ )` for runtime parameters. See Values, Code Points, and Conversion.

format() and opt-in access

Form	Scope	Result
using SmallFloats	exports 120 aliases, $K \leq 8$	alias directly in scope
using SmallFloats.Formats	brings all 504 aliases into scope	alias directly in scope
SmallFloats.Binary16p6se	unexported, always reachable	alias by qualified path
format(K, P, Σ , Δ)	programmatic, runtime parameters	the alias for those parameters
formatname(Binary{K,P, Σ , Δ })	programmatic	Symbol of the alias name

Format query table

Every query accepts either a Type or a value; `bitwidth(x) ≡ bitwidth(typeof(x))`, and both fold to a literal constant.

Query	Returns	Draft-named form	Related guide
<code>bitwidth(T)</code>	K	<code>BitwidthOf(T)</code>	Cheat Sheet
<code>precision(T)</code>	P	<code>PrecisionOf(T)</code>	Cheat Sheet
<code>issigned(T)</code>	Bool	<code>SignednessOf(T)</code>	Cheat Sheet
<code>isextended(T)</code>	Bool	<code>DomainOf(T)</code>	Cheat Sheet
<code>expbias(T)</code>	exponent bias	<code>ExponentBiasOf(T)</code>	Cheat Sheet
<code>expbitwidth(T)</code>	exponent field width	—	Cheat Sheet
<code>trailingsigbits(T)</code>	trailing significand bits	—	Cheat Sheet
<code>MaxFiniteOf(T)</code>	largest finite value, typed	—	Cheat Sheet
<code>MinFiniteOf(T)</code>	smallest-magnitude finite value, typed	—	Cheat Sheet
<code>MinPositiveOf(T)</code>	smallest positive value, typed	—	Cheat Sheet
<code>MinNormalOf(T)</code>	smallest positive normal value, typed	—	Cheat Sheet
<code>MaxSubnormalOf(T)</code>	largest subnormal value, typed	—	Cheat Sheet

Draft-named forms exist for every query in Group M, not only the four listed above (`BitwidthOf`, `PrecisionOf`, `SignednessOf`, `DomainOf`, `ExponentBiasOf`, and their neighbors); Julia-style and draft-named forms are interchangeable.

Value-and-type call forms

```

bitwidth(Binary8p4se)      # type form
x = Binary8p4se(1.6)
bitwidth(x)                 # value form, same answer, folds to 8
MaxFiniteOf(Binary8p4se)    # Binary8p4se(224.0 ≈ 0x7e)
MaxFiniteOf(x)              # same result

```

See also

[Core Model](#), [Values](#), [Code Points](#), and [Conversion](#), [Choose a Format](#).

Chapter 31

Projection Specifications

31.1 Projection Specifications

This page is the lookup contract for ProjSpec, the predefined deterministic grid, stochastic constructor families, and projection accessors. For the conceptual model, read [Rounding and Saturation](#).

ProjSpec constructor

```
ProjSpec(rounding_mode, saturation_mode)
```

A ProjSpec is the pair (rounding mode, saturation mode). Both components are zero-size singleton types; a ProjSpec costs nothing at runtime and fully specializes every call that reaches it.

Deterministic predefined specs (6 × 3)

Named `R<mode>_Sat<mode>`. Every cell below is an exported constant.

Rounding	SatFinite	SatPropagate	SatNone	Related guide
NearestTiesToEven	RNE_SF	RNE_SP	RNE_SN	Rounding and Saturation
NearestTiesToAway	RNA_SF	RNA_SP	RNA_SN	Rounding and Saturation
TowardPositive	RTP_SF	RTP_SP	RTP_SN	Rounding and Saturation
TowardNegative	RTN_SF	RTN_SP	RTN_SN	Rounding and Saturation
TowardZero	RTZ_SF	RTZ_SP	RTZ_SN	Rounding and Saturation
ToOdd	RT0_SF	RT0_SP	RT0_SN	Rounding and Saturation

`RNE_SN` is the package-wide initial default (`DefaultProjection()`).

Stochastic families

Constructors, not constants — the random-bit budget `N` is a type parameter. $N \in 1:60$; omitting `N` uses the default `N = 8` (`SmallFloats.DEFAULT_RBITS`).

Variant	SatFinite	SatPropagate	SatNone	Related guide
StochasticA	RSA_SF(N)	RSA_SP(N)	RSA_SN(N)	Rounding and Saturation
StochasticB	RSB_SF(N)	RSB_SP(N)	RSB_SN(N)	Rounding and Saturation
StochasticC	RSC_SF(N)	RSC_SP(N)	RSC_SN(N)	Rounding and Saturation

```
RSA_SN()      # StochasticA, N = 8, SatNone – the default form
RSA_SN(16)    # StochasticA, N = 16, SatNone
RSC_SF(16) == ProjSpec(StochasticC{16}(), SatFinite()) # true
```

Accessors

Accessor	Returns	Related guide
<code>roundingmode(p)</code>	the rounding-mode instance	Rounding and Saturation
<code>saturationmode(p)</code>	the saturation-mode instance	Rounding and Saturation
<code>isstochastic(p)</code>	Bool	Rounding and Saturation
<code>nrandbits(p)</code>	<code>N</code> for a stochastic spec	Rounding and Saturation

Saturation-mode meaning

Mode	Out-of-range behavior
<code>SatFiniteclamp</code>	everything to the finite range
<code>SatPropagate</code>	preserve representable infinities; clamp other overflow
<code>SatNone</code>	apply the draft's domain-, signedness-, and direction-dependent rows (nearest modes overflow to $\pm\text{Inf}$ or NaN for finite formats; directed modes pointing away from the overflow clamp to the extremal finite value)

Explicit R range rule

For an `N`-bit stochastic mode, an explicit draw `R` passed as a keyword must satisfy $0 \leq R \leq 2^N - 1$. Passing `R` outside that range is an error.

```
Add(T, σ, x, y; rng = Xoshiro(1)) # reproducible stream
Add(T, σ, x, y; R = 17)           # exact draw, ideal for tests
```

See also

[Rounding and Saturation](#), [Session Defaults](#).

Chapter 32

Operation Catalog

32.1 Operation Catalog

This is the canonical catalog of 52 operations: 30 unary, 18 binary, 3 ternary, and Convert. The tables distinguish P3109 operation names from Julia Base spellings. Use the named form when result format and projection must be explicit; use [Julia Compatibility Register](#) for the complete Base mapping.

Every operation here follows the same contract and it is not repeated per row: operands decode exactly, the operation evaluates exactly, and one projection produces the result code point. The tables therefore record names, arities, and the places where a draft operation's behavior differs from the Base spelling readers expect — not semantics, which live in [The Exact-Then-Project Contract](#).

Unary (30)

Operation	Arity	Base spelling	Notes	Related guide
Abs	unary	abs		Cheat Sheet
Negate	unary	unary -		Cheat Sheet
Sqrt	unary	sqrt		Cheat Sheet
RSqrt	unary	—	no Base counterpart; call by draft name	Julia Compatibility Register
Recip	unary	inv	Recip($\pm\text{Inf}$) = 0	Cheat Sheet
Exp	unary	exp		Cheat Sheet
Exp2	unary	exp2		Cheat Sheet
ExpMinusOne	unary	expm1		Cheat Sheet
Log	unary	log		Cheat Sheet
Log2	unary	log2		Cheat Sheet
LogOnePlus	unary	log1p		Cheat Sheet
Softplus	unary	—	no Base counterpart; call by draft name	Julia Compatibility Register
Sin	unary	sin		Cheat Sheet
Cos	unary	cos		Cheat Sheet
Tan	unary	tan		Cheat Sheet
ArcSin	unary	asin		Cheat Sheet
ArcCos	unary	acos		Cheat Sheet
ArcTan	unary	atan	single-argument form; see ArcTan2 for two-argument	Cheat Sheet
Sinh	unary	sinh		Cheat Sheet
Cosh	unary	cosh		Cheat Sheet
Tanh	unary	tanh		Cheat Sheet
ArcSinh	unary	asinh		Julia Compatibility Register
ArcCosh	unary	acosh		Julia Compatibility Register
ArcTanh	unary	atanh		Julia Compatibility Register
SinPi	unary	sinpi	reduces exactly mod 2 first	Cheat Sheet
CosPi	unary	cospi	reduces exactly mod 2 first	Cheat Sheet
TanPi	unary	tanpi	reduces exactly mod 2 first	Cheat Sheet
ArcSinPi	unary	—	no Base counterpart; call by draft name	Julia Compatibility Register
ArcCosPi	unary	—	no Base counterpart; call by draft name	Julia Compatibility Register
ArcTanPi	unary	—	no Base counterpart; call by draft name	Julia Compatibility Register

Trig of $\pm\text{Inf}$ is NaN for all rows above.

Binary (18)

Operation	Arity	Base spelling	Notes	Related guide
CopySign	binary	copysign		Julia Compatibility Register
Add	binary	+		Cheat Sheet
Subtract	binary	-		Cheat Sheet
Multiply	binary	*		Cheat Sheet
Divide	binary	/	$x / 0 \rightarrow \text{NaN}$ for every x , including $\pm\text{Inf}$	Cheat Sheet
Hypot	binary	hypot		Julia Compatibility Register
ArcTan2	binary	atan(y, x)	Base's (y, x) argument order	Cheat Sheet
ArcTan2Pi	binary	—	no Base counterpart; call by draft name	Julia Compatibility Register
Maximum	binary	max	NaN-propagating, exactly Base's float semantics	Cheat Sheet
Minimum	binary	min	NaN-propagating, exactly Base's float semantics	Cheat Sheet
MaximumNumber	binary	—	NaN-ignoring extremum; no Base counterpart	Julia Compatibility Register
MinimumNumber	binary	—	NaN-ignoring extremum; no Base counterpart	Julia Compatibility Register
MaximumMagnitude	binary	—	magnitude extremum; no Base counterpart	Julia Compatibility Register
MinimumMagnitude	binary	—	magnitude extremum; no Base counterpart	Julia Compatibility Register
MaximumMagnitudeNumber	binary	—	NaN-ignoring magnitude extremum; no Base counterpart	Julia Compatibility Register
MinimumMagnitudeNumber	binary	—	NaN-ignoring magnitude extremum; no Base counterpart	Julia Compatibility Register
MinimumFinite	binary	—	finite-preferring extremum; no Base counterpart	Julia Compatibility Register
MaximumFinite	binary	—	finite-preferring extremum; no Base counterpart	Julia Compatibility Register

$0 \cdot \infty \rightarrow \text{NaN}$. A NaN operand generally propagates, except the *Number/*Finite extremum variants, which prefer a non-NaN (and, for the Finite variants, finite) operand.

Ternary (3)

Operation	Arity	Base spelling	Notes	Related guide
FMA	ternary	fma, muladd	one rounding for both Base spellings	Cheat Sheet
FAA	ternary	—	fused add-add; no Base counterpart	Julia Compatibility Register
Clamp	ternary	clamp		Cheat Sheet

Conversion

Operation	Arity	Base spelling	Notes	Related guide
Convert	conversion	—	the one operation accepting non-Binary operands: Binary (any format), Float16/32/64, Float128, Integer, BigFloat	Values, Code Points, and Conversion

Catalog total: 30 unary + 18 binary + 3 ternary + 1 conversion = 52.

Base spellings that are componentwise composites

These Base functions are not single draft operations; each is a componentwise tuple of draft operations run under the session default.

Base spelling	Composite of
sincos	(Sin(x), Cos(x))
sincospi	(SinPi(x), CosPi(x))
minmax	(Minimum(x, y), Maximum(x, y))

Deliberately unmapped draft operations

No Base spelling exists for these; call them by draft name. Base has no counterpart for any row in this list.

Draft operation	Reason unmapped
RSqrt	Base has <code>inv(sqrt(x))</code> as two roundings, not one; no single-rounding Base spelling
Softplus	no Base equivalent
ArcSinPi	Base has only the <code>sinpi</code> family (no inverse)
ArcCosPi	Base has only the <code>sinpi</code> family (no inverse)
ArcTanPi	Base has only the <code>sinpi</code> family (no inverse)
ArcTan2Pi	Base has only the <code>sinpi</code> family (no inverse)
MaximumNumber, MinimumNumber	Base has no NaN-ignoring extremum pair
MaximumMagnitude, MinimumMagnitude	Base has no magnitude-extremum pair
MaximumMagnitudeNumber, MinimumMagnitudeNumber	Base has no NaN-ignoring magnitude-extremum pair
MinimumFinite, MaximumFinite	Base has no finite-preferring extremum pair
FAA	Base has no fused add-add

The mapping is a declarative partition over the full op list; the test suite asserts it is exhaustive — every operation is either mapped to a Base spelling or listed in `_NO_BASE_COUNTERPART`.

Call shapes

Explicit form, every operation:

```
Op(result_format, ρ, operands...; rng, R)
```

`rng` and `R` apply only to stochastic `ρ`.

Convenience form, same-format operands, `DefaultProjection()`:

```
Op(x...)
```

See also

[Cheat Sheet](#), [Why No Cross-Format Promotion](#), [Values](#), [Code Points](#), and [Conversion](#).

Chapter 33

Arrays, Blocks, and Packed Storage

33.1 Arrays, Blocks, and Packed Storage

This page collects signatures and contracts for bulk operations, block-scaled values, and packed storage. For procedures, use [Run Operations over Arrays](#), [Use Blocks for Dynamic Range](#), or [Pack Values for Storage](#).

Array operation forms

Form	Signature	Related guide
Unary array	<code>Op(fr, p, A)</code>	Cheat Sheet
Binary array	<code>Op(fr, p, A, B)</code>	Cheat Sheet
Ternary array	<code>Op(fr, p, A, B, C)</code>	Cheat Sheet
Generic map	<code>vmap(:Op, fr, p, operands...)</code>	Cheat Sheet
Generic map, in-place	<code>vmap!(C, Val(:Op), fr, p, operands...)</code>	Cheat Sheet
Callable operation object	<code>Op(:Op, fr, p) then .(A) or map(Op(:Op, fr, p), A, B)</code>	Run Operations over Arrays

Operands and destination must have matching axes.

Table cache introspection

Function	Signature	Notes	Related guide
<code>table_bytes</code>	<code>table_bytes()</code>	current cache footprint, bytes	Cheat Sheet
<code>empty_tables!</code>	<code>empty_tables!()</code>	reset the table cache	Cheat Sheet

Deterministic unary/binary array calls use cached result tables when the current signature satisfies the table policy (a `Binary8p4se` unary result table is 256 B; its binary table is 64 KiB). Ternary tables tier by combined operand bitwidth: eager (≤ 18 bits total), adaptive with LRU eviction (≤ 21 bits total), then compute beyond the adaptive band. Table entry width follows the result representation. Stochastic operations always compute per element and consume one draw per projection.

Block functions

Function	Signature	Related guide
<code>blocksize</code>	<code>blocksize(b)</code>	Cheat Sheet
<code>scaleformat</code>	<code>scaleformat(b)</code>	Cheat Sheet
<code>elemformat</code>	<code>elemformat(b)</code>	Cheat Sheet
<code>ConvertFromBlock</code>	<code>ConvertFromBlock(fr, p, b)</code>	Cheat Sheet
<code>ConvertToBlock</code>	<code>ConvertToBlock(fs, fr, p, values_tuple, scale)</code>	Cheat Sheet
<code>ConvertToBlockMaxAbsFinite</code>	<code>ConvertToBlockMaxAbsFinite(fs, fr, scale_p, element_p, values_tuple)</code>	Cheat Sheet
<code>BlockReduceAdd</code>	<code>BlockReduceAdd(fr, p, b)</code>	Cheat Sheet
<code>BlockReduceMultiply</code>	<code>BlockReduceMultiply(fr, p, b)</code>	Cheat Sheet
<code>BlockDotProduct</code>	<code>BlockDotProduct(fr, p, bx, by)</code>	Cheat Sheet

`BlockReduceAdd`, `BlockReduceMultiply`, and `BlockDotProduct` perform their lane products and accumulation exactly, projecting once at the end — no hidden intermediate rounding.

Generated BlockOp / ScaledOp pattern

Every scalar operation except `Convert` has a generated block form and a generated scaled form.

Pattern	Signature
<code>BlockOp</code>	<code>BlockOp(fr, p, b1, b2, result_scale)</code> (arity follows the underlying Op)
<code>ScaledOp</code>	<code>ScaledOp(fr, p, scale1, x1, scale2, x2, ...)</code> (block size 1; arity follows the underlying Op)

Examples:

```
BlockAdd(T, p, b1, b2, result_scale)
BlockExp(T, p, b, result_scale)

ScaledAdd(T, p, scale1, x1, scale2, x2)
ScaledExp(T, p, scale, x)
```

BlockVector

Stores many equal-shape blocks in a structure-of-arrays layout. No additional scalar API beyond the `Block` operations above; construct from a collection of same-shape `Blocks`.

PackedVector

Operation	Signature	Notes
Construct	<code>PackedVector(v)</code>	<code>v::AbstractVector{F}</code>
Index (read)	<code>pv[i]</code>	returns an <code>F</code>
Index (write)	<code>pv[i] = x</code>	<code>x::F</code>
Collect	<code>collect(pv)</code>	materializes a <code>Vector{F}</code>
sizeof rule	<code>sizeof(pv.data)</code>	<code>ceil(length(pv) * bitwidth(F) / 8)</code> bytes, rounded to the storage word
vmap acceptance	<code>vmap(:Op, F, ρ, pv)</code>	accepted directly; tiles unpack internally

`PackedVector` stores each code point in exactly `bitwidth(F)` bits. Computation unpacks tiles internally; packed arithmetic is deliberately not in-place — store packed, compute unpacked.

See also

[Cheat Sheet](#), [Run Operations over Arrays](#), [Use Blocks to Track Dynamic Range](#).

Chapter 34

Defaults, Randomness, and Display

34.1 Defaults, Randomness, and Display

This page collects the contracts for process defaults, random sources, and display policy. For their scope and concurrency consequences, read [Session State and Reproducibility](#).

Getter / setter table

Read with `DefaultX()`, set with `DefaultX!(v)`.

Default	Initial value	Setter accepts	Related guide
<code>DefaultType</code>	<code>Binary8p2se</code>	any fully-parameterized Binary type	Cheat Sheet
<code>DefaultReturnType</code>	<code>Binary8p2se</code>	any fully-parameterized Binary type	Cheat Sheet
<code>DefaultRoundingMode</code>	<code>NearestTiesToEven()</code>	mode instance or type	Cheat Sheet
<code>DefaultSaturationMode</code>	<code>SatNone()</code>	mode instance or type	Cheat Sheet
<code>DefaultProjection</code>	<code>RNE_SN</code>	a <code>ProjSpec</code> , or <code>(mode, sat)</code>	Cheat Sheet

```
DefaultType()      # Binary8p2se      DefaultType!(Binary8p4se)
DefaultReturnType() # Binary8p2se      DefaultReturnType!(Binary8p3se)
DefaultRoundingMode() # NearestTiesToEven()
DefaultSaturationMode() # SatNone()
DefaultProjection()  # RNE_SN
```

Coherence invariant

```
DefaultProjection() == ProjSpec(DefaultRoundingMode(), DefaultSaturationMode())
```

This holds at all times, in both directions:

Action	Effect
DefaultRoundingMode!(mode)	rebuilds DefaultProjection around the new mode, holding saturation fixed
DefaultSaturationMode!(mode)	rebuilds DefaultProjection around the new mode, holding rounding fixed
DefaultProjection!(p)	decomposes p back into both components

There is no DefaultAccumulatorType: removed after v0.1.0. Reduction accumulator carriers are picked by the span filter (provably exact for the operands), not by a session preference.

with_default_* combinators

Combinator	Signature	Related guide
with_default_type	with_default_type(f, args...) calls f(DefaultType(), args...)	Cheat Sheet
with_default_returntype	with_default_returntype(f, args...) calls f(DefaultReturnType(), args...)	Cheat Sheet
with_default_projection	with_default_projection(f, args...) calls f(DefaultProjection(), args...)	Cheat Sheet

```
with_default_type((T, x) -> T(x), 1.5)           # Binary8p2se(1.5 ≈ 0x41)
with_default_projection((p, x, y) -> Add(T, p, x, y), x, y)
with_default_returntype(f, args...)              # f(DefaultReturnType(), args...)
```

Each combinator restores nothing and changes nothing — it only *reads* the default at call time via `f(default, args...)`. There is no scoped mutate-and-restore form because the underlying state is process-global. Code that needs local numerical policy should pass `T` and `p` explicitly.

Cost model

- **Speculation guard.** Convenience forms `(x + y, Exp(x), T(2.1))` and the `with_default_*` combinators read the default through a speculation guard: allocation-free and concretely inferred while the default holds its initial value.
- **One barrier dispatch after a change.** Once a default is changed, each call crosses a function barrier: one dynamic dispatch at entry, everything inside still fully specialized.
- **Explicit forms are unaffected.** `Add(T, p, x, y)` with a `const` or argument-bound `p` never consults a session default and never pays this cost either way.
- **Boxing at escape.** When the combinator's result type *is* the default itself (`with_default_type` used as a constructor), the value computes on the specialized path but boxes once where it escapes the call — the irreducible cost of a runtime-chosen type. When the result type does not depend on the default, the call is zero-allocation with a concretely inferred result.

Randomness

Form	Entropy source	Projection control
<code>rand{T}</code>	task-local RNG	default uniform projection
<code>rand(rng, T)</code>	explicit rng	default uniform projection
<code>rand(rng, T; projection = p)</code>	explicit rng	explicit p
<code>randn(rng, T)</code>	explicit rng	default normal projection
stochastic operation with <code>rng = rng</code>	explicit stream	selected stochastic ProjSpec
stochastic operation with <code>R = r</code>	exact raw draw	selected stochastic ProjSpec

An omitted RNG uses Julia's task-local `Random.default_rng()`. The no-argument stochastic constructors use a compile-time bit budget of 8; select another budget explicitly, for example `RSA_SN(16)`.

An explicit `R` must lie in $0:(2^N - 1)$ for an N -bit stochastic mode. Array `rand`/`randn` forms use their documented defaults; construct arrays from scalar calls when each element must use another explicit projection.

Choose how values display

`get_show_style()` returns the process-wide style and `set_show_style!(style)` sets it. `VALID_SHOW_STYLES` contains four values:

Style	Example rendering
<code>:value</code>	1.625
<code>:codepoint</code>	0x45
<code>:datum</code>	(1.625 ≈ 0x45)
<code>:typed</code>	Binary8p4se(1.625 ≈ 0x45)

An `IIOContext` property, `:binary_show_style`, overrides the process default for one stream. Prefer that form in concurrent code and in libraries that should not change another caller's display.

See also

[First Session](#), [Julia Numeric Behavior](#), [Projection Specifications](#).

Chapter 35

Julia Compatibility

35.1 Julia Compatibility Register

Binary values speak ordinary Julia. The **Base register** (`src/juliacompat.jl`) maps every draft operation that has a Base counterpart onto its Base name, so generic Julia code — and your fingers — can use the spellings they already know. Every method there is exactly one same-format spec-register call under the session default projection; there is no third semantics: `x + y` is `Add(x, y)`, which is `Add(T, DefaultProjection(), x, y)`.

Using the overloads

```
julia> x, y, z = Binary8p4se(1.6), Binary8p4se(0.25), Binary8p4se(2.0);

julia> x + y                                     # Add
Binary8p4se(1.875 ≈ 0x47)

julia> exp(y)                                     # Exp — a 256-byte table lookup once warm
Binary8p4se(1.25 ≈ 0x42)

julia> atan(y, x)                                # ArcTan2, in Base's (y, x) argument order
Binary8p4se(0.15625 ≈ 0x2a)

julia> fma(x, y, z)                              # FMA — one rounding; muladd is the same op
Binary8p4se(2.5 ≈ 0x4a)

julia> sincos(y)                                 # composite: (Sin(y), Cos(y)) componentwise
(Binary8p4se(0.25 ≈ 0x30), Binary8p4se(1.0 ≈ 0x40))
```

Because the veneers follow the session default projection, changing it changes them (see [Session Defaults](#)):

```
julia> Binary8p4se(200.0) * Binary8p4se(2.0)      # RNE_SN: overflow → Inf
Binary8p4se(Inf ≈ 0x7f)

julia> DefaultProjection!(RTZ_SF);

julia> Binary8p4se(200.0) * Binary8p4se(2.0)      # now clamps to MaxFinite
Binary8p4se(224.0 ≈ 0x7e)

julia> DefaultProjection!(RNE_SN)                 # restore: this setter is PROCESS-wide
(NearestTiesToEven, SatNone)
```

Code that must be insensitive to the session default names its projection: `Multiply(T, RNE_SN, x, y)`.

Restore what you set

DefaultProjection! mutates state shared by every module in the process, and nothing scopes it for you — it changes `T(x::Real)` construction too, so a later `Binary8p4se(1.1)` would truncate rather than round to nearest. The `with_default_projection` combinator reads the configured projection; it does not scope mutation. Prefer explicit forms such as `Convert(T, ρ, x)` in library code. Every example after this point assumes the pristine (`NearestTiesToEven`, `SatNone`), which is why the demonstration above puts it back.

The mapping

Base spelling	draft operation
<code>+ - * /</code> , unary <code>-</code>	Add, Subtract, Multiply, Divide, Negate
<code>abs</code> , <code>inv</code> , <code>sqrt</code>	Abs, Recip, Sqrt
<code>exp</code> , <code>exp2</code> , <code>expm1</code> , <code>log</code> , <code>log2</code> , <code>log1p</code>	Exp, Exp2, ExpMinusOne, Log, Log2, LogOnePlus
<code>sin</code> , <code>cos</code> , <code>tan</code> , <code>asin</code> , <code>acos</code> , <code>atan</code>	Sin, Cos, Tan, ArcSin, ArcCos, ArcTan
<code>sinh</code> , <code>cosh</code> , <code>tanh</code> , <code>asinh</code> , <code>acosh</code> , <code>atanh</code>	Sinh, Cosh, Tanh, ArcSinh, ArcCosh, ArcTanh
<code>sinpi</code> , <code>cospi</code> , <code>tanpi</code>	SinPi, CosPi, TanPi
<code>atan(y, x)</code>	ArcTan2 (Base's argument order)
<code>copysign</code> , <code>hypot</code>	CopySign, Hypot
<code>min</code> , <code>max</code>	Minimum, Maximum (NaN-propagating, exactly Base's float semantics)
<code>fma</code> , <code>muladd</code>	FMA (both: one rounding)
<code>clamp</code>	Clamp
<code>sincos</code> , <code>sincospi</code> , <code>minmax</code>	componentwise composites of the draft ops

The mapping is a declarative partition over the op lists, and the test suite asserts it is exhaustive: every operation is either mapped above or listed in `_NO_BASE_COUNTERPART`. Adding an operation to the registry forces an explicit decision here.

What is *not* mapped, and why

Draft operations with no Base spelling — call them by their draft names: `RSqrt`, `Softplus`, `ArcSinPi`/`ArcCosPi`/`ArcTanPi`/`ArcTan2Pi` (Base has only the `sinpi` family), the NaN-ignoring/magnitude/finite extremum families (`MinimumNumber`, `MaximumMagnitude`, `MinimumFinite`, ...; Base has no NaN-ignoring pair), and `FAA` (no Base fused add-add).

```
julia> min(Binary8p4se(NaN), x)           # Base semantics: NaN propagates
Binary8p4se(NaN ≈ 0x80)

julia> MinimumNumber(Binary8p4se(NaN), x) # the NaN-ignoring draft op
Binary8p4se(1.625 ≈ 0x45)
```

Mixed Binary formats — deliberately a promotion ... failed to change error, never a silent widening. Mixing formats is an explicit `Convert`:

```
julia> x + Convert(Binary8p4se, RNE_SN, Binary5p3sf(1.0))
Binary8p4se(2.5 ≈ 0x4a)
```

Binary with ordinary numbers — promotes to the format's **promotion carrier** through the rules in `formats.jl`, so the result is an ordinary float, not a re-projected Binary.

For the 432 formats at rung 1 that carrier is `Float64`, which is what the example below shows. The 64 rung-2 formats promote to `Float128`; the 8 rung-3 formats promote to `BigFloat`. `Float64` cannot hold a `Binary16p1uf` datum, and promoting into it would return $\pm\text{Inf}$ from a perfectly finite value with no P3109 operation involved. `promotecarrier(T)` names the target.

```
julia> x + 2.0
3.625
```


Project it back explicitly when that is what you mean: `Convert{Binary8p4se, RNE_SN, decode(x) + 2.0}`.

The rest of the Base surface

Beyond arithmetic, Binary integrates where the other layers provide it: predicates and constants (`isnan`, `isfinite`, `signbit`, `zero`, `one`, `eps` in both its type and value forms, `typemin/typemax`, `floatmin/floatmax`), comparisons and `isless` on integer order keys, `sort` via an $O(n)$ counting sort (which places the single NaN **first**, below $-\text{Inf}$, per the draft's total order — not last as `Float64` sorting does), `nextfloat/prevfloat` as the draft Next operations, `codepoint`, and `rand/randn` (see [Random Values](#)). All of it goes through the same projection engine and decode tables as the draft-named API.

The AbstractFloat contract

Binary `<: AbstractFloat`, and generic Julia code reaches for that interface without asking. The routinely used numeric contracts are implemented, while operations with no P3109 meaning refuse explicitly:

verb	behaviour
<code>hash</code> , <code>Base.decompose</code>	exact — a datum is $S \cdot 2^Q$, so <code>Dict</code> keys and <code>Set</code> elements work
<code>round</code> , <code>floor</code> , <code>ceil</code> , <code>trunc</code>	one rounding: the integer value is formed exactly on the carrier and projected once, under the session default's saturation
<code>round(x, r::RoundingMode)</code>	<code>RoundNearest</code> , <code>RoundNearestTiesAway</code> , <code>RoundUp</code> , <code>RoundDown</code> , <code>RoundToZero</code>
<code>exponent</code> , <code>significand</code> , <code>frexp</code>	exact; <code>DomainError</code> on zero and the non-finites, as in <code>Base</code>
<code>eps(x)</code>	the ulp at <code>x</code> , and always exactly a datum of <code>x</code> 's own format — see below
<code>ldexp(x, n)</code>	exact: scaling by a power of two moves the exponent field and nothing else
<code>widen</code>	the format's promotion carrier — see No Automatic Promotion for why it is not another format
<code>reinterpret</code>	both directions, and the incoming one is checked against the representation invariant

`eps` has both forms: `eps(T)` is the constant $2^{(1-P)}$, and `eps(x)` is the spacing of the lattice at `x`. The second is always representable — the ulp is 2^Q with $2-P-B \leq Q \leq e_{\text{max}}-P+1$, and every power of two in a format's range is a datum for every $P \geq 1$ — so it is exact rather than a projected approximation, and it does not vary with the session projection. In the subnormal band the spacing is constant, and at zero it is the least positive datum.

```
julia> hash(Binary8p4se(1.5)) == hash(Binary8p4se(1.5))
true

julia> Base.decompose(Binary8p4se(1.5))           # 12 · 2-3 = 1.5, exactly
(12, -3, 1)

julia> floor(Binary8p4se(1.6)), ceil(Binary8p4se(1.1))
(Binary8p4se(1.0 ≈ 0x40), Binary8p4se(2.0 ≈ 0x48))

julia> eps(Binary8p4se(1.5)), eps(Binary8p4se(64.0)) # the lattice widens
(Binary8p4se(0.125 ≈ 0x28), Binary8p4se(8.0 ≈ 0x58))

julia> ldexp(Binary8p4se(1.5), 2), ldexp(Binary8p4se(1.5), -2)
(Binary8p4se(6.0 ≈ 0x54), Binary8p4se(0.375 ≈ 0x34))

julia> reinterpret(Binary8p4se, 0x44)             # checked, unlike `rawvalue`
Binary8p4se(1.5 ≈ 0x44)
```

Refusals name themselves. Where a `Base` verb has no draft counterpart, the error says which and why rather than surfacing as a bare `MethodError` that cannot be told apart from an oversight:

```
julia> rem(Binary8p4se(1.5), Binary8p4se(2.0))
ERROR: ArgumentError: rem is not defined for Binary8p4se: the draft defines no remainder; ...
```

That covers rem/mod, round under RoundFromZero and RoundNearestTiesUp (no draft μ), and significand on the one format whose binade $[1, 2)$ is truncated by its Inf encoding.

Chapter 36

Conformance and Approximation

36.1 Conformance and Approximation Registry

This page is the contract for exporting the live conformance declaration and for measuring and registering optional approximate implementations. The default operation API never substitutes a registered approximation.

Conformance functions

Function	Signature	Returns	Related guide
conformance	conformance()	the live declaration: formats, operations, mode vocabulary, table specializations instantiated this session, registered approximations	Read and Export Conformance
conformance_dict	conformance_dict()	a plain nested Dict, for JSON/TOML serialization	Read and Export Conformance
conformance_report	conformance_report()	prints the declaration	Read and Export Conformance
draft_revision	draft_revision()	the P3109 draft revision this package conforms to	Read and Export Conformance
draft_identity	draft_identity()	structured designation, upload date, retained source, and SHA-256 digest	Read and Export Conformance

κ registry

Function	Signature	Returns	Related guide
measure_ kappa	measure_kappa(fn, op, fr, (argformats...), ρ)	(κ, exhaustive) — measured max code-point deviation and whether the measurement was exhaustive	Register an Approximation
register_ approx!	register_approx!(name, op, fr, args, ρ, fn; κ)	registers fn as approximation name for op, declaring bound κ	Register an Approximation
approx	approx(name)	the registered implementation object	Register an Approximation
kappa	kappa(impl)	the declared κ bound	Register an Approximation
kappa_ measured	kappa_measured(impl)	the measured κ from registration	Register an Approximation
list_approx	list_approx()	all registered approximation names	Register an Approximation
unregister_ approx!	unregister_approx!(name)	removes a registered approximation	Register an Approximation

```
κ, exhaustive = measure_kappa(fn, :Exp, T, (T,), ρ)
register_approx!(:fast_exp, :Exp, T, (T,), ρ, fn; κ)
impl = approx(:fast_exp)
kappa(impl)
kappa_measured(impl)
list_approx()
unregister_approx!(:fast_exp)
```

Understatement-rejection rule

κ is measured exhaustively at registration time whenever the domain permits exhaustive enumeration. Declaring a κ smaller than the measured deviation throws — an implementation cannot register itself as more accurate than it measures.

κ = NaN acknowledgment rule

An implementation whose NaN behavior differs from the default API's (a NaN mismatch against the exhaustive reference) must be registered with an explicit κ = NaN. This is an acknowledgment, not a numeric bound: it cannot be produced by measurement alone and must be supplied by the caller of register_approx!.

Approximate implementations registered through this path are never substituted into the default API; the default API is bit-exact, always.

See also

[Read and Export Conformance](#), [Register an Approximation](#), and [Verification Strategy](#).

Chapter 37

Public API Index

37.1 Public API Index

This index renders source docstrings for the public surface. Begin with the curated family table: it states where each contract is maintained. The 120 format aliases at $K \leq 8$ are exported by using `SmallFloats`; the other 384 are opt-in through `SmallFloats.Format`s or available by qualification.

The index below groups the exported surface by family and points each family at the descriptive reference page that documents its full table. The `@autodocs` block that follows renders the docstrings themselves, alphabetically, as generated from source.

Index

Family	Scope	Descriptive reference page
Arrays, blocks, packed storage	<code>Op(fr,p,A,...)</code> , <code>vmap/vmap!</code> , <code>Block</code> , <code>BlockVector</code> , <code>PackedVector</code> , <code>table_bytes</code> , <code>empty_tables!</code>	Arrays, Blocks, and Packed Storage
Conformance and approximation	<code>conformance</code> , <code>conformance_dict</code> , <code>conformance_report</code> , <code>draft_revision</code> , <code>draft_identity</code> , <code>measure_kappa</code> , <code>register_approx!</code> , <code>approx</code> , <code>kappa</code> , <code>kappa_measured</code> , <code>list_approx</code> , <code>unregister_approx!</code>	Conformance and Approximation
Registry		Registry
Formats	504 aliases (120 exported at $K \leq 8$), <code>format</code> , <code>formatname</code> , and <code>format-query</code> accessors	Formats and Value Queries
Operations	30 unary, 18 binary, 3 ternary, and <code>Convert</code>	Operation Catalog
Projection specs	<code>ProjSpec</code> , deterministic grid, stochastic constructors, and accessors	Projection Specifications
Session state	numerical defaults, Julia/explicit RNG use, display styles, and default-consumption combinators	Defaults, Randomness, and Display
Julia integration	Base overloads, promotion carriers, sorting, and conversion behavior	Julia Compatibility Register

SmallFloats.SmallFloats – Module.

SmallFloats

A conforming, performance-oriented Julia implementation of the IEEE P3109 draft standard for arithmetic formats for machine learning (bitwidths 3–16).

Bit-exact defined results on every default path; the projection engine (`RoundToPrecision` → `Saturate` → `Encode`) is the single write path into a code point; approximate fast paths exist only behind the explicit `κ` registry (`register_approx!` / `approx`), never substituted silently. See `conformance()` for the live declaration and `SmallFloats.draft_revision()` for the implemented draft.

Performance note

Pass format types through `const` bindings, type parameters, or function arguments. All exported entry points fully specialize when the format type is statically known; the checked-in benchmark report records current timings and allocations. Calling through a **non-const global** format type forces dynamic dispatch on every call; this is Julia semantics, not package overhead. A single function barrier (`f(::Type{T}, args...)` where `{T}`) restores full speed. The test suite pins the specialization properties (concrete return types, zero warm-path allocation) as deterministic regressions.

`SmallFloats.RNA_SF` – Constant.

(`NearestTiesToAway`, `SatFinite`).

`SmallFloats.RNA_SN` – Constant.

(`NearestTiesToAway`, `SatNone`).

`SmallFloats.RNA_SP` – Constant.

(`NearestTiesToAway`, `SatPropagate`).

`SmallFloats.RNE_SF` – Constant.

(`NearestTiesToEven`, `SatFinite`).

`SmallFloats.RNE_SN` – Constant.

(`NearestTiesToEven`, `SatNone`) — the package-wide default `p`.

`SmallFloats.RNE_SP` - Constant.

(NearestTiesToEven, SatPropagate).

`SmallFloats.RTN_SF` - Constant.

(TowardNegative, SatFinite).

`SmallFloats.RTN_SN` - Constant.

(TowardNegative, SatNone).

`SmallFloats.RTN_SP` - Constant.

(TowardNegative, SatPropagate).

`SmallFloats.RTO_SF` - Constant.

(ToOdd, SatFinite).

`SmallFloats.RTO_SN` - Constant.

(ToOdd, SatNone).

`SmallFloats.RTO_SP` - Constant.

(ToOdd, SatPropagate).

`SmallFloats.RTP_SF` - Constant.

(TowardPositive, SatFinite).

`SmallFloats.RTP_SN` - Constant.

(TowardPositive, SatNone).

`SmallFloats.RTP_SP` - Constant.

(TowardPositive, SatPropagate).

`SmallFloats.RTZ_SF` - Constant.

(TowardZero, SatFinite).

`SmallFloats.RTZ_SN` - Constant.

(TowardZero, SatNone).

`SmallFloats.RTZ_SP` - Constant.

(TowardZero, SatPropagate).

`SmallFloats.ApproxImpl` - Type.

A registered k-approximate implementation of one operation specialization.

SmallFloats.Binary - Type.

```
Binary{K,P,SGN,EXT} <: AbstractFloat
```

A P3109 **format**: bitwidth $K \in 3:16$, precision P , signedness $SGN::\text{Bool}$ ($\text{true} = \text{Signed}$), domain $EXT::\text{Bool}$ ($\text{true} = \text{Extended}$, i.e. the datum set includes infinities).

Binary is **abstract**, and that is the draft's own distinction: a format is "a datum set and an encoding" (D1 §3.1), while the width of the machine word carrying a code point is a representation choice below the standard's level of description. Code8 is that representation. The parameterized name is the *format*; the draft-named alias is its *representation*:

```
Binary{8,4,true,true} # the format - abstract
Binary8p4se           # Code8{8,4,true,true}, its representation - concrete
```

Consequently `Binary8p4se == Binary{8,4,true,true}` is **false** and `Binary8p4se <: Binary{8,4,true,true}` is **true**. Use the alias, or `format(K, P,  $\Sigma$ ,  $\Delta$ )`, wherever a concrete type is needed — an array element type, a similar, a constructor target.

The code point occupies the low K bits of the storage unit; the high bits are maintained zero as a representation invariant. Construct code points with `T(c::Unsigned)` (validated code-point construction); the implementation uses an unexported unchecked primitive only after establishing the invariant. Construct from numeric values with `T(x::Real)` (default projection spec) or `Convert`. `Unsigned` is the one argument-type class meaning *code point*, at **every** width — `T(0x02)` is code point 2 whether `T`'s unit is a `UInt8` or a `UInt16`; all other `Reals` mean *value*.

There are $4K - 2$ formats at each K ($P \in 1:K \times \text{signed/unsigned} \times \text{finite/extended}$, less the two $P == K$ signed combinations, which have no exponent bits), so $K_{\text{MIN}}:K_{\text{MAX}} = 504$ in total.

SmallFloats.Block - Type.

```
Block{B, FS<:Binary, FE<:Binary}
```

A P3109 block (`s, [x1 ... xB]`): scale factor in format `FS`, $B \geq 1$ elements in format `FE` (draft §5). `isbits`; blocks live in registers/stack.

SmallFloats.BlockVector - Type.

```
BlockVector{B,FS,FE} <: AbstractVector{Block{B,FS,FE}}
```

Structure-of-arrays storage: `scales::Vector{FS}`, `elems::Matrix{FE}` ($B \times n$, column-major so each block's elements are one contiguous column).

SmallFloats.FPClass - Type.

Eight-way datum classification returned by `Class` (draft §4.13.1).

SmallFloats.NearestTiesToAway - Type.

Round to nearest; ties away from zero (draft §4.7.1).

SmallFloats.NearestTiesToEven - Type.

Round to nearest; ties to the even code point (IEEE default; draft §4.7.1).

`SmallFloats.Op` – Type.

```
Op(name::Symbol, FR::Type{<:Binary}, p::ProjSpec) -> Op
```

A registry operation bound to a result format and a projection, as a **callable** — so the idiomatic array verbs work directly:

```
map(Op(:Add, Binary8p4se, RNE_SN), A, B)
Op(:Exp, Binary8p4se, RNE_SN).(A)
```

`Op` is a singleton type: the operation, format and projection all live in the type parameters, so calls specialize exactly as `Add(Binary8p4se, RNE_SN, x, y)` does. Use `vmap!` when you have a destination array to fill — it is the in-place, table-aware kernel and this type is a thin front for it.

`SmallFloats.PackedVector` – Type.

```
PackedVector{F<:Binary} <: AbstractVector{F}
```

Bit-packed storage of `F` code points at `bitwidth(F)` bits per element. Construct with `PackedVector(A::AbstractVector{F})`; recover with `Vector(pv)` or `collect(pv)`. Memory: $\lceil n \cdot K / 64 \rceil$ words vs $n \cdot \text{sizeof}(\text{codeunit_type}(F))$ unpacked.

It saves nothing when $K == 8 \cdot \text{sizeof}(\text{codeunit_type}(F))$ — $K = 8$ and $K = 16$ — where the packed layout is bit-for-bit the unpacked one and every access pays the shift/mask/splice for no return. `packing_saves(F)` says so; see its docstring for why that is a query rather than an error.

`SmallFloats.ProjSpec` – Type.

```
ProjSpec{R<:RoundingMode3109, S<:SaturationMode}
```

A projection specification `p = (rounding mode, saturation mode)`, draft §4.2. Zero-size; construct as `ProjSpec(NearestTiesToEven(), SatNone())` or via the exported constants. Kernels specialize on its type.

`SmallFloats.StochasticA` – Type.

Stochastic rounding, variant A of draft §4.7.4: $\text{RoundAway} \iff \lfloor v \cdot 2^N \rfloor + R \geq 2^N$.

`SmallFloats.StochasticB` – Type.

Stochastic rounding, variant B: $\text{RoundAway} \iff \lfloor v \cdot 2^{(N+1)} \rfloor + (2R+1) \geq 2^{(N+1)}$.

`SmallFloats.StochasticC` – Type.

Stochastic rounding, variant C: $\text{RoundAway} \iff \text{RNITE}(v \cdot 2^N) + R \geq 2^N$.

`SmallFloats.ToOdd` – Type.

Round inexact results to the nearest *odd* code point (draft §4.7.3; sticky-friendly).

`SmallFloats.TowardNegative` – Type.

Directed rounding toward $-\infty$ (draft §4.7.2).

`SmallFloats.TowardPositive` – Type.

Directed rounding toward $+\infty$ (draft §4.7.2).

`SmallFloats.TowardZero` – Type.

Directed rounding toward zero — truncation (draft §4.7.2).

`SmallFloats.BlockDotProduct` – Method.

`BlockDotProduct` (draft §5.3.2). Lane products can carry 64 significant bits, so the products and their sum are formed in the exact accumulator, with the ∞/NaN fold algebra resolved on the `Float64` classifications first.

`SmallFloats.BlockReduceAdd` – Method.

`BlockReduceAdd` (draft §5.3.1): `project(reduce( $\omega$ Add, [0, X...]))`.

`SmallFloats.BlockReduceMultiply` – Method.

`BlockReduceMultiply` (draft §5.3.1): `project(reduce( $\omega$ Multiply, [1, X...]))`.

`SmallFloats.Class` – Method.

`Class(v) -> FPClass` (draft §4.13.1): classify `v` as `NaN`, $\pm\text{Inf}$, $\pm\text{normal}$, $\pm\text{subnormal}$, or zero.

`SmallFloats.Convert` – Method.

```
Convert(fr, p, A::AbstractArray{<:Union{Float16,Float32,Float64,BFloat16}}; rng) -> Array{fr}
```

Bulk ingestion: exact widening per element, then projection under `p`. Unlike the Binary-array `Convert` (a Shape-A table gather), external inputs are not enumerable, so this is a Shape-B loop; stochastic `p` draws per element.

`SmallFloats.Convert` – Method.

```
Convert(fr, p, x) -> fr
```

Draft §4.9 `Convert(fx, fr, p)`. Accepts Binary, IEEE binary16/32/64 floats (widened exactly to the `Float64` carrier), `Float128` (projected directly), Dyadic (the rung-3 datum carrier — what `decode` returns for the eight rung-3 formats; projected directly and exactly), `Integer` (exact via a sufficiently wide `BigFloat`), and `BigFloat` (projected directly; the caller warrants the value is exact).

`SmallFloats.ConvertFromBlock` – Method.

`ConvertFromBlock` (§5.2.1): `blockdecode`, then project each element (no scale division).

`SmallFloats.ConvertToBlock` – Method.

`ConvertToBlock` (§5.2.2): `decode` elements, `BlockProject` against the supplied scale.

`SmallFloats.ConvertToBlockMaxAbsFinite` – Method.

`ConvertToBlockMaxAbsFinite` (§5.2.3): `S = reduce( $\omega$ MaximumFinite, [NaN, |X1|...])`, project `S` under `ps`, then `BlockProject` the elements under `p`.

`SmallFloats.DefaultProjection` – Method.

```
DefaultProjection() -> ProjSpec
DefaultProjection!(p::ProjSpec) -> p
DefaultProjection!(m, s) -> ProjSpec
```

The session's default projection specification. Initialized to `(DefaultRoundingMode(), DefaultSaturationMode()) = RNE_SN`. Setting it directly decomposes `p` into `DefaultRoundingMode` and `DefaultSaturationMode`, so the three stay coherent in both directions.

SmallFloats.DefaultReturnType – Method.

```
DefaultReturnType() -> Type{<:Binary}
DefaultReturnType!(T::Type{<:Binary}) -> T
```

The session's default *result* format — the format an operation projects into when the caller does not name one. Initialized to Binary8p2se. Independent of DefaultType, which is the default *operand* format.

SmallFloats.DefaultRoundingMode – Method.

```
DefaultRoundingMode() -> RoundingMode3109
DefaultRoundingMode!(m) -> m
```

The session's default rounding mode. Initialized to NearestTiesToEven(). Setting it rebuilds [DefaultProjection](#) from the new mode and the current [DefaultSaturationMode](#). Accepts an instance or a fully-parameterized type (NearestTiesToAway, StochasticA{8}, ...).

SmallFloats.DefaultSaturationMode – Method.

```
DefaultSaturationMode() -> SaturationMode
DefaultSaturationMode!(s) -> s
```

The session's default saturation mode. Initialized to SatNone(). Setting it rebuilds [DefaultProjection](#) from the current [DefaultRoundingMode](#) and the new mode. Accepts an instance or a type.

SmallFloats.NextGreaterThan – Method.

[NextGreaterThan\(v\)](#) (draft §4.16): the least datum greater than *v* in the total order — one step up the code lattice, with NaN → NaN and MaxFinite/+Inf → NaN at the top. Base.nextfloat on Binary is this operation.

SmallFloats.NextLessThan – Method.

[NextLessThan\(v\)](#) (draft §4.16): the greatest datum less than *v* in the total order — one step down the code lattice, with NaN → NaN and MinFinite/-Inf → NaN at the bottom. Base.prevfloat on Binary is this operation.

SmallFloats.RoundOf – Function.

Draft-named accessor: [RoundOf\(p\)](#) — the rounding mode of *p* (§4.2).

SmallFloats.SatOf – Function.

Draft-named accessor: [SatOf\(p\)](#) — the saturation mode of *p* (§4.2).

SmallFloats.TotalOrder – Method.

[TotalOrder\(fx,fy\)](#) (draft §4.12.1): $x \leq y$ in the total order (single NaN smallest — $\text{NaN} \leq -\text{Inf} \leq \dots \leq +\text{Inf}$). Same-format comparisons run on the integer order key (bitops plan K1) — proven equivalent to the decode order exhaustively over all pairs; cross-format keys are not comparable, so mixed formats keep the decode path.

SmallFloats.approx – Method.

Retrieve a registered approximate implementation (callable via `.fn`).

SmallFloats.bitwidth – Method.

Format bitwidth *K* (3–8).

`SmallFloats.codedistance` – Method.

Code-point distance along the total order between two values of one format.

`SmallFloats.codeunit_type` – Method.

The storage unit of a format's code point: `UInt8` for $K \leq 8$, `UInt16` above.

`SmallFloats.conformance` – Method.

```
conformance() -> ConformanceDeclaration
```

The package's draft-§4.6-style conformance declaration, derived live from the operation registry, the table cache (the specializations actually instantiated), and the approximate-implementation registry. Serialize with `conformance_dict` or render with `conformance_report`.

`SmallFloats.conformance_dict` – Function.

Nested `Dict{String,Any}` form of the declaration — serializable by any JSON/TOML writer.

`SmallFloats.conformance_report` – Function.

Human-readable conformance report.

`SmallFloats.decode!` – Method.

```
decode!(dest::AbstractArray{Float32}, A::AbstractArray{<:Binary}) -> dest
decode!(dest::AbstractArray{Float64}, A::AbstractArray{<:Binary}) -> dest
```

Exact bulk `ωDecode` into an external float array.

Exactness is the whole contract here — it is what distinguishes `decode!` from `Float32.(A)`, which rounds like any other Julia conversion — so the destination element type is **gated on the format**, not assumed. `datumsexact(X, F)` decides it, and a format whose datums do not all fit `X` raises rather than silently rounding a bulk array. Every $K \leq 8$ format passes both gates, so this is a new refusal only for formats that did not exist before the $K \leq 16$ extension.

Where the gate passes and $K \leq \text{KSPLIT}$, the implementation is unchanged: a Shape-A gather of the (narrowed) constant decode table.

`SmallFloats.decode` – Method.

```
decode(v::Binary) -> datumcarrier(typeof(v))
```

`ωDecode` (draft §4.7.2). **Exact for every datum of every format**, which is why the return type is the format's carrier rather than a fixed one: `Float64` where the bias allows (432 of the 504 formats), `Float128` to `B = 8192`, `BigFloat` above. `Float64(v)` is the way to ask for a `Float64` and accept the rounding.

Two implementations, selected by `decodepolicy` — dispatch on the representation, not a branch on `K` (invariant 9, §2 R-F):

- `TableDecode` ($K \leq \text{KSPLIT}$): a 2^K -entry constant-tuple lookup (bitops plan $K2$), generated once from `_decode_compute`, so the two are correct by construction and asserted equivalent exhaustively; constant inputs fold.
- `ComputeDecode` ($K > \text{KSPLIT}$): `_decode_compute` directly. No table — 65 536 entries is what invariant 10 exists to forbid.

`SmallFloats.draft_identity` - Method.

Structured identity and SHA-256 digest of the retained P3109 draft transliteration.

`SmallFloats.draft_revision` - Method.

Human-readable designation of the retained P3109 draft implemented by this release.

`SmallFloats.empty_tables!` - Method.

Drop every cached table and adaptive counter (they rebuild lazily on next use).

`SmallFloats.expbias` - Method.

Exponent bias: $2^{(K-P-1)}$ signed, $2^{(K-P)}$ unsigned.

`SmallFloats.expsbitwidth` - Method.

Width of the exponent field in bits: $(K - \text{signbit}) - (P - 1)$.

`SmallFloats.f32_exact` - Method.

```
f32_exact(op, f1, f2) -> Bool
```

True iff `op` $\in$ (`:Add`, `:Subtract`, `:Multiply`) on Float32-decoded operands is exact for every finite pair of (`f1`, `f2`) datums — checked once by exhaustive enumeration against a 300-bit BigFloat oracle, then cached. When true, a Float32 intermediate is an exact carrier under every projection mode.

Defined at **every** bitwidth. A signature whose pair count exceeds `F32_EXACT_MAX_PAIRS[]` raises rather than enumerating; that is a cost refusal naming the budget, never a statement about the formats.

`SmallFloats.f32_impl` - Method.

```
f32_impl(op, fr, p) -> fn
```

The `decode32` $\rightarrow$ guarded-`op32` $\rightarrow$ project kernel for `op(...  $\rightarrow$  fr, p)`, shaped for `register_approx!`. Nearest-ties `p` only: directed modes after a nearest-rounded Float32 compute measure $\kappa \geq 1$ (and $\kappa = \text{NaN}$ at saturation boundaries), and stochastic `p` is rejected by κ measurement itself.

`SmallFloats.format` - Method.

```
format(K, P, Σ, Δ) -> Type
```

The concrete representation type for a format, from **runtime** parameters — `format(8, 4, true, true) == Binary8p4se`. The programmatic route to a type that can be an array element or a constructor target, and the supported replacement for spelling `Binary{K,P,Σ,Δ}` where a concrete type is meant.

Deliberately a Dict lookup: this path is reflection, it is not performance-sensitive, and it must not pretend to be. Where the parameters are static, the aliases and `reptype` are the zero-cost routes.

`SmallFloats.formatname` - Method.

`formatname(T)` — the draft §3.2 name of a format type.

`SmallFloats.ftz_variant` - Method.

```
ftz_variant(op, fr, fl, p) -> fn
```

Build the Annex-style approximate unary implementation: compute the defined result, then flush subnormal results to zero or `MinNormalOf(fr)`, whichever is nearer, ties toward zero (magnitude-symmetric for signed formats). Returned `fn` is suitable for `register_approx!`; for `fr` with precision P its κ is $2^{(P-2)}$ when $P \geq 2$ (largest flushed-to-zero subnormal) and 0 when $P = 1$ (no subnormals exist).

`SmallFloats.get_show_style` - Method.

```
get_show_style()      -> Symbol
get_show_style(io::IO) -> Symbol
```

The style show will use: the process default, or — given an `IO` — the style that `io` resolves to, which is its `:binary_show_style` property when set and the process default otherwise.

The no-argument form previously read an undefined `io` and threw `UndefVarError` on every call.

`SmallFloats.isextended` - Method.

Whether the format's domain is Extended (datum set includes infinities).

`SmallFloats.issigned` - Method.

Whether the format is Signed (has a sign bit and negative datums).

`SmallFloats.list_approx` - Method.

Sorted names of all registered κ -approximate implementations.

`SmallFloats.measure_kappa` - Method.

```
measure_kappa(fn, op, fr, argformats, p; max_exhaustive=2^22, rng, samples=2^20)
-> (κ::Float64, exhaustive::Bool)
```

Measure κ of `fn(x::argformats[1], ...):fr` against the defined results of draft operation `op` under `p`. Enumerates the full input cross-product when it has at most `max_exhaustive` points; otherwise measures on `samples` uniform draws and reports `exhaustive = false`.

The parenthetical here used to read "always true for arity ≤ 2 ", which was a $K \leq 8$ fact: two $K = 16$ formats are 2^{32} points against a 2^{22} budget, so a wide binary signature samples. That is the mechanism working — κ is a *measurement*, `exhaustive = false` is reported, and `conformance_report` prints "(κ sampled — not exhaustive)" — but the claim was no longer true.

`SmallFloats.nrandbits` - Method.

Number of random bits N consumed per projected element; 0 for pure modes.

`SmallFloats.projmode` - Method.

Translate a `Base.RoundingMode` to its `RoundingMode3109` counterpart (identity on the latter). Used only at API boundaries. Internals always carry `RoundingMode3109`.

`SmallFloats.register_approx!` - Method.

```
register_approx!(name, op, fr, argformats, p, fn; κ=nothing, kwargs...) -> ApproxImpl
```

Register `fn` as a named κ -approximate implementation of `op(argformats → fr, p)`. κ is measured by enumeration (see `measure_kappa`); a declared κ smaller than the measured value is **rejected**. Omitting κ declares the measured value. NaN- κ implementations (non-finite mismatches) may be registered only by declaring κ =NaN.

`SmallFloats.register_f32!` - Method.

```
register_f32!(op, fr, argformats, p; κ=nothing) -> ApproxImpl
```

Build, measure, and register the Float32 kernel for `op(argformats → fr, p)` under the canonical name `f32/op/args=fr/rm_sm`. Registration measures κ by exhaustive enumeration and fails loudly on understatement — that is the point. Measured baseline (2026-07-24): $\kappa = 0$ for all 120 same-format signatures $\times \{\text{Add, Subtract, Multiply, Divide, Sqrt}\} \times \text{RNE}_{\{\text{SatNone, SatFinite, SatPropagate}\}}$.

`SmallFloats.reptype` - Method.

```
reptype(F) -> Type
```

The concrete representation of a format. The identity on a concrete input, and the `Binary{K,P,S,E} → Code8|Code16` map on an abstract one.

Julia will **not** exploit "this abstract type has exactly one subtype": there are no sealed types, inference cannot narrow an abstract element type to its unique subtype, and nothing stops a third party declaring another `<: Binary`. So `reptype` carries real weight rather than being belt-and-braces — it is how a method that holds only the format parameters reaches a constructible type.

The `Val` barrier is deliberate (invariant 9). The obvious spelling, `K <= KSPLIT ? Code8{...} : Code16{...}`, returns a `Type`, so its inferred return type is a small Union unless the compiler folds the comparison. It usually will; "usually" cannot underwrite a correctness-critical selection that also gates a zero-allocation guarantee. Behind a `Val`, each `_rep` method has exactly one concrete return type and there is nothing to fold. Gate G9 checks it per format.

`SmallFloats.table_bytes` - Method.

Total bytes currently held by every table cache.

`SmallFloats.table_count` - Method.

Number of cached unary/binary specializations, both code units.

`SmallFloats.table_policy` - Method.

```
table_policy(op, fr, f1[, f2[, f3]], p) -> (; shape, entries, bytes, reason)
```

Which shape an array call on this signature will take, and why — `:A` for the table gather, `:B` for the per-element scalar path.

§4 Stage 5 item 6: "this signature is running the compute kernel" should be *discoverable*, not inferred from a benchmark that came out slower than expected. Reads the same predicates the kernels do, so it cannot drift from them.

`SmallFloats.ternary_count` - Method.

Number of cached ternary specializations, both code units.

`SmallFloats.trailing_sigbits` – Method.

Trailing-significand width $P - 1$ (the stored fraction bits).

`SmallFloats.unregister_approx!` – Method.

Remove a registered k-approximate implementation (no-op if absent).

`SmallFloats.vmap!` – Function.

```
vmap!(dest, Val(op), fr, p, A[, B[, C]]; rng) -> dest
vmap(op, fr, p, A...; rng)
```

Elementwise draft operation over arrays of Binary values, projecting into `fr` under `p`. Pure `p` runs the Shape-A gather whenever a table exists or the ternary bitwidth policy grants one; stochastic specializations (and untabled ternary signatures) run the scalar path per element (each element drawing its own `R`).

`SmallFloats.vmap` – Method.

```
vmap(op, fr, p, pv::PackedVector; rng=nothing) -> Vector{fr}
```

Elementwise `op` over packed storage: unpack a 256-element tile into a byte scratch, run the ordinary byte kernel (`vmap!`), emit, repeat. Never computes on packed words directly — the deliberate compute-unpacked/store-packed boundary.

`SmallFloats.with_default_projection` – Function.

```
with_default_type(f, args...)      -> f(DefaultType(), args...)
with_default_returntype(f, args...) -> f(DefaultReturnType(), args...)
with_default_projection(f, args...) -> f(DefaultProjection(), args...)
```

Call `f` with the named session default as its first argument — the supported way to *consume* a default. While the default still holds its initial value (the overwhelmingly common case) the call is statically compiled against that constant: no dynamic dispatch, `f` fully specialized. After the default is changed, the call crosses a function barrier instead: one dynamic dispatch, everything inside fully specialized.

Allocation contract: the combinator's return type unions both paths. When `f`'s result type does not depend on the default — e.g. `with_default_projection` with the formats fixed by the caller — the union is concrete and the call is **zero-allocation with a concretely inferred result**. When the result's type *is* the default (`with_default_type((T, x) -> T(x), ...)`), the value is computed on the specialized path but boxes once at escape — the irreducible cost of a runtime-chosen type.

```
julia> with_default_type((T, x) -> T(x), 1.5)
Binary8p2se(1.5 ≡ 0x41)

julia> with_default_projection((p, x, y) -> Add(Binary8p4se, p, x, y),
                               Binary8p4se(1.5), Binary8p4se(0.25))
Binary8p4se(1.75 ≡ 0x46)
```


SmallFloats.with_default_returntype - Function.

```
with_default_type(f, args...)      -> f(DefaultType(), args...)
with_default_returntype(f, args...) -> f(DefaultReturnType(), args...)
with_default_projection(f, args...) -> f(DefaultProjection(), args...)
```

Call `f` with the named session default as its first argument — the supported way to *consume* a default. While the default still holds its initial value (the overwhelmingly common case) the call is statically compiled against that constant: no dynamic dispatch, `f` fully specialized. After the default is changed, the call crosses a function barrier instead: one dynamic dispatch, everything inside fully specialized.

Allocation contract: the combinator's return type unions both paths. When `f`'s result type does not depend on the default — e.g. `with_default_projection` with the formats fixed by the caller — the union is concrete and the call is **zero-allocation with a concretely inferred result**. When the result's type *is* the default (`with_default_type((T, x) -> T(x), ...)`), the value is computed on the specialized path but boxes once at escape — the irreducible cost of a runtime-chosen type.

```
julia> with_default_type((T, x) -> T(x), 1.5)
Binary8p2se(1.5 ≡ 0x41)

julia> with_default_projection((p, x, y) -> Add(Binary8p4se, p, x, y),
                               Binary8p4se(1.5), Binary8p4se(0.25))
Binary8p4se(1.75 ≡ 0x46)
```

SmallFloats.with_default_type - Method.

```
with_default_type(f, args...)      -> f(DefaultType(), args...)
with_default_returntype(f, args...) -> f(DefaultReturnType(), args...)
with_default_projection(f, args...) -> f(DefaultProjection(), args...)
```

Call `f` with the named session default as its first argument — the supported way to *consume* a default. While the default still holds its initial value (the overwhelmingly common case) the call is statically compiled against that constant: no dynamic dispatch, `f` fully specialized. After the default is changed, the call crosses a function barrier instead: one dynamic dispatch, everything inside fully specialized.

Allocation contract: the combinator's return type unions both paths. When `f`'s result type does not depend on the default — e.g. `with_default_projection` with the formats fixed by the caller — the union is concrete and the call is **zero-allocation with a concretely inferred result**. When the result's type *is* the default (`with_default_type((T, x) -> T(x), ...)`), the value is computed on the specialized path but boxes once at escape — the irreducible cost of a runtime-chosen type.

```
julia> with_default_type((T, x) -> T(x), 1.5)
Binary8p2se(1.5 ≡ 0x41)

julia> with_default_projection((p, x, y) -> Add(Binary8p4se, p, x, y),
                               Binary8p4se(1.5), Binary8p4se(0.25))
Binary8p4se(1.75 ≡ 0x46)
```

Chapter 38

Internal API Index

38.1 Internal API Index

This index renders docstrings for unexported machinery used by the implementation and contributor guides. These names are not covered by the public compatibility contract. Begin with [Architecture and Invariants](#) and follow its source map before depending on an internal entry point.

Public contracts are collected in [Public API Index](#).

`SmallFloats.DEFAULT_RBITS` – Constant.

Default random-bit budget used by no-argument stochastic projection constructors.

`SmallFloats.ExactExternalFloat` – Type.

External float element types whose widening to the `Float64` carrier is exact.

`SmallFloats.F32_EXACT_MAX_PAIRS` – Constant.

Largest $2^{(K1+K2)}$ operand-pair count `f32_exact` will enumerate.

A bound on **time**, not memory: each pair costs two decodes and two 300-bit `BigFloat` operations, so the pair count — not any byte figure — is what governs it. The same distinction `TABLE_EAGER_BITS` draws against `TABLE_MAX_BITS`.

The floor is 2^{16} , the largest $K \leq 8$ signature ($8 + 8$), so **no answer that existed before the $K \leq 16$ extension can change** — the fixed-point argument `TABLE_EAGER_BITS` makes for itself. 2^{20} above it admits every signature through $K1 + K2 = 20$ while keeping the worst case seconds rather than hours; 2^{32} (two $K = 16$ formats) is four billion pairs and is what this exists to refuse.

`SmallFloats.KMAX` – Constant.

Highest bitwidth this package implements.

`SmallFloats.KMIN` – Constant.

Lowest bitwidth the draft defines ($K = 3$: sign, one exponent bit, one significand bit).

`SmallFloats.KSPLIT` – Constant.

Largest K whose code point fits a `UInt8` — the `Code8/Code16` seam.

`SmallFloats.MaybeR` – Type.

An explicit stochastic draw R , or nothing to draw one from the `rng`.

SmallFloats.TABLE_BUILD_MIN_ENTRIES – Constant.

Minimum entry count before a table build parallelizes.

Its own threshold rather than `THREAD_MIN_ELEMS`, because the units differ: an array element costs a decode plus one operation, while a table entry costs a full scalar-engine trip including any MPFR escalation. One number serving both would be conflating two different costs that happen to be counted in the same kind of integer.

SmallFloats.TABLE_EAGER_BITS – Constant.

Largest table the package will *build*, as $\log_2(\text{entries})$. $16 = 65\,536$ entries.

A bound on **time**, not memory: a table costs one scalar trip per entry, at a measured $\sim 1\,\mu\text{s}/\text{entry}$, so this band caps first-call latency at $\sim 0.1\,\text{s}$ while 2^{24} entries would be tens of seconds — regardless of how comfortably 32 MiB fits in RAM.

16 is not a guess — it is exactly the largest table the $K \leq 8$ grid already builds (a binary 8×8 table, 65 536 entries), so **no $K \leq 8$ decision changes** and G5 stays byte-identical by construction. It also admits the case that matters most at wide K : a unary table for a $K = 16$ format is 2^{16} entries, the same count, now 128 KiB instead of 64 KiB because the entries are `UInt16`.

SmallFloats.TABLE_MAX_BITS – Constant.

Hard ceiling on a single table, as $\log_2(\text{bytes})$ — the byte budget invariant 10 requires. $25 = 32\,\text{MiB}$.

Compared **as bits**. Computing $1 \ll \Sigma K$ in order to decide whether $1 \ll \Sigma K$ is too large is exactly the bug the budget exists to prevent: at $\Sigma K = 64$ the shift wraps to zero and the impossible allocation looks free.

SmallFloats.TERNARY_ADAPTIVE_BITS – Constant.

$K_1+K_2+K_3$ up to which a ternary table may be built adaptively once the signature has processed `TERNARY_BUILD_ELEMS[]` elements (default 21 bits = 2 MiB — the $K=7$ band). Above this, ternary ops always run the compute kernel.

SmallFloats.TERNARY_BUILD_ELEMS – Constant.

Cumulative element count at which an adaptive-band signature earns its table.

SmallFloats.TERNARY_CACHE_BYTES – Constant.

Byte budget for the ternary cache; least-recently-used tables evict first.

SmallFloats.TERNARY_EAGER_BITS – Constant.

$K_1+K_2+K_3$ up to which a ternary table is built eagerly on first array call (default 18 bits = 256 KiB — covers every all- $K \leq 6$ signature).

SmallFloats.THREADED_KERNELS – Constant.

Master switch for threaded compute loops, at every arity.

SmallFloats.THREADED_TABLE_BUILDS – Constant.

Master switch for parallel table builds.

SmallFloats.THREAD_MIN_ELEMS – Constant.

Minimum element count before a pure-p compute loop threads.

Core.Type - Method.

```
(T::Type{<:Binary})(c::Unsigned) -> T
```

Construct a value from its **code point** (validated: *c* must satisfy the representation invariant for the format's *K*, or an `ArgumentError` is thrown; the check folds away when *K* equals the storage width and for constant codes). `Binary5p3sf(0x08) == Binary5p3sf(1.0)`.

`Unsigned` is the argument-type *class* with code-point semantics, at every width — every other `Real` (including signed `Integers`) constructs by projecting the numeric value. `Binary8p4se(0x02)` is code point 2 (the second-smallest subnormal), while `Binary8p4se(2)` is the number 2.0. Widening an `Unsigned` is lossless, so the only way a wrong-width argument can be wrong is by being out of code range — which throws. `rawvalue(T, c)` remains the unchecked kernel-internal route.

SmallFloats.BigExactF - Type.

Deferred exact result: `f()` returns the value as an exact `BigFloat`.

SmallFloats.Code16 - Type.

```
Code16{K,P,SGN,EXT} <: Binary{K,P,SGN,EXT}
```

The two-byte representation — every format with $KSPLIT < K \leq KMAX$. Differs from its `Code8` sibling only in the width of the word carrying the code point; all semantics live in the `Binary` parameters.

SmallFloats.Code8 - Type.

```
Code8{K,P,SGN,EXT} <: Binary{K,P,SGN,EXT}
```

The one-byte representation of a `P3109` format — every format with $K \leq 8$. Differs from its `Code16` sibling only in the width of the word carrying the code point; all semantics live in the `Binary` parameters.

SmallFloats.ComputeDecode - Type.

$2^K > 256$: compute from the code point; a 65 536-element tuple literal is not a table, it is a compile-time hazard (invariant 10).

SmallFloats.DecodePolicy - Type.

Decode strategy for a format: a 2^K -entry constant tuple, or computed.

SmallFloats.Enclose128F - Type.

Correctly-rounded `Float128` bracket (plan Site C, Class R): IEEE mandates correct rounding for `Float128` /, `sqrt`, `fma`, so a nearest-CR result of a value known inexact brackets the truth in the *open* interval $(lo, hi) = (prevfloat(q), nextfloat(q))$ with no envelope assumption. *f* is the MPFR fallback for (near-impossible at $P \leq 8$) grid-straddling brackets.

SmallFloats.EncloseF – Type.

Deferred enclosure, resolved by `_finish` through up to three stages, cheapest first:

1. `yd` — an *eager* Float64 estimate (`NaN` $\Rightarrow$ absent) whose libm evaluation is faithful (≤ 1 ulp); the true value is taken to lie within $|yd| \cdot 2^{-45}$ (`_F64_RELEXP`), $\geq 2^7$ slacker than any faithful-libm error.
2. `fq` — a zero-argument Float128 estimator (plan Site D, Class E; `nothing` $\Rightarrow$ absent): true value within $|y| \cdot 2^{-90}$ of $y = fq()$ (`_F128_RELEXP`), a bound $\geq 2^{18}$ slacker than any published libquadmath transcendental error, discharged empirically by the differential-build tests.
3. `f(prec)` — the rigorous MPFR ladder: directed `(lo, hi)::NTuple{2, BigFloat}` strictly bracketing the true value.

Each estimate stage returns only when the two-sided sticky projection of its envelope agrees on a single code point; any disagreement (or an absent/degenerate estimate) falls through to the next stage. Stage 3 always decides.

SmallFloats.Head – Type.

Head

Evaluation-carrier selector: a singleton tag, so carrier choice is a method lookup rather than a value the compiler must fold. Rungs are ordered `HeadF64 < HeadF128 < HeadExact` by carrier width.

SmallFloats.HeadExact – Type.

Rung 3 — unbounded exponent. `BigFloat` until Stage 7, Dyadic after; the tag is the seam that makes that swap a one-line change (gate G7 pins that the two agree).

SmallFloats.HeadF128 – Type.

Rung 2 — Float128 arithmetic (the existing escalation machinery, promoted to entry point). Allocation-free: Float128 is a bitstype.

SmallFloats.HeadF64 – Type.

Rung 1 — Float64 arithmetic. Every $K \leq 8$ format and 432 of the 504.

SmallFloats.Rounded – Type.

Result of `ωRoundToPrecision` in canonical integer form: a kind tag plus $\text{sign} \cdot S \cdot 2^Q$ for finite results (design §5.2). The saturation stage consumes this without ever re-touching the carrier value.

SmallFloats.StickyF – Type.

Sticky-head exact result (wide-spread tail, Float128 revision follow-on): the true value is $v + \text{sgn} \cdot \varepsilon$ for an infinitesimal $\varepsilon > 0$ — v carries every bit the projection can consume and $\text{sgn} \in \{-1, +1\}$ the direction of the neglected tail. Sound because v is either exactly on a rounding threshold of the target grid (where *sticky* decides, for every mode including stochastic sub-grids) or strictly farther from the nearest threshold than the tail magnitude; the emitting sites in `oracle.jl` discharge that bound (operand significands ≤ 17 bits, target grids $\leq P-1+N \leq 67$ fractional bits, spreads $>$ the DE^* thresholds). Non-allocating: replaces the `BigFloat` escalation for FMA/FAA.

SmallFloats.TableDecode – Type.

$2^K \leq 256$: the generated constant tuple folds and costs one load.

`SmallFloats.TableKey` – Type.

Value key for the table cache: `op` + format parameters + `p` parameters (all values, so the cache itself is type-stable and non-allocating to query).

`SmallFloats.TernaryKey` – Type.

Value key for a ternary table: `op` + result + three operand formats + `p`.

`Base.nextfloat` – Method.

`NextGreaterThan(v)` (draft §4.16): the least datum greater than `v` in the total order — one step up the code lattice, with `NaN` → `NaN` and `MaxFinite/+Inf` → `NaN` at the top. `Base.nextfloat` on `Binary` is this operation.

`Base.prevfloat` – Method.

`NextLessThan(v)` (draft §4.16): the greatest datum less than `v` in the total order — one step down the code lattice, with `NaN` → `NaN` and `MinFinite/-Inf` → `NaN` at the bottom. `Base.prevfloat` on `Binary` is this operation.

`Base.widen` – Method.

The public promotion target for `T` against external numeric types. Always a Julia `AbstractFloat`, never the internal exact carrier.

`SmallFloats._bp_element` – Method.

One element of `ωBlockProject`: the draft's `S`-special rows, then `ωDivide` + `ωProject`.

Fast-path cascade after the special rows, ordered cheapest-first by result kind: exact `Float64` quotient → `CR Float128` quotient (exact or one-ulp bracket) → `CR-divided Enclose128F` bracket → `fq-composition` envelope (2^{89}) — each resolved by the two-sided sticky gate; any miss falls through to the rigorous MPFR interval at the bottom.

`SmallFloats._check_tabulable` – Method.

Stochastic `p` is a distribution over `R`, never a table (design §5.4).

`SmallFloats._comparable` – Method.

Numeric comparison is defined only when neither operand is `NaN` — every `==/</>=` below returns `false` otherwise, per IEEE unorderedness.

`SmallFloats._extexprange` – Method.

(least subnormal exponent, greatest normal exponent) of an external float.

`SmallFloats._resolve_rng` – Method.

Resolve a caller-supplied `rng`, falling back to the task-local default. Call sites hoist this out of loops so array kernels resolve once per call, not per element.

`SmallFloats._rng_for` – Method.

The `rng` an operation under `p` will actually draw from — nothing for pure `p`, which must never touch RNG state.

`SmallFloats._rungindex` – Method.

Per-format rung index from the exponent bias. A pure `Int` function of the type parameters, so `Val(_rungindex(F))` is a compile-time constant.

Stated on `2B` rather than `B` so it is visibly the $n = 2$ case of `_rungindex_span` and not an independent constant pair. $B \leq 512 \iff 2B \leq 1024$ and $B \leq 8192 \iff 2B \leq 16384$, so the 432/64/8 partition is unchanged — `G9` enumerates it.

`SmallFloats._rungindex_span` – Method.

Rung index for a monomial whose factor biases sum to ΣB . This is the primitive form of the boundary rule stated above; the per-format answer is its $n = 2$ same-format case.

`SmallFloats._ternary_table_for` – Method.

```
_ternary_table_for(op, fr, f1, f2, f3, p, nelems) -> Union{Nothing, Memory{UInt8}}
```

The kernel-facing policy gate, called once per array operation with the call's element count. Eager band ($\sum K_i \leq \text{TERNARY_EAGER_BITS}[]$): fetch/build now. Adaptive band ($\leq \text{TERNARY_ADAPTIVE_BITS}[]$): return the cached table if present, otherwise accumulate `nelems` against the signature and build only once it has earned `TERNARY_BUILD_ELEMS[]`. Beyond the adaptive band (the $K=8$ territory): always nothing — the compute kernel is the right tradeoff there.

`SmallFloats._unitmask` – Method.

```
_unitmask(U, K) -> U
```

The low- K -bits mask of storage unit U , built by **complement** — shift down from `typemax`, never up from one.

Julia defines a shift at or beyond a type's width as zero, so `UInt8(1) << 8` and `UInt16(1) << 16` are both 0, and 2^K computed in a unit of width K is 0. The format grid **always contains a format whose K equals its storage width** ($K = 8$ on `UInt8`, $K = 16$ on `UInt16`), so that case is structural, not an edge. Here the shift amount is $\text{width} - K \in [0, \text{width}-1]$, in which the pathological amount is unreachable *by construction*, at every K , for every unit — including a future `Code32`.

`SmallFloats._unsupported` – Method.

Refuse a Base verb with a reason. Absence and refusal must not look alike.

`SmallFloats.apply_op` – Method.

`apply_op(Val(op), fr, p, R, xs...)` — evaluate `op`'s ω -semantics on decoded operands and project the result into `fr` under `p` (with random bits `R`).

The evaluation carrier comes from the **operands**, not from `fr`. That is `carriers.jl`'s stated rule — the carrier constrains what `weval` may form, never what may be projected into a format, because `round_to_precision` works in (P, B) integer space and is exact for any exact input. Stage 3 selected on `rung(fr)` instead, as a deliberate stand-in while decode still refused $K \geq 9$; keeping that would have refused `Add(Binary16p2se, p, x8, y8)`, whose operands are ordinary `Float64` datums and whose exact sum needs no wide carrier at all.

Completeness of the operand-only rule rests on one premise, stated because Stage 3's version of it is what M28 was: the spec register's maximum factor count is **2**, and the `rung` boundaries are stated on 2B precisely so that a datum on carrier C guarantees its square is also on C . Beyond two factors — `ScaledMultiply`'s four — the monomial bound stops following from the operand types and `rung(op, Fs...)` must be used with the formats in hand. That is why `blocks.jl` calls the join and this does not.

`SmallFloats.bigprec_prod` – Method.

```
bigprec_prod(xs...) -> Int
```

`bigprec` for an expression whose leading term is the **product** $x_1 \cdot x_2$ and whose remaining terms are added to it — the FMA shape. The product's exponent is $e_1 + e_2$, which lies outside the range of the operands' own exponents, so the sum method would understate the spread. Its significand is $P_1 + P_2$ bits wide.

`SmallFloats.blockdecode` – Method.

ω BlockDecode (draft §5.1.1): per lane ω Multiply(decode(s), decode(x_i)), exact on the carrier the block's two formats jointly require.

The carrier is keyed on `rung(Val(:Multiply), FS, FE)` — **not** on either format's own rung, and that distinction is the whole content of this method. A block's scale and element formats are independent, and what has to fit is their *product*: `carriers.jl`'s rule is $\Sigma B_i \leq e_{\max}$, so two formats that are each rung 1 at $2B \leq 1024$ can compose into a lane product needing $B_S + B_E$ well past it. Asserting `carriertype(rung(FE))` would be wrong in the worst way available here — the overflow produces $\pm\text{Inf}$, a plausible special value, rather than an error.

The assertion stays because it is load-bearing for more than documentation: it is what keeps the returned `NTuple` concretely typed, and with it the reductions downstream allocation-free. It now names a carrier computed from both formats instead of a constant.

`SmallFloats.blockproject` – Method.

ω BlockProject (draft §5.1.2) over a tuple of result kinds, scale supplied as a value.

`SmallFloats.codemask` – Method.

The low-K-bits mask of a format's storage unit (representation invariant 3): the code point occupies the low K bits, the high bits are maintained zero.

`SmallFloats.datumcarrier` – Method.

The exact carrier for the *datums* of F: what `decode(v::F)` returns.

`SmallFloats.datumsexact` – Method.

```
datumsexact(X, F) -> Bool
```

true when **every** datum of format F is exactly representable in the external float type X.

Three conditions, one per way a conversion can lose information:

- $P \leq \text{precision}(X)$ — the significand fits;
- the largest finite datum's exponent is within X's normal range. That exponent is $B - 1$ for **both** signedness classes, and it is $B - 2$ when $P == 1$, so $B - 1$ is a sound upper bound at every format: `Binary8p4uf`'s largest datum is $14 \cdot 2^{12} = 57\,344$, exponent $15 = B - 1$, and `Binary16pluf`'s is 2^{32766} with a least positive datum of 2^{-32767} ;
- $2 - P - B \geq X$'s least subnormal exponent — and the smallest datum being a *subnormal* of X is fine, because a datum is $\text{sig} \cdot 2^{(2-P-B)}$ with integer `sig`, so once the step is a multiple of X's subnormal step every datum lands on X's grid exactly.

`SmallFloats.decodepolicy` – Method.

Decode strategy: a $2^K \leq 256$ constant tuple, or computed from the code.

`SmallFloats.encode` – Method.

```
encode(T, sign, S, Q) -> UInt8 (private; design §3.3)
```

ω Encode from canonical integer form: $\text{value} = \text{sign} \cdot S \cdot 2^Q$, with $S \in 0:2^P$ ($S == 2^P$ is the next-binade carry, draft §4.7.4 NOTE 4). Precondition: the value is in the datum set of T (guaranteed by `RoundToPrecision ◦ Saturate`).

`SmallFloats.get_table` – Function.

```
get_table(op, fr, f1, [f2,] p) -> Memory{codeunit_type(fr)}
```

Fetch (building and caching on first use) the complete result table for the pure- p specialization $op(fr; f1[, f2])$ under p : entry $c + 1$ (unary) or $(c1 \ll K2) + c2 + 1$ (binary) holds the result code point for those operand code points. Throws for stochastic p , and throws when the table would exceed the byte budget — its contract is to return a table, so it says so rather than returning nothing. Callers that can fall back want `table_for`.

@noinline by design: kernels call this once per array operation and index the returned table in their hot loop.

`SmallFloats.get_table` – Method.

```
get_table(op, fr, f1, f2, f3, p) -> Memory{UInt8}
```

Ternary form: entry $((c1 \ll K2 \mid c2) \ll K3) + c3 + 1$ holds the result code point. Builds unconditionally (policy lives in `_ternary_table_for`); pure p only.

`SmallFloats.joinhead` – Method.

The wider of two heads. The carrier join is a max over this order, which is what lets block and scaled operations compose without a second rule.

`SmallFloats.nan_order_key` – Method.

```
nan_order_key(F)
```

The order key reserved for the single NaN — **below** every finite key and $\pm\text{Inf}$.

Key 0 is free by construction: every datum's key is at least 1 (`order_key` adds 1 above the code arithmetic on both the signed and unsigned paths), so `zero(orderkeytype(F))` sits strictly under the whole datum lattice without shifting any other key.

Per format, not a constant, and the key type is `orderkeytype(F)` rather than `UInt16`. The reason is arithmetic, not tidiness: a format has 2^K code points mapping to keys $1 \dots 2^K$, so `UInt16` runs out at exactly $K = 16$ — `UInt16(c) + UInt16(1)` for $c = 0xffff$ wraps **silently to 0**, which is now precisely the NaN key: the largest datum would compare below NaN and below every other datum. `orderkeytype` returns `UInt32` for `Code16`, which has room for keys $1 \dots 2^{16}$ above the NaN sentinel.

(§4 Stage 4 item 3, pulled forward into Stage 3 — §11 M16. The rest of Stage 3 makes wide formats fail loudly; a wraparound in the total order is the one $K \geq 9$ defect that would have been silent, so it is closed here rather than carried for a commit.)

`SmallFloats.orderkeytype` – Method.

Unsigned type wide enough for a format's $2^K + 1$ total-order keys.

`SmallFloats.packing_saves` – Method.

```
packing_saves(F) -> Bool
```

Whether `PackedVector{F}` is smaller than `Vector{F}` — false exactly when `bitwidth(F)` equals the storage unit's width, i.e. $K = 8$ and $K = 16$.

Why this is a query and not a refusal. §4 Stage 4 item 5 called for `PackedVector` to refuse loudly at $K = 16$, on the grounds that packing is the identity there and a silent identity invites benchmark confusion. The grounds are right and the boundary is not: packing is *equally* the identity at $K = 8$, which has shipped since before the extension and is enumerated by G5's packed section over all 120 formats. A rule that fires on one of two identical cases is not a rule, and the alternative — refusing both — would be a breaking change justified by tidiness.

So the fact is exposed instead of enforced. Benchmarks and the docs read it; `PackedVector` at those widths stays correct, stays supported, and stays pointless. (§11 M21.)

`SmallFloats.project` – Method.

```
project(T, ρ, X; R=0, sticky=0) -> T
```

Project a closed-extended-real carrier value into format `T` under ρ . $X::\text{Float64}$ must be *exact* (the `Float64` fast path); $X::\text{BigFloat}$ likewise (or an interval endpoint with `sticky` $\in \{-1, +1\}$). R is the stochastic draw.

`SmallFloats.promotecarrier` – Method.

The public promotion target for `F` against external numeric types. Always a Julia `AbstractFloat`, never the internal exact carrier.

`SmallFloats.rawvalue` – Method.

Unsafe raw constructor used by kernels after invariants are established.

Accepts an abstract format and normalizes through `reptype`; the type is passed *as an argument* to `_rawvalue` rather than assigned to a local, so it becomes a type parameter of the callee instead of a value the compiler has to fold (technique T6). Kernel call sites pass concrete types and pay nothing.

`SmallFloats.rung` – Method.

```
rung(op::Val, Fs::Type{<:Binary}...) -> Head
```

The carrier for evaluating `op` over operands of formats `Fs`. A **join of two lower bounds**, and both are necessary:

- the **monomial bound** — `op`'s exact result is a sum of monomials in at most `opfactors(op)` datum factors, so it needs `_rungiindex_span(nf · maxB)` ($n_f \cdot \max B$ rather than $\sum B_i$ over the actual operands: it dominates every monomial shape, including x^2 from `Hypot`, which $\sum B_i$ does not, and it coincides with the tight bound in the same-format case that is virtually all traffic);
- the **operand bound** — `decode(v::Fi)` has already produced a value on `datumcarrier(Fi)`, so the head can never be narrower than any operand's own `rung` regardless of what the monomial needs.

Omitting the second is the subtle error. `Abs` has one factor, so the monomial bound alone would pick `rung 1` for a $B = 1024$ format — whose datums arrive as `Float128`, for which `weval(::HeadF64, ...)` has no row. The join makes `lift`'s missing narrowing method a consequence rather than a hazard: the head dominates every operand, so nothing ever needs narrowing.

For the same format at two factors this agrees with `rung(F)` exactly, which is what keeps the 432/64/8 partition and every $K \leq 8$ result unchanged.

`SmallFloats.rung` – Method.

```
rung(F) -> Head
```

The starting carrier for evaluating on datums of format `F` alone. For an operation over several formats use `rung(op, Fs...)`, which additionally accounts for the operation's factor count — a `ScaledMultiply` of two Group A formats can need a wider carrier than either operand does.

`SmallFloats.rungindex` – Method.

Rung index 1/2/3 of a head — for ordering and for the carrier join.

`SmallFloats.table_for` – Method.

```
table_for(op, fr, f1[, f2], p) -> Union{Nothing, Memory{codeunit_type(fr)}}
```

The kernel-facing gate for unary and binary specializations: the table if policy grants one, nothing if the caller should run the scalar path per element.

Separate from `get_table` on purpose — **one name per question**. `get_table`'s contract is "return the table", so it throws when it cannot honour that, and that is what the suite, conformance and the precompile workload want. `table_for` asks the different question "should there be a table at all", and answers nothing without prejudice. Its `Union` return costs nothing because kernels call it once per array operation and branch outside the loop.

`SmallFloats.table_keys` – Method.

Keys of every cached unary/binary specialization, in a deterministic order.

`SmallFloats.tablebits` – Method.

$\log_2$ of the bytes a table over `Fs` with results in `fr` would occupy.

`SmallFloats.tablecache` – Method.

The unary/binary table cache holding results of format `fr`.

`SmallFloats.ternarycache` – Method.

The ternary table cache holding results of format `fr`.

`SmallFloats.unpack_tile!` – Method.

Unpack `pv[first:first+len-1]` into `buf` (byte scratch); tile-granular kernel entry.

`SmallFloats.within_byte_budget` – Method.

Would this table fit the byte budget? (Memory, not build time.)

Part VII

Technical Guide

Chapter 39

Guide Overview

39.1 Technical Guide

This is the maintainer's route through SmallFloats.jl. One semantic promise — compute the defined result, then project it once — is preserved across two code-unit representations, three carrier rungs, and every acceleration layer, and each page of this guide answers one question that preservation raises. Read the [User Guide](#) first; this guide assumes its vocabulary.

Two questions organize the implementation

[Architecture and Invariants](#) is the entry point. It separates two questions that must never be conflated:

1. **What carries an exact datum or intermediate?** The format selects a carrier rung — Float64, Float128, or the exact dyadic carrier — from its range requirements.
2. **Why is an operation result justified?** The oracle attaches a rigor class, from unconditional exactness through validated envelopes.

That separation is what lets Code8 and Code16 share public semantics while using different storage, decode, table, and compute policies.

Follow one datum through the package

How does a code point become a value, and a value a code point? [Encoding and Decoding](#) owns the format-to-representation seam, decodepolicy, the exact carriers, encode's code-unit-generic output, and the total-order keys. [Projection Engine](#) is the other direction: RoundToPrecision → Saturate → Encode, the only route that writes a result code point.

Where does the exact result come from? [Oracle and Rigor Classes](#): exact arithmetic, the Dyadic/Rational bridge, interval agreement, and conservative fallback. The relationship is directional — the oracle establishes a value or an enclosure; the projection engine decides the code point. Neither is a second definition of the other.

Performance layers are semantic layers first

Why is caching sound? [Function Tables and Array Kernels](#): a table entry is the scalar path memoized, and admission is a resource policy, never a semantic one. *Why are fused reductions exact without a superaccumulator?* [Exact Block Reductions](#) makes that argument in full. [Performance Model](#) and [Benchmark Correctly](#) turn those internals into measurement practice without promoting one machine's timings to a guarantee.

Keep an extension honest

Use [Add an Operation](#) for the implementation checklist and its worked examples, [Verify Custom Code](#) for an independent audit route, and [Verification Strategy](#) for the package-wide doctrine — enumerate rather than sample, and say "sampled" in so many words where enumeration is unaffordable. The [Documentation Claim Ledger](#) tracks each implementation-sensitive claim in these pages back to its source authority. An extension is complete only when its semantics, projection, applicable table policy, evidence, documentation, and benchmark scope agree.

See also

[Technical Examples](#), [Performance Evidence](#), and the [Internal API Index](#).

Chapter 40

Architecture and Invariants

40.1 Architecture and Invariants

The implementation is organized around two independent classifications: the carrier rung needed to hold a format exactly and the rigor class attached to a computed result. The remaining internal pages follow values through those classifications and into the one projection path.

The first axis is the **carrier rung** a format's datums decode to, chosen by the format's exponent bias: Float64 for the 432 formats whose full dynamic range fits Float64's normal range (rung 1), Float128 for formats that need more exponent room but stay in binary floating point (rung 2), and an exact dyadic carrier for the formats — 72 of them — whose range or precision breaks even Float128 (rung 3). The choice is a dispatch decision derived from the exponent bias, never a caller's, and every rung is exact for the datums it carries. The second axis is the **rigor class** attached to every non-Float64 computed result: Class R (unconditional — provable without any accuracy assumption) and Class E (envelope-conditional — faithful-but-not-CR estimates validated by a generous, measured slack). Together the two axes answer, for any operation on any format, "what carries the value" and "what justifies the result."

The pages that follow work through the machinery each axis produces: the encoding and projection engine that every carrier rung feeds into ([Encoding and Decoding](#)), the oracle and its two rigor classes ([Oracle and Rigor Classes](#)), the table and kernel layer that makes the common rung fast ([Function Tables and Array Kernels](#)), the block layer's exactness argument ([Exact Block Reductions](#)), and the [Verification Strategy](#) that keeps all of it honest.

Layer map

Source files load in dependency order, each layer speaking only downward:

layer	file	provides
formats	formats.jl	Binary{K,P,SGN,EXT}, the 504 named aliases, Group M queries
display	show.jl	value/code-point rendering and IOContext display policy
carriers	carriers.jl	exact in-flight carrier rungs selected from format range
specs	projspec.jl	rounding/saturation singletons, ProjSpec{R,S}, the predefined spec grid
de-	defaults.jl	settable session defaults (DefaultType, DefaultProjection, ...) behind Refs; consumed via the speculation guard
faults		guard
codec	decode_ encode.jl	decode (generated tables + bit-composed compute), encode, order keys, counting sort, Class, Next ops
engine	project.jl	round_to_precision (mask-based Float64 core + generic core), saturate, project, project_interval
ops	ops_scalar.jl	result-kind protocol, apply_op, the operation registry, the spec register
compat	juliacompat.jl	the Base register: declarative op $\Rightarrow$ Base-function map
oracle	oracle.jl	ω -semantics for all 52 operations
tables	tables.jl	the pure-p result-table cache (unary/binary + bitwidth-gated ternary)
kernels	kernels.jl	Shape-A gathers, Shape-B scalar/threaded loops, vmap
blocks	blocks.jl	Block, block/scaled ops, exact reductions
packed	packed.jl	sub-byte PackedVector storage
approx	approx.jl	κ measurement/registry, conformance declaration
rand	rand.jl	Random-API hooks: uniform-[0,1) floor projection, clamped normal

See also

[Encoding and Decoding](#), [Oracle and Rigor Classes](#), [Function Tables and Array Kernels](#), [Exact Block Reductions](#), [Verification Strategy](#).

Chapter 41

Encoding and Decoding

41.1 Encoding and Decoding

The public promise is simple: a code point identifies one exact datum, and decoding that identity never rounds. The implementation chooses different carriers and decode strategies by format, but those choices may change cost, never meaning.

Invariant

For every valid code point c of a concrete representation T :

```
codepoint(SmallFloats.rawValue(T, c)) == c
```

and `decode(SmallFloats.rawValue(T, c))` equals the P3109 datum assigned to c . For every non-NaN datum, encoding the exact decoded value returns the corresponding code point under an exact-grid projection.

Representation and storage

`Binary{K,P,SGN,EXT}` describes the abstract format. A concrete alias such as `Binary8p4se` supplies the machine representation. Code points use `UInt8` at $K \leq 8$ and `UInt16` above that split; unused high bits are not part of the format's identity.

The bit layout follows the draft: sign where present, biased exponent, and trailing significand. The negative-zero slot names the format's single NaN; there is one zero rather than signed zeros.

Decode policies

`decodpolicy(T)` selects an implementation shape.

For $K \leq 8$, decoding uses a generated constant tuple containing all datums. The tuple is built from the computational decoder, so lookup and computation have one source of meaning. A runtime decode becomes one indexed load, while a constant input can fold completely.

Above $K = 8$, materializing a 65,536-entry constant for every format is not a useful trade. The computational decoder runs directly.

Datum carriers

The decoded datum must be exact even when its magnitude exceeds `Float64`. `datumcarrier(T)` selects among three carrier rungs from format properties:

- rung 1 uses `Float64` where its normal range and precision cover every datum;

- rung 2 uses Float128 when greater exponent room is required;
- rung 3 uses an exact dyadic carrier when a hardware-style binary float cannot cover the format exactly.

The caller does not select the datum carrier. The invariant is stronger than “most values fit in Float64”: every selected carrier is exact for every datum of its format.

Ordering keys

Comparisons and counting sort use an unsigned integer key obtained by a sign-magnitude fold. Key zero is reserved for the single NaN, which P3109's total order places below the other datums. Every ordinary datum key is at least one.

The key type has room for the entire code lattice plus the NaN sentinel. The mapping is monotone, so comparison and sort can operate on integers without decoding floating values.

Source map

- format parameters, representations, masks, and carrier traits: `src/formats.jl`;
- computational decoding and order keys: `src/decode_encode.jl`;
- wide exact datum carriers: `src/dyadic.jl` and carrier support files;
- constructor normalization and representation aliases: format and conversion entry points in `src/`.

Evidence

The verification suite compares computational decoding against an independent big-float transliteration over the supported lattice. It also checks lookup versus computation, round trips, carrier widening, ordering monotonicity, and the format/code-point algebra.

Extension seam

A new representation or carrier policy must preserve the public invariant and add equivalence tests against an existing wider path. Do not make a carrier choice from runtime caller preference; that would turn representation into semantics.

Chapter 42

Projection Engine

42.1 Projection Engine

Projection is the only write path into a code point. Constructors, conversions, operations, table construction, and block results all converge here.

```
exact result → RoundToPrecision → Saturate → Encode
```

Public promise

Given a result format, projection specification, and exact mathematical value, the engine returns the P3109-defined code point. No faster path is allowed to substitute a second rounding routine.

RoundToPrecision

The rounding stage produces a canonical `Rounded(kind, sign, S, Q)`. *S* and *Q* describe a scaled significand exactly; *kind* distinguishes finite and special results.

Two implementations cover the supported carriers:

- the generic core scales an exact or enclosed carrier and compares the discarded fraction with the selected mode's decision points;
- the `Float64` mask core extracts sign, exponent, and significand fields and performs the same decisions with integer operations.

The mask core is an optimization only. Specials, subnormal carrier inputs, or cases outside its proof obligations fall back to the generic core. Equivalence tests establish that the two implementations choose the same rounded result.

Symbolic sticky

An interval or asymptotic computation sometimes knows that the true value is immediately above or below an exactly representable endpoint without possessing a distinct carrier value for that infinitesimal displacement.

`sticky ∈ {-1,0,+1}` carries that relation into every rounding comparison. The delicate “just below a datum” case crosses into the preceding binade and uses a fraction of `1-`; treating it as the endpoint itself would give the wrong answer for directed modes and ties. The observable effect: projecting carrier value `1.0` with `sticky = -1` under `TowardNegative` lands on `0.9375` — exactly `NextLessThan(1.0)`, across the binade boundary — where without `sticky` the projection of an exactly representable value is the identity. [The Exact-Then-Project Contract](#) walks that example.

Saturate

The saturation stage classifies the rounded result against the result format's range and applies the draft's saturation rows. A disposition can be:

- as-is;
- minimum or maximum finite;
- positive or negative infinity;
- NaN.

Classification uses canonical integer fields rather than an inexact carrier comparison. Signed and unsigned formats, finite and extended domains, rounding direction, and special operation results all participate in the row choice.

Encode

Encoding assembles the final code point with integer operations. It handles significand carry, normal/subnormal boundaries, exponent fields, special encodings, and the single-zero rule. Encode does not make another numerical decision; all policy has already been resolved.

Source map

- projection orchestration and saturation: `src/project.jl`;
- rounding cores and rounded representation: the rounding implementation files under `src/`;
- operation-to-exact-result dispatch: `src/oracle.jl` and operation files;
- encoding and raw value construction: `src/decode_encode.jl`.

Evidence

Required gates include deterministic and stochastic rounding predicates, generic/mask equivalence, boundary saturation rows, carrier escalation, symbolic-sticky cases, and exhaustive code-point results where the input space is affordable. Array tables are compared entry-for-entry with this scalar path.

Extension seam

An optimization may bypass work only after proving that it returns the same canonical rounded form or final code point. A new rounding or saturation mode belongs in the shared ProjSpec machinery and its exhaustive predicate tests, not in an operation-specific shortcut.

Chapter 43

Oracle and Rigor Classes

43.1 Oracle and Rigor Classes

Every operation's defined result is computed by `weval`, which returns one of seven result kinds; `apply_op` fast-splits the common one and finishes the rest:

kind	meaning	finished by
Float64	exact (specials; representable arithmetic)	direct project
Dyadic	exact at rung 3: an Int128 significand and an Int64 exponent, isbits	direct project (dyadic carrier)
Float128	exact by width analysis	direct project (Float128 carrier)
StickyF	wide-spread FMA/FAA tail: head value (Float64 or Float128) + tail sign	direct project with sticky set — no allocation
BigExactF	exact at a precision derived from the operands — bigprec, not a constant (wide-spread tail for Add; the near-impossible FAA distillation miss)	project on BigFloat
Enclose128	directly-rounded Float128 bracket	sticky agreement, MPFR fallback
EncloseF	MPFR directed enclosure <code>f(prec)</code> , optional Float128 pre-filter <code>fq</code> , optional eager Float64 estimate <code>yd</code>	three-stage: <code>yd</code> → <code>fq</code> → interval protocol

The Dyadic↔Rational bridge

The dyadic carrier converts to and from Rational **exactly**, and that bridge is the cleanest way to reason about a rung-3 value:

```
using SmallFloats.DyadicNumbers: dyadic_to_rational, rational_to_dyadic, isdyadic

dyadic_to_rational(x)           # Rational{BigInt}, exact — a Dyadic IS  $5 \cdot 2^Q$ 
rational_to_dyadic(3 // 4)      # exact, or throws
isdyadic(1 // 3)                # false — no Dyadic represents it
```

Three properties are worth knowing. The infinities convert ($\pm 1/\theta$, matching `Rational{BigInt}(Inf)`) and NaN does not, because Rational has no NaN slot. A non-dyadic rational is **refused rather than rounded** — rounding it here would be a rounding outside project. And the round trip is a law about *values*, not fields: Dyadic's significand is deliberately unnormalized, so $6 \cdot 2^{-2}$ and $3 \cdot 2^{-1}$ are the same value and round-trip to the canonical one.

The full design, with every edge case and the three laws, is in `other/dyadic_rational.md`.

The two rigor classes

Two **rigor classes** govern every non-Float64 path, and their arguments are never mixed:

Class R (unconditional) rests on two facts that hold without any accuracy assumption:

- **Width analysis.** Sums of decoded datums are *exactly representable* in Float128 whenever operand bits + exponent spread fit 113 bits — checked in advance by integer exponent arithmetic (Add at $\Delta E \leq 100$, FMA ≤ 92 , FAA span ≤ 98). Beyond the threshold, Add escalates to the exact MPFR path — at a precision **derived from the operands** by bigprec, not the retired 2200-bit constant, which was ample at $K \leq 8$ and insufficient for 50 of the 504 formats (gate G2 measures this) — while FMA/FAA take a non-allocating **sticky-head** shortcut: past that spread the smaller term is provably too small (by construction of the threshold) to affect anything but the tail direction of the larger one, so the result is StickyF(head, sign) — no BigFloat involved. FAA's three-term case runs a bounded Float128 2sum-distillation (Priest-style, ≤ 6 sweeps) to find that head/tail split directly; the residual BigExactF MPFR fallback exists only for the near-impossible case the distillation doesn't converge.
- **IEEE correct rounding.** Division and sqrt are correctly rounded at every binary width, by mandate. Divide/Recip therefore use the **Float64** quotient directly: it is within half an ulp of truth, so it serves as EncloseF's eager yd estimate with no Float128 arithmetic on the path at all. Sqrt/RSqrt use the **Float128** CR result: an inexact nearest-CR value q brackets the truth in the *open* interval ($\text{prevfloat}(q)$, $\text{nextfloat}(q)$). Exactness itself is detected by an fma residual test.

Class E (envelope-conditional). Faithful-but-not-CR evaluations stand in for an enclosure only inside a generous envelope, resolved by the two-sided sticky gate. Two stages, cheapest first:

- *Float64 stage:* for operations whose Float64 libm is faithful (≤ 1 ulp $\approx 2^{-52}$ relative — the exp/log families, the hyperbolics, forward trig inside the $|x| \leq 10^{15}$ reduction window, atan, asinh, the softplus composition), the eager estimate yd carries envelope $E = |yd| \cdot 2^{-4^5}$, $\geq 2^7$ slacker than any faithful-libm error (measured margin on this machine: $\geq 2^{6.6}$).
- *Float128 stage:* libquadmath's estimate $y = fq()$ carries $E = |y| \cdot 2^{-9^0}$, at least 2^{18} slacker than any published libquadmath bound.

In both stages, if the sticky-projected endpoints agree, that code point is the answer; if not, the next stage (ultimately the MPFR ladder) decides. An envelope failure can therefore only cost speed, never correctness — unless both endpoints agree *and* the envelope is wrong, which the differential-build tests rule out empirically by building every standard table twice (Float128-first and pure-MPFR) and by byte-diffing, and which the exhaustive oracle cross-check re-verifies per operation against the 3072-bit protocol.

The interval protocol

The interval protocol (`project_interval`) is the termination backbone: evaluate the MPFR enclosure at precision p ; if $lo == hi$ the value is exactly representable — project it; otherwise project lo with `sticky = +1` and hi with `sticky = -1`; agreement (projection is monotone at fixed R) yields the answer; disagreement doubles p , clamped to the maxprec ceiling (**derived from the format**, $\max(4096, \text{bigprec}(T) + 64)$ — a flat 4096 was ample at $K \leq 8$ and could not decide ArcCosPi at Binary15p2ue under a directed or stochastic rule, honored exactly even when it is not a power-of-two multiple of the 256-bit start). Termination requires that the enclosure not be chasing a value the grid can actually hit — which is why the π -scaled operations carry **Niven peels**: for dyadic arguments, $\tan(\pi r)$ takes rational values only at the quarter-integers (± 1) and $\text{atan2}/\pi$ only on the diagonals ($\pm 1/4$, $\pm 3/4$); those cases are answered exactly before any enclosure is built, and Niven's theorem proves the peel set complete.

Source and gates

The operation registry and ω evaluation live in `src/oracle.jl`; exact dyadic carriers and their rational bridge live in `src/dyadic.jl`, with carrier-rung selection in `src/carriers.jl`. Changes here must preserve enclosure containment, termination, and agreement with a wider independent path. The gate and tier files under `test/` record whether each check is exhaustive or sampled.

See also

[Encoding and Decoding](#), [Function Tables and Array Kernels](#), [Add an Operation](#).

Chapter 44

Function Tables and Array Kernels

44.1 Function Tables and Array Kernels

Tables are a cache for defined scalar semantics, not a second arithmetic implementation. For a pure projection specification, a unary or binary operation is a finite function on code points. When the policy admits that function, SmallFloats builds it through the scalar path once and gathers from it thereafter. When the policy declines it, the kernel calls that same scalar path per element. The choice changes cost; it must never change a code point.

Two representations, two concrete caches

A table entry is a result code point. Code8 results therefore use `Memory{UInt8}` and Code16 results use `Memory{UInt16}`. The package keeps separate caches rather than a Union-typed cache, so the representation is chosen by dispatch before the hot loop indexes its local memory.

```
pure p + admitted table → Shape A: gather from Memory{codeunit_type(fr)}  
stochastic p or declined table → Shape B: compute one defined scalar result
```

The branch happens once per array call, not once per element. `table_policy` reports the chosen shape, estimated entries and bytes, and the reason; it is the right diagnostic when an array call takes the compute path.

Admission is a resource policy

The table policy has separate limits for memory and build time. The byte budget accounts for the result code-unit width, so a Code16 table with the same number of entries as a Code8 table occupies twice the storage. The eager build limit bounds first-call latency rather than pretending that every table that fits in RAM is useful to build.

Unary and binary pure-spec tables are admitted only within both limits. Ternary tables have a third, reuse-sensitive policy:

- **Eager band.** Small ternary signatures build on their first array call.
- **Adaptive band.** Medium signatures accumulate use and build only after they have processed enough elements to amortize construction; the cache is byte-bounded and least-recently-used tables may be evicted.
- **Compute band.** Larger signatures remain Shape B.

These are current policy knobs, not semantic limits. `table_policy` and the cache-accounting API expose them so a benchmark or application can report what actually happened instead of inferring it from a timing.

Why Shape A and Shape B agree

One table entry is one call through `_scalar_code`. Shape B reaches that same defined scalar route with the element's decoded operands. The table builders and compute kernels are therefore two scheduling strategies for the same operation, and the verification gates compare them wherever a table is affordable.

This matters most outside the byte-sized sweet spot. A Code16 format can still participate in an affordable unary or mixed-width table; it is not categorically “compute only.” Conversely, a byte format can take Shape B when its arity, signature, projection, or cache policy makes a table unsuitable. Describe the actual signature and `table_policy` result, never a broad K-only slogan.

Threading and stochastic calls

Large pure-spec table builds partition independent output regions. Pure Shape-B loops may also thread when the index set, element count, and runtime thread configuration meet the kernel policy. The safety argument is concrete: each pure entry is independent, and exact fallback state is task-local.

Stochastic projection remains sequential. It consumes one RNG stream in index order, and preserving a seeded experiment's result takes priority over a scheduler-dependent speedup.

Sorting belongs to the same performance story

Same-format comparisons use integer order keys, not decoded floating values. The key type is representation-aware (UIInt16 for Code8, UIInt32 for Code16) and reserves zero for the NaN that P3109 orders first. Counting sort is available when the input is large enough to amortize its key-space setup; smaller inputs use Base's ordinary sort. This is another policy decision whose semantics are the total order, not a timing claim.

Evidence and extension seam

`src/tables.jl` owns keys, budgets, cache lifecycle, and table construction; `src/kernels.jl` owns Shape-A gathering and Shape-B execution. A change is complete only when it preserves concrete code-unit storage, byte accounting, scalar/table equivalence, stochastic ordering, and the relevant threaded versus sequential result checks. Record cold build, warm hit, and compute-path results separately in [Performance Evidence](#).

See also

[Technical Guide](#), [Encoding and Decoding](#), [Performance Model](#), and [Verification Sessions](#).

Chapter 45

Exact Block Reductions

45.1 Exact Block Reductions

A block reduction has to add B products exactly and round **once**. The textbook way to do that is a *superaccumulator* — a fixed-point register wide enough to hold every representable product without loss, typically hundreds of bits, updated per lane. `SmallFloats.jl` does not have one, and does not need one.

This page explains what it does instead, why that is exact rather than merely accurate, and where the two guards that make it exact actually live.

The problem

A block is `(scale, elements)` — one scale factor shared by B elements (draft §5):

```
julia> b = Block(Binary8pluf(0.5), (Binary8p4se(80.0), Binary8p4se(-0.375),  
                                   Binary8p4se(-5.5), Binary8p4se(2.25)));  
  
julia> scaleformat(b), elemformat(b), blocksize(b)  
(Binary8pluf, Binary8p4se, 4)
```

Reducing it means summing the four *lane values* `scale × elementi`. Do that lane by lane in the result format and every partial sum rounds:

```
julia> SmallFloats.blockdecode(b)           # the exact lane values  
(40.0, -0.1875, -2.75, 1.125)  
  
julia> BlockReduceAdd(Binary8p4se, RNE_SN, b) # exact sum 38.1875, projected ONCE  
Binary8p4se(40.0 ≈ 0x6a)  
  
julia> lanes = SmallFloats.blockdecode(b);  
  
julia> foldl((a, v) -> Add(Binary8p4se, RNE_SN, a, Convert(Binary8p4se, RNE_SN, v)),  
            lanes; init = zero(Binary8p4se)) # rounding at every step  
Binary8p4se(36.0 ≈ 0x69)
```

One code point apart, from four lanes. The exact sum is 38.1875; the ulp in `[32, 64)` is 4, so the neighbours are 36 and 40 and the correct answer is 40. Accumulated rounding lands on 36 — not catastrophically wrong, but *not the defined result*, and the error grows with B .

So the reduction must be exact before it is projected. The question is what to accumulate in.

Step 1: decode each lane exactly

`blockdecode` produces `wMultiply(decode(s), decode(xi))` per lane, exactly, on a carrier chosen from **both** formats.

That last point is the whole content of the function, and it prevents a subtle misapplication. A block's scale and element formats are independent, and what has to fit is their *product*. The carrier rule is $\Sigma B_i \leq e_{\max}$, so two formats that are each comfortable alone can compose into a lane product that is not: a `Binary16p1uf` scale reaches 2^{32768} . Selecting the carrier from either format's own rung would overflow to $\pm\text{Inf}$ — a plausible-looking special value rather than an error, which is the worst available failure mode. The carrier is therefore keyed on `rung(Val(:Multiply), FS, FE)`.

Step 2: accumulate exactly — two strategies

Once the lanes are exact, the sum must be too. There are two ways to get there, and the package tries the cheap one first.

The fast path: a span filter

Float128 carries a 113-bit significand. A sum of B values is exactly representable in it when

$$(\text{lane significand bits}) + (\text{exponent span}) + \lceil \log_2 B \rceil \leq 113$$

— the significand bits of the widest lane, plus the number of binades the lanes straddle, plus the carries from adding B of them. Every term is an *integer*, and the exponent span is one pass over the lanes. So the package computes the span and decides:

reduction	shipped guard
BlockReduceAdd	$\text{span} + \lceil \log_2 B \rceil \leq 92$
BlockDotProduct	$\text{span} + \lceil \log_2 B \rceil \leq 76$

If the guard passes, a plain Float128 accumulation **is** the exact answer — no rounding occurs at any step, so there is nothing to be careful about. The dot product's budget is tighter because each of its terms is a product of *four* datums (two scales, two elements) rather than two.

Why the filter is sound

The guards look generous against the worst case you would compute on paper. Two Binary16p11se datums multiply to 22 significand bits, which leaves 91 for $\text{span} + \lceil \log_2 B \rceil$ — one less than the 92 the sum guard admits.

The resolution is that **lane width and exponent span are anti-correlated, not independent**. A datum only reaches the bottom of a format's exponent range by being *subnormal*, and a subnormal datum carries proportionally fewer significand bits. In Binary16p11se, full 11-bit significands exist only at exponents $-15\dots15$; below that, every bit of extra reach costs a bit of width. The two quantities in the inequality cannot both be maximal.

That argument is easy to get wrong, so it is checked rather than trusted: an audit over wide format pairs admitted **17 040** sums and **16 479** dot products through the guards and compared every one against exact BigFloat truth. Zero were inexact.

Two guards, not one

The span filter is about the **significand**. It says nothing about whether the lanes *fit in Float128 at all* — and at rung 3 they do not, because a Binary16p1uf scale reaches 2^{32768} , which is Inf in Float128 before any accumulation begins.

So the accumulator is selected by the **head** as well: `_f128acc` answers false for HeadExact unconditionally. Width and range are independent conditions and both must hold. An earlier version checked only the width, and the exponent condition was a premise nobody had stated.

The fallback: precision derived from the formats

When either guard fails, an exact BigFloat accumulation takes over at a precision **derived from the block's two formats and its length** — never a constant:

$$\begin{aligned} \text{sums} & \quad 2(B_S + B_E + P_S + P_E) + 64 + \lceil \log_2 B \rceil \\ \text{products} & \quad B(P_S + P_E) + 128 \end{aligned}$$

Two different formulas because two different analyses. For a sum, the lanes straddle $2(B_S + B_E)$ binades and each carries $P_S + P_E$ significand bits, plus carries. For a product reduction the exponent merely shifts, but the significand is the *sum* over lanes.

Both replaced magic numbers, and the history is the reason they are written out. The product form was $16B + 128$, where $16 = 8 + 8$ is $P_S + P_E$ with $P = 8$ hardcoded — correct at $K \leq 8$ and understating by up to $2\times$ above it. Its sibling $_REDPREC = 2400$ was ample through $B \approx 512$ and silently truncating above. Both failure modes produce a plausible wrong number rather than an error, which is exactly the class of defect this package treats as unacceptable.

Step 3: project once

Whichever path produced it, the exact result meets the projection engine exactly once. `ωBlockProject` follows the draft's special rows for the scale (NaN , 0 , $\pm\text{Inf}$), then divides each element result by the scale through a cheapest-first cascade:

```
exact Float64 quotient → correctly rounded Float128
                        → bracket / pre-filter → MPFR interval
```

mirroring the scalar quotient group's rigor arguments. Each rung is tried only if the one before it cannot decide the code point.

P = 1 scales divide exactly, for free

A $P = 1$ format has no stored significand bits, so every datum is $\pm 2^e$. Dividing by a power of two is exact, and the first rung of the cascade always succeeds — the draft calls this out in its own NOTE. This is why the MX-style scale format `Binary16pluf` is not just conventional but cheap: the whole enclosure ladder collapses to one correctly rounded division.

Why no superaccumulator

A superaccumulator is the right answer when you cannot inspect the data — a fixed register sized for the worst case, paid for on every lane. Here the data *is* inspectable: one integer pass over the lane exponents answers the question the accumulator width was insuring against, and answers it exactly. The common case then runs in hardware `Float128`, and the uncommon case gets arbitrary precision sized to the actual formats.

A register-resident superaccumulator remains a tracked optimization for the Phase-3 work — as a *speed* change on the fallback path, not a correctness one.

Source and gates

`src/blocks.jl` owns block construction, scale/element operations, span filtering, exact reductions, and `BlockVector`. Review changes against three obligations: decoded lane values are exact, the chosen accumulator covers the proven span, and exactly one final projection occurs. Compare fused reductions with a deliberately wider reference, including mixed scales and special values.

See also

- [Blocks and Scaled Operations](#) — the API contract.
- [Use Blocks for Dynamic Range](#) — the task-oriented guide.
- [Oracle and Rigor Classes](#) — where the enclosure cascade comes from.
- [Proving your fused kernel exact](#) — `BlockDotProduct` against 512-bit truth.
- [Add an Operation](#) — adding block variants.

Chapter 46

Verification Strategy

46.1 Verification Strategy

The doctrine in one sentence: **enumerate rather than sample wherever the input space allows it, and where it does not, say "sampled" in so many words.** The formats are finite, so most correctness questions here are theorems checkable by exhaustion rather than statistics — and a tier the caller asks for is honoured, never downgraded, because a gate that quietly runs less than it was told to is worse than no gate.

What Is Verified

The suite combines exhaustive finite-domain gates with explicitly sampled axes. The final roll-call is the authoritative record for a run: it names the selected quick, default, or release tier, the units compared by every gate, and whether each gate was exhaustive or sampled. Major coverage includes:

- formats against an independent draft transliteration (14 679);
- ordering over all 2.5 M same-format pairs plus Next-op edge tables (7.6 M);
- unary, divide, and ternary operations against high-precision reference paths, with the exact format/input coverage reported by their gates;
- stochastic R-sweeps with directed-asymptote pins;
- table $\equiv$ scalar over every entry, including the ternary tiers (eager and adaptive) against the scalar path;
- the sticky-head FMA/FAA escalation against the MPFR reference across every rounding-mode family and adversarial cancellation cases;
- blocks against a from-scratch reference composition;
- Float128 carrier $\equiv$ Float64 carrier, and Float128-first $\equiv$ MPFR differential builds;
- the mask rounding core $\equiv$ the generic core over datums and reachable sums/products at boundary stochastic budgets $N \in \{45, 60\}$;
- packed round-trips, and κ /conformance behavior.

Deterministic **specialization regressions** (concrete inferred return types at the public entry points, zero warm-path allocation) stand in for timing assertions.

The Approach to Benchmarking

Correctness is asserted here; timing methodology belongs to [Benchmark Correctly](#), which records the four measurement post-mortems and the rules `benchmarking/benchmarking.jl` enforces structurally. The division matters: a gate that fails means a wrong answer, a benchmark that misleads means a wrong decision, and conflating them is how a harness bug gets filed as a performance regression.

Deliberate limitations

No implicit cross-format arithmetic exists. Arithmetic with an ordinary Julia number promotes to the format's public promotion carrier, which may be wider than `Float64`; callers who want a particular result format must name it and a projection explicitly. There is no in-place packed arithmetic. Threading is opt-in for pure Shape-B compute loops at every arity, only above the configured size cutoff and when `Threads.nthreads() > 1`; stochastic paths remain sequential to preserve RNG-stream order. Irrational/Rational inputs to `Convert` are rejected rather than double-rounded silently. The `Float128` machinery never changes results — disabling it (`SmallFloats_Float128=disable`) is a tested no-op semantically.

Reading a verification result

Treat the roll-call label as part of the result. **Exhaustive** means every member of the stated finite domain was checked. **Sampled** means the report must also state the generator, seed, sample count, and why that sample addresses the risk. A pass without its domain and method is not a portable claim.

The gate files under `test/` own correctness assertions; benchmark scripts own timing methodology. Documentation examples are checked separately so a green package test cannot conceal stale printed output.

See also

[Function Tables and Array Kernels](#), [Benchmark Correctly](#), [Verify Custom Code](#).

Chapter 47

Add an Operation

47.1 Add an Operation

How to add a new operation to `SmallFloats.jl` so that it inherits everything the existing catalog has: bit-exact defined results, the full projection-mode vocabulary, generated scalar/array/block surfaces, table caching, exhaustive test coverage, benchmark pickup, and a truthful conformance declaration. Read [Architecture and Invariants](#) and [Oracle and Rigor Classes](#) first — this page assumes their vocabulary (the result-kind protocol, the two rigor classes, the interval protocol, the `yd` envelope stage).

What's generated vs what you write

Generated for you, from one registry line: the spec-register method `Op(fr, p, operands...; rng, R)`, the same-format convenience method `Op(x...)`, array methods (Shape-A table gathers or Shape-B loops), the Block/Scaled variants, the export, table cacheability, the conformance listing, and test-harness pickup.

You must write: the ω -mathematics (`weval` in `src/oracle.jl`) and the six universal duties below — special rows, the result-kind protocol, termination (the Niven duty), envelope justification, ladder validity, and running the gates.

The architecture is registry-driven: for the supported arities, **most of the surface is generated**. Adding an operation is mostly (1) one registry line and (2) one rigorous `weval` method — plus the *duties* that keep the package's guarantees true. Every subsection below states the duties explicitly; they are not optional.

Every new operation

1. **Special rows are spec, not optimization.** Write out the NaN/±Inf/zero rows explicitly in `weval`, in draft style, before any general evaluation. Mark any behavior the draft text under-determines with `[interp]`.
2. **Result-kind protocol.** `weval` must return one of `Float64` | `Float128` | `StickyF` | `BigExactF` | `EncloseF` | `Enclose128F`, with `Float64`/`Float128` reserved for *provably exact* results and `StickyF` for exact-head-plus-tail-direction results whose soundness bound is discharged at the emitting site (see the FMA/FAA wide-spread escalations). Keep the return union narrow (ideally `{Float64, EncloseF}`) — a wide union defeats inference and costs allocations (measured: a 4-type union turned a 50 ns op into 240 ns with 2 allocations).
3. **Termination (the Niven duty).** `project_interval` terminates only if the enclosure is not chasing a value the projection grid can actually hit. Any input whose *exact* result is a representable dyadic must be **peeled or detected** before an enclosure is built. For transcendental ops this is usually a finite set of special rows (prove completeness, as the π -family does with Niven's theorem); for algebraic ops it is an exactness test.

4. **Envelope justification.** If you supply an eager `yd` estimate, you owe an error bound: the true value must lie within $|yd| \cdot 2^{-45}$ (`_F64_RELEXP`). Acceptable bases, strongest first: IEEE-CR hardware ops ($\div$, `sqrt`: $\leq \frac{1}{2}$ ulp, unconditional); short compositions of CR ops ($\leq \sim 1.5$ ulp); faithful `libm` (≤ 1 ulp, an empirical-but-published claim); short faithful compositions with condition number ≤ 1 ($\leq \sim 3$ ulp). **Measure it:** sample your estimator against ≥ 256 -bit MPFR truth over the op's domain, including the treacherous regions (near zeros of the result, domain boundaries), and demand $\geq 2^6$ slack under 2^{-45} . The same duty applies to a `Float128` fq estimator against its 2^{-90} envelope. Beware **argument-error amplification**: an estimate computed as $f(g(x))$ where g rounds (e.g. $\sin(\pi \cdot r)$ in `Float64`) can be catastrophically worse near zeros of f than a native result-domain evaluation (`sinpi(r)`) — measure the composition you actually ship.
5. **Ladder validity.** The MPFR closure must return a genuine directed enclosure. Whole-expression set rounding is valid only when the composition is monotone in every intermediate (state the argument in a comment, as `Softplus` does); otherwise round each step in the correct direction explicitly (as `_mpfr_sqrt` does for the anti-monotone $1/\sqrt{x}$).
6. **Gates.** Run the full suite. The unary exhaustive `xhp` gate and the coverage check pick up registry ops automatically; extend the hand-curated lists where applicable (monotonicity ops, binary exhaustive sweeps). Then regenerate the benchmark report — `bench_scalar_ops` iterates the registry lists, so your op appears in all four operand-class tables automatically; **add a `_SAFE_DOMAINS` entry** in `benchmarking/benchmarking.jl` if the op is argument-restricted, or the `safe-args` table will use the oracle-derived fallback pool.

Scalar operations

The registry lives in `src/ops_scalar.jl`. Adding a name to `_UNARY_OPS` / `_BINARY_OPS` / `_TERNARY_OPS` causes, with no further work:

- registration in `OP_REGISTRY` (group `:B/:C/:A` per the loop defaults — override with an explicit `register_op!` if the default group is wrong);
- the generated spec-register method `Op(fr, p, operands...; rng, R)` and the same-format convenience `Op(x...)`;
- generated array methods routing through `vmap` (Shape-A tables for pure p — always at arity ≤ 2 , bitwidth-gated at arity 3 — Shape-B scalar/threaded loops otherwise);
- generated `BlockOp` and `ScaledOp` variants (`src/blocks.jl`);
- the export (`src/SmallFloats.jl` loops the registry);
- table cacheability, conformance listing, and test-harness coverage.

What is *not* generated is the mathematics: `weval` in `src/oracle.jl`.

Adding a unary operation

Worked example: **Cbrt** (cube root), chosen because it exercises the hardest duty — exactness detection for an algebraic op.

Step 1 — registry (`src/ops_scalar.jl`): append `:Cbrt` to `_UNARY_OPS`. The loop registers it as group `:B` with arity 1.

Step 2 — ω -semantics (`src/oracle.jl`), with every duty discharged.

Type the row `::AbstractFloat`, not `::Float64`

An enclosure-ladder row is reached at **every rung**. `_LADDER_OPS` is derived by subtraction in `oracle.jl` — whatever is neither an exact selection, nor exact arithmetic, nor `Convert` — so a new registry operation automatically gets `_wf128`, `_wexact` and `_wdyadic` rows that forward to *this* method with `Float128`, `BigFloat` or converted-from-Dyadic operands.

A row typed `::Float64` therefore works for the 432 formats at rung 1 and raises a `MethodError` for the other 72. The package types 22 of its own ladder rows `::AbstractFloat` for exactly this reason; the 6 that are `::Float64` are the exact-arithmetic rows, which have per-head siblings written separately.

Gate G10 catches the omission — it calls every registry operation at every rung and asserts only that the call returns — and G4 asserts that `_EXACT_SELECTION`, `_EXACT_ARITH`, `_LADDER_OPS` and `Convert` partition `OP_REGISTRY` exactly, so an operation cannot be classified into none of them and silently lose its wide-rung rows.

```
# Exact-cube detection (termination duty): x = ±m·2^(3k) with m an odd perfect
# cube ⇒ cbrt(x) is a Float64-exact dyadic that project_interval must never
# chase. Oracle-path cost is irrelevant; correctness of the peel is everything.
function _exact_cbrt(x::Float64)
    fr, e = frexp(abs(x))          # |x| = fr·2^e, fr ∈ [0.5, 1)
    m = Int64(ldexp(fr, 53)); e -= 53 # |x| = m·2^e exactly
    tz = trailing_zeros(m); m >= tz; e += tz
    e % 3 == 0 || return nothing
    r = round(Int64, cbrt(Float64(m)))
    for c in max(r - 1, 1):(r + 1) # cbrt is faithful ⇒ true root ∈ r ± 1
        Int128(c)^3 == Int128(m) && return flpsign(ldexp(Float64(c), e ÷ 3), x)
    end
    nothing
end

function weval(::Val{:Cbrt}, x::AbstractFloat) # ::AbstractFloat, NOT ::Float64
    isnan(x) && return NaN # special rows first: they are spec
    isinf(x) && return x # cbrt(±∞) = ±∞
    iszero(x) && return 0.0
    c = _exact_cbrt(x)
    c === nothing || return c # exact dyadic result – peeled
    # yd: Float64 cbrt is faithful (≤ 1 ulp ≈ 2-52; measure it) ⇒ ≥ 27 slack
    # under the 2-45 envelope. fq: libquadmath cbrtq under the 2-90 envelope.
    # Ladder: cbrt is monotone ⇒ whole-expression directed rounding encloses.
    _encl1(cbrt, x; fq=() -> cbrt(Float128(x)), yd=cbrt(x))
end
```

Step 3 — optional surface: a Base veneer (`Base.cbrt(x::Binary) = Cbrt(x)` next to the others in `ops_scalar.jl`), a docstring, and — since `cbrt` is monotone — its symbol in the test suite's monotonicity list.

Step 4 — validate: envelope measurement for `cbrt` (include tiny and huge `x`; `cbrt` has no dangerous regions — the result is never near zero except at zero, which is peeled), then the full suite: the exhaustive `xhp` gate now checks `Cbrt` on every code point of the reference format under four modes against the 3072-bit protocol, automatically.

Adding a binary operation

Worked example: **LogAddExp** — $\log(e^x + e^y)$ — chosen because it exercises composed estimators and the ladder-monotonicity argument.

Step 1 — registry: append `:LogAddExp` to `_BINARY_OPS` (default group `:C`; the special five arithmetic names get `:A`).

Step 2 — ω -semantics. One generic evaluator serves all three precisions — Float64 (yd), Float128 (fq), and BigFloat (the ladder):

```
# Stable form: h + loglp(exp(l - h)), h = max, l = min. Every step is monotone
# nondecreasing in (x, y), so whole-expression directed rounding in _mpfr2 is a
# valid enclosure (the same argument Softplus records).
_logaddexp(a, b) = (h = max(a, b); l = min(a, b); h + loglp(exp(l - h)))

function weval(::Val{:LogAddExp}, x::AbstractFloat, y::AbstractFloat)
    (isnan(x) | isnan(y)) && return NaN
    (x == Inf || y == Inf) && return Inf          # ∞ dominates
    x == -Inf && return y                        # log(0 + ey) = y, exactly
    y == -Inf && return x
    # yd: two faithful libm calls plus additions, condition ≤ 1 throughout
    # (result ≥ max(x, y)) ⇒ ≤ ~3 ulp ≈ 2-50 - 25 slack; measure it, including
    # x ≈ y (where loglp(exp(0)) dominates) and widely separated operands.
    _encl2(_logaddexp, x, y;
        fq=() -> _logaddexp(Float128(x), Float128(y)),
        yd=_logaddexp(x, y))
end
```

Termination note: can $\logaddexp(x, y)$ be exactly dyadic for dyadic operands with the special rows peeled? $\log(e^x + e^y) = d$ requires $e^x + e^y = e^d$ — by Lindemann-Weierstrass this forces relations that dyadic arguments cannot satisfy except on the peeled rows, so no further peel is needed; **record that argument in a comment**. When you cannot prove such a statement for your op, hunt for exact cases the way `_exact_cbrt` does.

Step 3 — tests: binary ops are not in the unary exhaustive gate; add a divide-style exhaustive sweep (256×256 operand pairs against a ladder-only reference, a handful of modes) to the suite, and an envelope-sanity loop entry if you introduced a new libm dependence.

Adding a ternary operation

Worked example: **FMS** — fused multiply-subtract, $x \cdot y - z$ — chosen because the cheapest correct implementation is *delegation*, and delegation is a first-class technique here.

Step 1 — registry: append `:FMS` to `_TERNARY_OPS` (group `:A` by the loop). The generated array method routes through the same ternary bitwidth policy every other ternary op uses (`tables.jl`'s `_ternary_table_for`): small operand formats table (eagerly or adaptively), $K = 8$ runs the — optionally threaded — Shape-B scalar loop, and stochastic p always draws per element. No op-specific work needed; the policy keys on formats and p , not on which op it is.

Step 2 — ω -semantics by delegation to FMA's already-verified analysis (width thresholds, `_twosum` exactness, wide-spread fallback):

```
# x·y - z ≡ FMA(x, y, -z). Negation of z follows the Subtract convention:
# preserve NaN, and map -0 to +0 so the single-zero datum model is respected.
weval(::Val{:FMS}, x::Float64, y::Float64, z::Float64) =
    weval(Val{:FMA}, x, y, isnan(z) ? z : (iszero(z) ? 0.0 : -z))
```

Delegation inherits FMA's result kinds, its exactness thresholds, *and* its test pedigree — but verify the sign algebra of your reduction on the special rows (here: $\text{FMS}(\infty, 1, \infty)$ must be NaN, which the delegation produces because $\text{FMA}(\infty, 1, -\infty)$ is $\infty - \infty$). Add the delegated op to the ternary exhaustive sweep in the suite.

Adding a quaternary operation

Arity 4 is **not yet plumbed** — the generation loops, kernels, and block layer handle arities 1-3, so adding one is two tasks: extend the plumbing once, then add the op like any other. Four sites need the arity-4 branch: the spec-register generation loop in `src/ops_scalar.jl` (register manually, since the arity loops stop at 3 — `apply_op` and `weval` are already variadic); a four-array `vmap!` in `src/kernels.jl` (the plain Shape-B loop suffices — table-tiering at arity 4 means 2^{4K} growth, so extend it only for a measured win); a `Block/Scaled` branch or a deliberate `op.arity > 3` && `continue skip` in `src/blocks.jl`; and `conformance_report's` for `a in 1:3` in `src/approx.jl`.

The op itself follows the FMA/FAA width-analysis playbook unchanged. Two duties get harder at this arity. The exact-escalation precision must be **derived** — `bigprec_prod(x, y, z, w)`, never a constant:

Derive the precision from the operands — never a constant

This escalation used to read `_BIGP`, a 2200-bit constant. It was ample for every format that existed at $K \leq 8$ and **insufficient for 50 of the 504**, and it did not announce that: it returned a plausible wrong number, which gate G2 was written red to record (1//1 where the exact answer is 2//1, the operand returned unchanged because a residual 30 000 binades below the rounding position fell off the accumulator).

`bigprec(F...)` derives the precision from the formats; `bigprec(x...)` and `bigprec_prod(x...)` derive it from the operands in flight. Gate G2 asserts the derivation is sufficient across the whole grid, so a new operation that uses it inherits that proof. A new constant would need its own.

And the testing duty changes in kind: the full cross-product is 2^{32} tuples, so the suite entry must be *sampled* (seeded, ladder-referenced, $\geq 10^5$ tuples across several modes) plus hand-picked exhaustion of the special-row algebra and the width-threshold boundary. Document the sampling in the test itself — it breaks the suite's "enumerate, never sample" norm, and that exception must be visible rather than silent.

Block operations

The block layer (`src/blocks.jl`) implements the draft's elementwise schema **once** — `blockdecode` → ω -op `lanewise` → `blockproject` — and generates every registry op's `BlockOp/ScaledOp` from it. The reductions are hand-written on a second shared pattern. Which subsection you need depends on whether your operation fits one of those two patterns.

Adding a scaled block operation

If the scalar operation is (or can be) a **registry op, the scaled form is free**: registering `:Cbrt` above already produced `ScaledCbrt(fr, p, s, x)` and `BlockCbrt(fr, p, b, sr)` with the draft's §5.4 semantics — decode `scale × element` exactly per lane, apply the ω -op, `BlockProject` against the result `scale` through the division cascade. **Prefer this route**; it is the package's non-divergence mechanism.

Write a scaled op by hand only when it is *not* an elementwise image of a registry op. The primitive to compose with is `_bp_element(fr, p, R, res, Sdat)` — it owns the draft's scale-special rows (NaN/0/±Inf scale) and the entire division-by-scale rigor cascade. Example, a scaled affine $s_1 \cdot x + s_2 \cdot y$ projected against a *unit* result scale:

```
function ScaledAffine(fr::Type{<:Binary}, p::ProjSpec,
                    s1::Binary, x::Binary, s2::Binary, y::Binary;
                    rng::MaybeRNG=nothing)
    # Each (scale, element) pair is a TWO-FACTOR MONOMIAL and gets its own head
    # from the two formats — the same rule `blockdecode` applies to a block.
    h1 = rung(Val(:Multiply), typeof(s1), typeof(x))
    h2 = rung(Val(:Multiply), typeof(s2), typeof(y))
    a = weval(h1, Val(:Multiply), lift(h1, decode(s1)), lift(h1, decode(x))):: carriertype(h1)
    b = weval(h2, Val(:Multiply), lift(h2, decode(s2)), lift(h2, decode(y))):: carriertype(h2)
    # The two products may land on DIFFERENT carriers, so the add runs at their
    # join with each operand lifted into it.
    h = _joinheads(a, b)
    res = weval(h, Val(:Add), lift(h, a), lift(h, b))
    _bp_element(fr, p, _drawR(p, rng, nothing), res, 1.0)
end
```

The assertion is `::carriertype(h)`, and writing `::Float64` there is a real defect rather than a stylistic choice. It is true for the 432 formats at rung 1 and an outright `TypeError` for the other 72, whose lanes decode to `Float128` or to the exact dyadic carrier. The generated `Scaled*` surface in `blocks.jl` carried exactly that mistake and it is recorded as §11 M44 in the execution log — the guidance here used to teach it.

What the assertion *is* load-bearing for survives the correction: `Multiply` is closed on the carrier at every rung (gate G4 pins this), so the row cannot escalate to a lazy `BigExactF` here, and saying so is what lets `_joinheads` below see carrier values only. JET finds the missing statement before any test can.

The `Add` result may be any kind in the protocol — `_bp_element` finishes all of them. Dividing by a scale of `1.0` is the identity fast path.

Adding an elementwise block operation

Same schema, whole-block granularity. The three obligations: decode operand blocks with `blockdecode` (exact), evaluate lanewise with `weval` (never with ad-hoc `Float64` math), and finish with `blockproject` against the supplied result scale (never with per-lane project — the scale-division rows belong to `blockproject`). Example — `BlockRelu`, an elementwise $\max(x, 0)$ that reuses the verified `MaximumNumber` semantics:

```
"""BlockRelu(fr, p, b, sr): lanewise max(xi, 0) — NaN lanes propagate NaN per
MaximumNumber's number-preferring rows only when both operands are NaN, i.e.
never here; state your lane semantics this explicitly for any new op."""
function BlockRelu(fr::Type{<:Binary}, p::ProjSpec, b::Block{B}, sr::Binary;
                  rng::MaybeRNG=nothing) where {B}
    X = blockdecode(b)
    Z = ntuple(i -> weval(Val(:MaximumNumber), X[i], 0.0), Val(B))
    blockproject(fr, p, sr, Z; rng)
end
```

Add the name to the block-surface export list and to `conformance()`'s `blocknames` extension vector so the declaration stays truthful. Test against a from-scratch reference composition (the suite's existing block tests show the pattern) — element format $\times$ scale format $\times$ B, exhaustively where feasible.

Adding a reductive block operation

Reductions do **not** fit the elementwise schema; they follow the second pattern, visible in `BlockReduceAdd/BlockDotProduct`: (1) resolve the ∞/NaN **fold algebra** on `Float64` classifications first; (2) an integer **span filter** decides whether the whole reduction is exactly representable in `Float128` — if so, plain `Float128` accumulation is the exact answer; (3) otherwise a `BigExactF` exact accumulator at provably sufficient precision; (4) exactly **one** projection, at the very end, via `_finish`. Never round intermediates.

Worked example — `BlockSumOfSquares`, $\sum x_i^2$:

```

"""BlockSumOfSquares(fr, p, b): project( $\sum$  decode-lane $i^2$ ) – one rounding total.
Fold algebra: any NaN lane  $\Rightarrow$  NaN; else any  $\pm\text{Inf}$  lane  $\Rightarrow$   $+\text{Inf}$  (squares cannot
cancel – no opposite-infinity NaN case exists, unlike ReduceAdd); else exact."""
function BlockSumOfSquares(fr::Type{<:Binary}, p::ProjSpec, b::Block{B};
    rng::MaybeRNG=nothing, R::Union{Nothing,Int}=nothing) where {B}
    X = blockdecode(b)
    res = if any(isnan, X)
        NaN
    elseif any(isinf, X)
        Inf
    elseif all(iszero, X)
        0.0
    else
        # span filter: lanes carry  $\leq 17$ -bit significands  $\Rightarrow$  squares  $\leq 34$  bits and
        # a square's exponent is  $2 \cdot (\text{lane exponent})$ ; the sum is exact in Float128
        # when  $34 + 2 \cdot \text{span} + \lceil \log_2 B \rceil + 1 \leq 113 \Rightarrow 2 \cdot \text{span} + \lceil \log_2 B \rceil \leq 78$ .
        if _f128() && 2 * _expspan(X) + _log2ceil(B) <= 78
            acc = Float128(0)
            for v in X
                q = Float128(v)
                acc += q * q                # exact squares, exact sum
            end
            acc
        else
            # precision DERIVED from the formats (the retired-`_REDPREC` lesson):
            # squares carry  $2(P_S+P_E)$  bits and span  $4(B_S+B_E)$  binades, plus
            # 64 slack and the carry term – `_lane_sum_prec`'s analysis, doubled
            # where squaring doubles it.
            p = 4 * (expbias(scaleformat(b)) + expbias(elemformat(b))) +
                2 * (precision(scaleformat(b)) + precision(elemformat(b))) +
                64 + _log2ceil(B)
            BigExactF() -> setprecision(BigFloat, p) do
                acc = BigFloat(0)
                for v in X
                    acc += BigFloat(v)^2    # exact at derived precision
                end
                acc
            end
        end
    end
    _finish(fr, p, _drawR(p, rng, R), res)
end

```

The two lines that demand your care in any new reduction are the **fold algebra** (derive it from the draft's reduce definition on the extended reals — here the absence of an opposite-infinities NaN case is a *theorem about squares*, stated in the docstring) and the **span-filter inequality** (derive it from operand widths, write the derivation in the comment, and add boundary-condition tests at the threshold, exactly as the suite does for the existing `_DE_*` constants). Everything else — the exact accumulator, the single `_finish`, stochastic R plumbing — is pattern.

Checklist

- Full test suite green, including your new sweep entries.
- Envelope measurements recorded for any new yd/fq estimator.
- Benchmark report regenerated; `_SAFE_DOMAINS` entry added if the op is argument-restricted; sanity-check the op's row in the safe-args table.
- Docstring on the public name; `[interp]` markers where you interpreted.
- `conformance_report()` shows the op (automatic for registry ops; manual block additions extend blocknames).

See also

[Oracle and Rigor Classes](#), [Exact Block Reductions](#), [Verify Custom Code](#).

Chapter 48

Verify Custom Code

48.1 Verify Custom Code

Trust, then verify — these patterns turn "is my code exact?" into an enumeration. Formats are small enough that "test on a few points" is never necessary; the recipes below enumerate the input space (or, where the space is too large to enumerate, sweep it exhaustively along the dimension that matters) and compare against an independently computed reference. Several examples use unexported internals (`SmallFloats.round_to_precision`, `SmallFloats.project`, ...); those are stable enough to learn from but are not covered by semantic-versioning guarantees.

Verify a quantizer

Here every in-range datum of `Binary8p3se` is projected into `Binary8p4se` and checked against a brute-force nearest search under 256-bit arithmetic (distance agreement; the draft's tie rule is the projection engine's job and is pinned elsewhere in the suite):

```
using SmallFloats

function ref_nearest_distance(::Type{T}, x) where {T}
    fins = [v for v in (SmallFloats.rawvalue(T, UInt8(c)) for c in 0:255) if isfinite(decode(v))]
    minimum(abs(setprecision(() -> BigFloat(decode(v)) - BigFloat(x), BigFloat, 256))
            for v in fins)
end

function verify_quantizer()
    ok = true
    for c in 0x00:0xff
        x = decode(SmallFloats.rawvalue(Binary8p3se, c))
        (isfinite(x) && abs(x) <= decode(MaxFiniteOf(Binary8p4se))) || continue
        got = Convert(Binary8p4se, RNE_SN, x)
        ok &= abs(decode(got) - x) == Float64(ref_nearest_distance(Binary8p4se, x))
    end
    ok
end
verify_quantizer()
```

```
true
```

This pattern — enumerate the inputs, compare against an independently written reference — is how the entire package is tested, and it is available to *your* quantization code at trivial cost.

Prove a fused kernel

The block layer promises *one* projection with exact lane arithmetic. Trust, then verify — against big-float truth, over random blocks with mixed scales and an honest special-value mix:

```
using SmallFloats, Random

rng = Xoshiro(5)
mk() = Block(Binary8pluf(2.0^rand(rng, -3:3)),
             ntuple(_ -> SmallFloats.rawvalue(Binary8p4se, UInt8(rand(rng, 0:255))), 32))

function verify_dots(trials)
  for _ in 1:trials
    bx, by = mk(), mk()
    lx = [decode(bx.s) * decode(v) for v in bx.x]
    ly = [decode(by.s) * decode(v) for v in by.x]
    (any(!isfinite, lx) || any(!isfinite, ly)) && continue
    truth = setprecision(() -> sum(BigFloat(lx[i]) * BigFloat(ly[i]) for i in 1:32),
                                BigFloat, 512)
    got = BlockDotProduct(Binary8p4se, RNE_SN, bx, by)
    ref = setprecision(() -> SmallFloats.project(Binary8p4se, RNE_SN, truth), BigFloat, 512)
    codepoint(got) == codepoint(ref) || return false
  end
  true
end
verify_dots(200)
```

```
true
```

The same shape verifies any custom fused kernel you write: compose the truth in BigFloat, project once, compare code points.

Audit stochastic rounding

Every stochastic projection is a deterministic function of the draw R , so its *distribution* is checkable exactly. For $x = 2 + 3/64$ in Binary8p4se the fraction is $v = 3/16$ of an ulp, so StochasticA{4} must round up for exactly 3 of the 16 draws:

```
σ4 = RSA_SN(4) # ≡ ProjSpec(StochasticA{4}(), SatNone())
x = 2.0 + 3/64
count(decode(SmallFloats.project(Binary8p4se, σ4, x; R)) == 2.25 for R in 0:15)
```

```
3
```

Sweeping R like this turns "is my stochastic pipeline unbiased?" from a statistical question into an exhaustive one — the pattern the shipped test suite uses.

Applying monotonicity

Decision and ranking systems consume scores through *order*. If scores are produced in one format and compared in another, the conversion must be monotone — and formats are small enough to prove it by enumeration rather than trust it. Walk every Binary8p3se datum in total order and check that its Binary8p4se image never goes backward on the integer order keys:

```
using SmallFloats
using SmallFloats: order_key

function monotone_conversion(::Type{From}, ::Type{To}) where {From,To}
    prev = nothing
    for v in sort(From.(0x00:UInt8(2^bitwidth(From) - 1)))
        isnan(decode(v)) && continue
        g = Convert(To, RNE_SN, v)
        prev !== nothing && order_key(g) < order_key(prev) && return false
        prev = g
    end
    true
end
monotone_conversion(Binary8p3se, Binary8p4se)
```

```
true
```

Run this for every (From, To, ρ) triple your system actually uses; it is a few hundred integer comparisons per triple. A ranking pipeline whose conversions all pass this audit cannot invert a preference by changing formats — a guarantee no amount of spot-testing provides.

Measuring Approximation

Search under a compute budget often wants a cheap, approximate evaluation function. The κ registry turns "how approximate?" into a *measured* code-point bound — and code-point bounds compose into a decision rule: **if two defined evaluations differ by more than 2κ code points along the total order, the approximate evaluator cannot invert their comparison.** Verify the rule exhaustively for a $\kappa = 2$ evaluator:

```
using SmallFloats: order_key

fast(x) = (r = Exp(Binary8p4se, RNE_SN, x);
           isfinite(decode(r)) ? NextGreaterThan(NextGreaterThan(r)) : r)
κ, exhaustive = measure_kappa(fast, :Exp, Binary8p4se, (Binary8p4se,), RNE_SN)

function margin_audit(fast, κ)
    codes = [SmallFloats.rawvalue(Binary8p4se, UInt8(c)) for c in 0:255]
    safe = violations = 0
    for a in codes, b in codes
        da, db = Exp(Binary8p4se, RNE_SN, a), Exp(Binary8p4se, RNE_SN, b)
        (isnan(decode(da)) || isnan(decode(db))) && continue
        codedistance(da, db) > 2κ || continue
        safe += 1
        (order_key(fast(a)) < order_key(fast(b))) ==
            (order_key(da) < order_key(db)) || (violations += 1)
    end
    (safe, violations)
end
(κ, exhaustive, margin_audit(fast, κ)...)

(2.0, true, 49522, 0)
```

49,522 operand pairs clear the 2κ margin, and the approximate evaluator agrees with the defined ordering on every one of them — zero violations, exhaustively. Inside the margin, comparisons are genuinely undecidable at this κ ; a search can treat sub-margin comparisons as ties to expand, or escalate those few nodes to the exact evaluator. Either way the pruning is *provably* sound, with κ measured at registration rather than promised.

See also

[Add an Operation](#), [Benchmark Correctly](#), [Verification Sessions](#).

Chapter 49

Benchmark Correctly

49.1 Benchmark Correctly

The benchmark doctrine is the record of four measurement post-mortems. The first two are about the *closure*: a closure over a non-const global measures Julia's dispatch machinery as well as the code under test, and that distortion changes ratios, not only absolute timings, because dispatch cost varies with call shape. The other two are about the *harness*, and they are below. Treat any quoted timing as a record of a particular format, policy, machine, and benchmark harness.

The rules

The rules the shipped benchmarking/benchmarking.jl enforces structurally:

- format types enter as type parameters, never as globals;
- operands come from untimed setup;
- functions retrieved reflectively pass through argument barriers to specialize;
- a preflight aborts the run if warm scalar paths allocate — including the wide-spread FMA/FAA sticky-head path.

The table-build section reports both the cold build (cache evicted per sample) and the steady-state warm cache hit, since callers amortize the former through the latter. Vary one binding per variant; verify specialization before believing a number.

Two ways the harness itself can lie

- **One variant per process, alternating, or nothing.** A harness that compiles several variants in one process is not a measurement. It reported an `add_sticky_dy` rewrite as a $1.87\times$ *win*; alternating single-variant processes reversed the same comparison into a 31% *loss*. A related form: timing two kernels sequentially in one process let the first perturb the second into a 39% phantom regression on code that had not changed.
- **A loop that overwrites its accumulator measures dead-code elimination.** `z = f(xs[i])` leaves every iteration but the last dead, and whether the loop survives depends on whether the compiler can prove the body pure — so it compares *elimination* between implementations rather than work. Measured: a pure kernel under that loop reported 8 ns for 4096 elements. Reduce into an accumulator and end with `Base.dontdelete`; `benchmarking/benchmark_dyadics.jl` shows the pattern on every loop.

Set up a measurement that is actually measuring your code

The package's benchmark doctrine, in one snippet (needs the benchmarking/ environment for Chairmarks). Format types enter as **type parameters**; operands come from Chairmarks' *untimed* setup; and if you retrieve functions reflectively (getfield(SmallFloats, op)), pass them through an argument barrier so they specialize:

```
using Chairmarks, SmallFloats, Random
using Statistics: median

function bench_add(::Type{T}) where {T<:Binary}          # T: type parameter, not a global
    pool = [SmallFloats.rawvalue(T, rand(UInt8)) for _ in 1:4096]
    @be (rand(pool), rand(pool)) (t -> Add(T, RNE_SN, t[1], t[2]))(_)
end

b = bench_add(Binary8p4se)
round(median(b).time * 1e9; digits=1), median(b).allocs
```

Record that result with the Julia version, CPU/thread configuration, source revision, operand pool, and projection. The value is evidence for that setup; it is not the package's scalar-Add contract.

Do not benchmark unresolved globals

If T comes from a non-const global, the result includes dynamic dispatch and cannot fairly compare projection, carrier, table, or compute paths. The shipped benchmarking/benchmarking.jl checks specialization before it accepts a result; your benchmark should do the same.

See also

[Verification Strategy](#), [Verify Custom Code](#), [Verification Sessions](#) — whose *Benchmarking without measuring the dispatcher* is the complete version of the snippet above, with the measurement post-mortems that produced it.

Chapter 50

Documentation Claim Ledger

50.1 Documentation claim ledger

This ledger keeps implementation-sensitive prose reviewable. It is a maintenance artifact: readers should use the guides and references, while authors use this file to prevent a correct Code8 statement from becoming an incorrect claim about every Binary format.

ID	Claim and scope	Primary source authority	Documentation owner	Review status
FMT-1	Legal formats span $K = 3:16$; Code8 and Code16 are distinct concrete representations selected from the abstract format.	src/formats.jl (KMAX, KSPLIT, reptime)	format anatomy, format reference	reviewed 2026-08-18
CODEC-1	decode(v) returns datumcarrier(typeof(v)); the result is exact for every datum.	src/carriers.jl, src/decode_encode.jl	values, core model, codec guide	reviewed 2026-08-18
CODEC-2	Code8 uses a generated decode tuple; Code16 uses computational decode and refuses a wide constant tuple.	src/formats.jl, src/decode_encode.jl	codec guide, technical guide	reviewed 2026-08-18
CODEC-3	Encode returns codeunit_type(F) and documents a caller-side datum-set precondition.	src/decode_encode.jl, src/project.jl	projection/codec guides	reviewed 2026-08-18
ORDER-1	NaN is the smallest total-order member; order-key type scales from UInt16 to UInt32.	src/formats.jl, src/decode_encode.jl	values, code-point algebra, codec guide	reviewed 2026-08-18
TABLE-1	Tables store result code units of the result representation and are admitted by the current byte-budget policy; compute paths remain semantically identical.	src/tables.jl, src/kernels.jl	tables guide, performance model	reviewed 2026-08-18
CARRIER-1	Carrier selection and oracle rigor are separate: a datum carrier answers representation range; a rigor class justifies an operation result.	src/carriers.jl, src/oracle.jl	architecture, oracle guide	reviewed 2026-08-18
PERF-1	Timings, allocation figures, and thread scaling are measurements, not package guarantees; every reported number names its format, path, environment, and measurement script.	benchmarking/, docs/ src/examples_performance.md	performance pages	pending measurement refresh

When source changes, update the row before changing prose. A new numerical claim requires a new row or a refreshed PERF-1 entry with its command, environment, and date.

Part VIII

Help

Chapter 51

Troubleshooting

51.1 Troubleshooting

Symptom-first reference. Find what you are seeing, not what you did wrong — each entry names the symptom, the cause, the fix, and where to read the full explanation.

Values and Construction

Symptom: `T(0x02)` gives a different value than `T(2)`, and it looks like a bug. Cause: an Unsigned argument is a *code point*, at every bitwidth — `T(0x02)` constructs code point 0x02, while `T(2)` projects the numeric value two. Signedness of the integer type is what distinguishes them, so this holds for `UInt8`, `UInt16`, and any other Unsigned. Fix: use `T(Int(c))` when you mean the number, `T(c::Unsigned)` when you mean the code point, or call `Convert` when intent should be unmistakable. Full explanation: [Cheat Sheet](#), Values, Code Points & Conversion.

Symptom: `Binary8p4se(0x02)` and `convert(Binary8p4se, 0x02)` disagree. Cause: `convert` is the deliberate exception to the Unsigned-is-a-code-point rule and is value-preserving — `Binary8p4se(0x02)` is code point 2, while `convert(Binary8p4se, 0x02)` is the value 2.0. Promotion, similar, and fill depend on `convert` behaving this way. Fix: know which one you are calling; prefer `T(x)` or `Convert(T, p, x)` in your own code and reserve bare `convert` for places Base itself calls it. Full explanation: [Cheat Sheet](#), Values, Code Points & Conversion.

Symptom: constructing from a Rational throws. Cause: `Rational` is not one of the supported construction carriers. Fix: convert explicitly to an exact supported carrier (a `Float64`, `Dyadic`, or similar) or knowingly to `Float64` first, then construct. Full explanation: [Cheat Sheet](#), Values and code points.

Symptom: `T(-0.0) == zero(T)` surprises code that expects signed zero. Cause: every P3109 format has exactly one zero; there is no negative zero to distinguish. Fix: stop testing for negative zero; if you need "approached from below," use the projection engine's sticky argument instead of relying on a stored sign. Full explanation: [Core Model](#), [Encoding and Decoding](#).

Symptom: a value built with `SmallFloats.rawvalue` decodes to garbage or a non-existent datum. Cause: `SmallFloats.rawvalue(T, c)` is an internal unchecked constructor — it skips the range validation that `T(c::Unsigned)` performs, so an out-of-range integer produces a bit pattern with no corresponding intent-checked datum. Fix: use validated `T(c::Unsigned)` outside package implementation code. Full explanation: [Cheat Sheet](#), Values and code points.

Symptom: dividing by zero gives NaN, not Inf as IEEE-754 code expects. Cause: P3109 defines $x / 0 \rightarrow \text{NaN}$ (and $\text{Recip}(\pm\text{Inf}) = 0$), which is not IEEE division-by-zero semantics. Fix: do not port IEEE-754 zero-division assumptions across; check for zero divisors explicitly if your algorithm depends on signed infinities from division. Full explanation: [Cheat Sheet](#), Values and code points; [P3109 in One Chapter](#).

Arithmetic and Semantics

Symptom: Multiply/Add overflowed and you expected a clamp, but got Inf (or vice versa). Cause: `SatNone` does not mean "no saturation" in the clamping sense — it applies the draft's domain-, signedness-, and direction-dependent rows, which can still produce `Inf`, `MaxFinite`, or `NaN` depending on the case. Only `SatFinite` guarantees clamping to the finite range. Fix: choose `SatFinite` whenever clamping is actually required; do not assume `SatNone` behaves like `clamp`. Full explanation: [Cheat Sheet](#), Projection specifications.

Symptom: mixing two Binary formats in one expression throws a promotion error instead of silently widening. Cause: this is deliberate — the package refuses silent cross-format promotion so that no rounding step happens somewhere nobody wrote. Fix: convert explicitly, `Convert(T, p, x)`, to the format you actually want before combining. Full explanation: [Cheat Sheet](#), Values and code points.

Symptom: `rem`, `mod`, or a similar Base numeric function refuses to run, or errors in a way that looks like a missing method. Cause: some generic Base numeric functions are defined in terms of operations or promotion behavior this package deliberately does not supply by default (for example, functions that would require an implicit cross-format promotion or an IEEE-754-only special case). This is a refusal, not an omission bug. Fix: express the computation with the package's explicit operation catalog (`Add`, `Subtract`, `Multiply`, `Divide`, `FMA`, ...) and `Convert` calls instead of relying on the generic fallback. Full explanation: [Cheat Sheet](#), Scalar operation catalog.

Arrays and Performance

Symptom: an array of Binary values is slow, or every element looks boxed. Cause: the element type is the *abstract* format `Binary{K,P,Σ,Δ}`, not a concrete alias — the abstract format is not `isbits`, so a generic-element array boxes every entry. Fix: use a concrete element type, `Vector{Binary8p4se}` or `Vector{format{K,P,Σ,Δ}}`. Full explanation: [Cheat Sheet](#), Common pitfalls (below), and [Architecture and Invariants](#).

Symptom: `Vector{Binary{K,P,Σ,Δ}}(undef, n)` silently boxes every element, even though you thought you had a concrete type. Cause: `Binary{K,P,Σ,Δ}` with type-parameter variables is abstract, and `Vector{...}(undef, n)` cannot be intercepted or redirected to a concrete representation the way constructors can. Fix: use `Vector{Binary8p4se}` (a concrete alias) or `Vector{format{K,P,Σ,Δ}}` (concrete once `K,P,Σ,Δ` are literal values); `similar` normalizes correctly and is safe. Full explanation: [Cheat Sheet](#), Common pitfalls.

Symptom: benchmarks report ~1 microsecond per call for an operation you expect to be much faster on a fixed, warm signature. Cause: the format type or projection spec was read from a non-const global, so every call pays Julia's dynamic dispatch instead of running the specialized method. This is a benchmarking-methodology bug, not a package performance regression — it has been traced to this exact mistake in real post-mortems. Fix: keep format types and projection specs in const bindings, function arguments, or type parameters; put runtime-selected-format code behind a function barrier; verify zero warm-path allocation before trusting a number. Full explanation: [Verification Sessions](#), Benchmarking without measuring the dispatcher; [Verification Strategy](#).

Symptom: the first array call for a given format/operation is much slower than every call after it. Cause: deterministic unary/binary operations build a cached result table on first use for that specialization; you measured the build, not the steady state. Fix: warm the specialization before benchmarking, and benchmark warm calls separately from the first call. Use `table_bytes()` to inspect the cache footprint and `empty_tables!()` to reset it. Full explanation: [Cheat Sheet](#), Performance checklist (pointer); [Function Tables and Array Kernels](#).

Random and Stochastic

Symptom: a stochastic-rounding pipeline gives different results every run, even though you want to compare against a fixed baseline. Cause: stochastic projections (`RSA`, `RSB`, `RSC`) consume random bits from an RNG (or an explicit draw), and without a fixed source that draw changes every run. Fix: supply a seeded rng (for example `Xoshiro(1)`), or pass an explicit `R` in `0:(2^N - 1)` for exact, reproducible draws — ideal in tests. Full explanation: [Cheat Sheet](#), Stochastic projections; Reproducible Stochastic Rounding.

Symptom: you passed a seed but two calls in the same pipeline still diverge. Cause: the uniform draw and the stochastic rounding draw must come from the *same* rng stream to stay reproducible together; if you construct a second Xoshiro mid-pipeline instead of threading the same rng through, the streams desynchronize. Fix: thread one rng object through every call in the pipeline that needs to agree. Full explanation: [Cheat Sheet](#), Stochastic projections.

Platform note: disabling Float128

On platforms where libquadmath misbehaves, set an environment variable **before** loading the package:

```
ENV["SmallFloats_Float128"] = "disable"  
using SmallFloats
```

This selects the pure-MPFR configuration. Results are **bit-identical** — that equivalence is itself part of the test suite; only oracle and build speed change.

See also

[Cheat Sheet](#), [Glossary](#), [Performance Model](#), [Verification Strategy](#).

Chapter 52

Cheat Sheet

52.1 Cheat Sheet

A compact reference for the SmallFloats.jl operations you are most likely to write. For semantics and explanation, follow the links out of each section; for implementation detail, see [Insights](#).

Start here

```
using SmallFloats
using Random: Xoshiro

T = Binary8p4se
ρ = RNE_SN

x = T(1.6)
y = T(0.25)
z = Add(T, ρ, x, y)
```

The explicit operation shape is always:

```
OperationName(result_format, projection_spec, operands...; rng, R)
```

For same-format operands under the default projection, ordinary Julia syntax is available:

```
z == x + y
exp(x)
fma(x, y, z)
```

Format names

BinaryKpP(s|u)(e|f) — bitwidth, precision, signedness, domain. The field-by-field grammar is in [Formats and Values](#); read it off the examples:

Name	Meaning
Binary8p4se	8-bit, precision 4, signed, extended
Binary8p4sf	8-bit, precision 4, signed, finite
Binary6p3ue	6-bit, precision 3, unsigned, extended
Binary5p5uf	5-bit, precision 5, unsigned, finite
Binary16p6se	16-bit, precision 6, signed, extended (opt-in export)
Binary16p1uf	16-bit, precision 1, unsigned, finite — the MX scale shape

The alias is the *representation* of the parametric format, related by <: and **not** by ==:

```
Binary8p4se <: Binary{8, 4, true, true}    # true
Binary8p4se == Binary{8, 4, true, true}    # FALSE — `Binary` is abstract
format(16, 6, true, true)                  # the alias, from runtime parameters
```

using SmallFloats exports the 120 aliases at $K \leq 8$. For the other 384: using SmallFloats.Formats, or SmallFloats.Binary16p6se, or format(...).

Useful format queries:

```
bitwidth(T)      # K
precision(T)     # P
issigned(T)
isextended(T)
expbias(T)
expbitwidth(T)
trailingsigbits(T)

MaxFiniteOf(T)
MinFiniteOf(T)
MinPositiveOf(T)
MinNormalOf(T)
MaxSubnormalOf(T)
```

Every query above also takes a **value** instead of a type — $\text{bitwidth}(x) \equiv \text{bitwidth}(\text{typeof}(x))$ — and folds to the same constant:

```
x = Binary8p4se(1.6)
bitwidth(x)      # 8
MaxFiniteOf(x)   # Binary8p4se(224.0 ≈ 0x7e)
```

Values and code points

```
x = T(1.6)                # numeric value: project into T
x = Convert(T, p, 1.6)    # explicit projection

c = codepoint(x)          # the code point, in the format's storage unit
x = T(c)                  # an Unsigned means validated CODE POINT
x = T(c)                  # checked code-point construction (`c::Unsigned`)

d = decode(x)             # the exact datum, on the format's own carrier
```

'Unsigned' means code point — at every width

`T(0x02)` constructs code point `0x02`; `T(2)` projects the numeric value two. Signedness of the integer type is what distinguishes them, so the rule holds for `UInt8`, `UInt16` and any other Unsigned: `Binary16p6se(0x0002)` is code point 2, not the value 2.0. Use `Convert` when intent should be unmistakable.

`codepoint` returns the format's *storage unit* — `UInt8` at $K \leq 8$, `UInt16` above — and that is `codeunit_type(T)`. A code-point argument is range-checked against 2^K whatever its width, so passing a wider Unsigned than necessary is lossless and safe.

`convert` is the deliberate exception and is value-preserving: `Binary8p4se(0x02)` is code point 2, while `convert(Binary8p4se, 0x02)` is the value 2.0. Promotion, similar and `fill` depend on that.

Random values (uniform `[0,1]` floor-projected; normal round-to-nearest, tails clamped to `MaxFinite`; `randn` signed formats only):

```
rand(T); rand(T, dims); rand!(A)      # uniform on [0, 1)
randn(T); randn(T, dims); randn!(A)  # standard normal, always finite
rand(Xoshiro(1), T)                  # explicit rng as usual

rand(rng, T; projection=RTP_SN)      # scalar forms: land under any ProjSpec
randn(rng, T; projection=RTZ_SF)     # (stochastic p draws R from same rng)
```

Classification and stepping:

```
isnan(x); isinf(x); isfinite(x); iszero(x)
signbit(x); issubnormal(x)
Class(x)
NextGreaterThan(x); NextLessThan(x)
nextfloat(x); prevfloat(x)
TotalOrder(x, y)
```

Projection specifications

A projection specification is (rounding mode, saturation mode):

```
p = ProjSpec(TowardPositive(), SatFinite())
roundingmode(p)
saturationmode(p)
```

Deterministic projections

Rounding	SatFinite	SatPropagate	SatNone
nearest, ties even	RNE_SF	RNE_SP	RNE_SN
nearest, ties away	RNA_SF	RNA_SP	RNA_SN
toward $+\infty$	RTP_SF	RTP_SP	RTP_SN
toward $-\infty$	RTN_SF	RTN_SP	RTN_SN
toward zero	RTZ_SF	RTZ_SP	RTZ_SN
round to odd	RTO_SF	RTO_SP	RTO_SN

RNE_SN is the package-wide default. Saturation modes mean:

Mode	Out-of-range behavior
SatFinite	clamp everything to the finite range
SatPropagate	preserve representable infinities; clamp other overflow
SatNone	apply the draft's domain-, signedness-, and direction-dependent rows

Stochastic projections

Constructors are grouped by stochastic variant:

Variant	SatFinite	SatPropagate	SatNone
A	RSA_SF(N)	RSA_SP(N)	RSA_SN(N)
B	RSB_SF(N)	RSB_SP(N)	RSB_SN(N)
C	RSC_SF(N)	RSC_SP(N)	RSC_SN(N)

N is the random-bit budget, $1 \leq N \leq 60$. Omitting it uses $N = 8$.

```

σ = RSA_SN(8)

Add(T, σ, x, y; rng=Xoshiro(1))  # reproducible stream
Add(T, σ, x, y; R=17)            # exact draw, ideal for tests

isstochastic(σ)                  # true
nrandbits(σ)                     # 8

```

For an N-bit mode, explicit R must be in $0:(2^N - 1)$.

Session defaults

Read with `DefaultX()`, set with `DefaultX!(v)`:

```

DefaultType()           # Binary8p2se    DefaultType!(Binary8p4se)
DefaultReturnType()     # Binary8p2se    DefaultReturnType!(Binary8p3se)
DefaultRoundingMode()   # NearestTiesToEven()
DefaultSaturationMode() # SatNone()
DefaultProjection()     # RNE_SN

```

Setting a rounding/saturation component rebuilds `DefaultProjection`; setting `DefaultProjection!` directly updates both components. Always: `DefaultProjection() == ProjSpec(DefaultRoundingMode(), DefaultSaturationMode())`.

The convenience methods `(a + b, Exp(x), T(2.1))` **do** consult `DefaultProjection`, so changing it changes their results globally. They read it through a speculation guard: allocation-free while the default holds its initial value, one dispatch per call after you change it. Explicit forms `(Add(T, p, x, y))` are unaffected by the session default.

Consume a default via the combinators — never by computing on a bare `DefaultX()` read. No dispatch while the default is unchanged, one barrier dispatch after a change; zero-alloc when `f`'s result type doesn't depend on the default (a default-typed result boxes once at escape):

```
with_default_type((T, x) -> T(x), 1.5)           # Binary8p2se(1.5 ≈ 0x41)
with_default_projection((p, x, y) -> Add(T, p, x, y), x, y)
with_default_returntype(f, args...)              # f(DefaultReturnType(), args...)
```

Scalar operation catalog

```
# explicit result format and projection
Add(T, p, x, y)
FMA(T, p, x, y, z)
Convert(T, p, external_value)

# same-format default-projection convenience
Add(x, y)
FMA(x, y, z)
```

Arity	Operations
Unary	Abs, Negate, Sqrt, RSqrt, Recip, Exp, Exp2, ExpMinusOne, Log, Log2, LogOnePlus, Softplus, Sin, Cos, Tan, ArcSin, ArcCos, ArcTan, Sinh, Cosh, Tanh, ArcSinh, ArcCosh, ArcTanh, SinPi, CosPi, TanPi, ArcSinPi, ArcCosPi, ArcTanPi
Binary	CopySign, Add, Subtract, Multiply, Divide, Hypot, ArcTan2, ArcTan2Pi, Maximum, Minimum, MaximumNumber, MinimumNumber, MaximumMagnitude, MinimumMagnitude, MaximumMagnitudeNumber, MinimumMagnitudeNumber, MinimumFinite, MaximumFinite
Ternary	FMA, FAA, Clamp
Conversion	Convert

Common Base spellings under the session projection (initially `RNE_SN`) — the full register is in [Julia Compatibility](#):

```
x + y; x - y; x * y; x / y
-x; abs(x); inv(x); sqrt(x)
exp(x); exp2(x); expm1(x)
log(x); log2(x); log1p(x)
sin(x); cos(x); tan(x)
asin(x); acos(x); atan(x); atan(y, x)
sinh(x); cosh(x); tanh(x)
min(x, y); max(x, y); clamp(x, y, z)
fma(x, y, z); muladd(x, y, z)
```

Arrays

Every registered operation has elementwise array methods:

```
A = T.([1.0, 1.5, 2.0])
B = T.([0.25, 0.5, 0.75])

C = Add(T, p, A, B)
E = Exp(T, p, A)
F = FMA(T, p, A, B, C)

C = vmap(:Add, T, p, A, B)
vmap!(C, Val(:Add), T, p, A, B)
```


Operands and destination must have matching axes. For table-eligible $K \leq 8$ formats, deterministic unary/binary operations use cached result tables and affordable ternary signatures may also use tables. Wider signatures compute directly. Stochastic operations always compute each element and consume one draw per projection.

```
table_bytes()
empty_tables!()
```

Blocks and scaled operations

```
FS = Binary8p1uf
FE = Binary8p4se

b = Block(FS(4.0), FE(1.5), FE(-0.75), FE(2.0), FE(0.5))

blocksize(b)
scaleformat(b)
elemformat(b)

ConvertFromBlock(T, ρ, b)
BlockReduceAdd(T, ρ, b)
BlockReduceMultiply(T, ρ, b)
BlockDotProduct(T, ρ, b, b)
```

Every scalar operation except Convert has generated block and scaled forms:

```
BlockAdd(T, ρ, b1, b2, result_scale)
BlockExp(T, ρ, b, result_scale)

ScaledAdd(T, ρ, scale1, x1, scale2, x2)
ScaledExp(T, ρ, scale, x)
```

Quantize against a supplied or automatically selected scale:

```
ConvertToBlock(FS, FE, ρ, values_tuple, scale)
ConvertToBlockMaxAbsFinite(FS, FE, scale_ρ, element_ρ, values_tuple)
```

BlockVector stores many equal-shape blocks in a structure-of-arrays layout.

Packed storage

```
v = Binary5p2se.([0.5, 1.0, 1.5, 2.0])
pv = PackedVector(v)

pv[2]
pv[2] = Binary5p2se(0.75)
collect(pv)

out = vmap(:Exp, Binary5p2se, ρ, pv)
```

PackedVector stores each code point in exactly $\text{bitwidth}(F)$ bits. Computation unpacks tiles internally; packed arithmetic is deliberately not in-place.

Conformance and approximations

```
conformance()
conformance_dict()
conformance_report()
draft_revision()
draft_identity()

κ, exhaustive = measure_kappa(fn, :Exp, T, (T,), ρ)
register_approx! (:fast_exp, :Exp, T, (T,), ρ, fn; κ)
impl = approx (:fast_exp)
kappa(impl)
kappa_measured(impl)
list_approx()
unregister_approx! (:fast_exp)
```

Approximate implementations are never substituted into the default API.

Common missteps

Misapplication	Correct pattern
Treating an Unsigned as a number	$T(\text{Int}(c))$ for a numeric integer; $T(c::\text{Unsigned})$ for a code point — at every width
$\text{Vector}\{\text{Binary}(K, P, \Sigma, \Delta), n\}$	Silently boxes every element — the abstract format is not isbits. Use $\text{Vector}\{\text{Binary8p4se}\}$ or $\text{Vector}\{\text{format}(K, P, \Sigma, \Delta)\}$. similar normalizes; $\text{Vector}\{\dots\}(\text{undef}, n)$ cannot be intercepted
Silently mixing formats	Convert explicitly: $\text{Convert}(T, \rho, x)$
Assuming SatNone always clamps	Choose SatFinite when clamping is required
Expecting IEEE division-by-zero	P3109 semantics here define $x / 0 \rightarrow \text{NaN}$
Expecting negative zero	Every format has one zero; $T(-0.0) == \text{zero}(T)$
Expecting sort to put NaN last	The draft's total order puts the single NaN first , below $-\text{Inf}$ (§4.12.1) — the opposite end from Float64
Passing a Rational	Convert explicitly to an exact supported carrier or knowingly to Float64
Reproducibility with stochastic rounding	Supply a seeded rng, or an explicit R in tests
Using the internal unchecked constructor	Prefer validated $T(c::\text{Unsigned})$ outside kernels

Symptom-first versions of these, with the error text you would actually see, are in [Troubleshooting](#).

Performance checklist

- Keep format types and projection specs in const bindings, function arguments, or type parameters.
- Put code using runtime-selected formats behind a function barrier.
- Expect the first deterministic array call for a specialization to build a table; benchmark warm calls separately.
- Use PackedVector when storage bandwidth matters more than direct byte access.
- Use `table_bytes()` to inspect cache footprint and `empty_tables!()` to reset it.
- Do not replace explicit projection semantics with `@fastmath`.

The reasoning behind each is in the [Performance Model](#); the measurement rules are in [Verification Strategy](#).

For worked applications, continue to the [Examples Index](#). For correctness and benchmark methodology, continue to [Architecture](#).

Chapter 53

Glossary

53.1 Glossary

Alphabetical, one or two lines per term, with links to the page where the term is defined or used in full.

See also

[Cheat Sheet](#), [Troubleshooting](#), [Architecture and Invariants](#).

Term	Definition	See
Carrier	The concrete numeric type an exact intermediate or datum is held on during computation — Float64 at rung 1, Float128 at rung 2, or an exact dyadic carrier at rung 3, chosen by the format's exponent bias, never by the caller.	Architecture and Invariants; Projection Engine
Code point	The 2^K -valued integer identity of a value within its format — what a Binary value actually stores, in bijection with the format's datum set.	Core Model
Code unit	The storage word a format's code points are packed into — UInt8 at $K \leq 8$, UInt16 above — returned by <code>codeunit_type(T)</code> and by <code>codepoint(x)</code> .	Cheat Sheet, Values, Code Points, and Conversion
Datum	The exact mathematical value a code point denotes — a member of the closed extended reals (finite dyadic values, $\pm\text{Inf}$, NaN).	Core Model
Datum carrier	The specific carrier type that is exact for a given format's datums, named by <code>datumcarrier(T)</code> ; not to be confused with the promotion carrier used for mixed-type Base arithmetic.	Projection Engine
Dyadic	An exact rational-like carrier, $S \cdot 2^Q$ with an Int128 significand and Int64 exponent, <code>isbits</code> , used at rung 3 where even Float128 cannot hold a format's range or precision; converts to and from Julia's Rational exactly.	Oracle and Rigor Classes; Standard and Design Notes
κ (kappa)	The maximum code-point distance, measured along the total order, between an approximate implementation's result and the defined result — measured by exhaustive enumeration at registration time, never declared unchecked.	Oracle and Rigor Classes; Defined, Stochastic, and Approximate Results
Niven peel	The termination argument for π -scaled trigonometric operations: a finite set of exactly-representable cases split off ("peeled") before interval enclosure begins, with Niven's theorem proving the peel set is complete.	Oracle and Rigor Classes
Omega-operation (ω)	The auxiliary, mathematically exact operation (<code>weval</code>) that a defined operation's result is the projection of — the thing that is computed <i>before</i> rounding and saturation ever see it.	Oracle and Rigor Classes
Order key	An integer, sign-magnitude-folded encoding of a value used for comparisons and sorting — monotone with the total order, NaN mapped to key 0 (below every datum, per §4.12.1), enabling $O(n)$ counting sort.	Projection Engine
Projection	The single write path that produces every code point in the package: exact result $\rightarrow$ RoundToPrecision $\rightarrow$ Saturate $\rightarrow$ Encode. There is no second rounding routine anywhere in the package.	Projection Engine
Projection specification (ProjSpec)	The pair (rounding mode, saturation mode) that parameterizes a projection, constructed as <code>ProjSpec(rounding, saturation)</code> or via a named constant like <code>RNE_SN</code> .	Cheat Sheet, Projection Specifications
Promotion carrier	The ordinary Julia float type (Float64 at rung 1, BigFloat at rung 2/3) that a Binary value promotes to when combined with an ordinary number — named by <code>promotecarrier(T)</code> ; distinct from the datum carrier used internally.	Standard and Design Notes; Julia Compatibility Register
Result kind	One of the five categories <code>weval</code> returns a computed result as (for example exact-Float64, exact-Dyadic, or an enclosure needing further work), which <code>apply_op</code> dispatches on to finish the projection.	Oracle and Rigor Classes
Rigor class (R/E)	The two grades of justification for a non-Float64 computed result: Class R (unconditional, provable without any accuracy assumption) and Class E (envelope-conditional, a faithful-but-not-correctly-rounded estimate validated by a measured slack).	Architecture and Invariants; Oracle and Rigor Classes
Rounding mode	One of the six deterministic modes (NearestTiesToEven, NearestTiesToAway, TowardPositive, TowardNegative, TowardZero, ToOdd) or three stochastic variants (RSA, RSB, RSC) that make up half of a projection specification.	Cheat Sheet, Projection Specifications
Rung	One of three carrier tiers a format is assigned to by its exponent bias: rung 1 (Float64, 432 formats), rung 2 (Float128, formats whose range needs more exponent room), rung 3 (exact dyadic carrier, formats whose range or precision breaks Float128).	Architecture and Invariants
Saturation mode	One of three policies for out-of-range projection results: <code>SatFinite</code> (clamp to the finite range), <code>SatPropagate</code> (preserve representable infinities, clamp other overflow), <code>SatNone</code> (apply the draft's domain-, signedness-, and direction-dependent rows).	Cheat Sheet, Projection Specifications
Saturation rows	The draft's twenty-row table mapping a rounded value's range classification to its final disposition (as-is, MaxFinite, MinFinite, $\pm\text{Inf}$, NaN), consulted by the Saturate stage of projection.	Projection Engine
Span filter	An integer pre-pass over lane exponents in a block reduction that proves a narrower carrier already covers every value the reduction can produce, licensing a faster accumulation path without loss of exactness.	Exact Block Reductions
Sticky (symbolic)	A sticky $\in \{-1, 0, +1\}$ argument to the projection engine meaning "the true value is the carrier value plus an infinitesimal of this sign" — how enclosure endpoints and asymptotes ($\tanh \rightarrow 1^-$) project exactly without inventing a fake carrier value.	Projection Engine
Table gather	The array-kernel strategy of computing a result by indexing a precomputed, cached lookup table (built once through the scalar path) rather than recomputing per element.	Function Tables and Array Kernels
Total order	The draft's total ordering over a format's value set, with the single NaN placed first, below $-\text{Inf}$ (§4.12.1); the basis for <code>TotalOrder</code> , <code>isless</code> , comparisons, and counting sort.	Projection Engine

Chapter 54

Standard and Design Notes

54.1 Standard and Design Notes

An index to the curated docs/other/ collection: the standard's concept map, the design rationales behind specific implementation choices, and the working papers that planned the $K > 8$ extension. Each entry carries a status label — read it before treating a document as current.

Designing to the Draft Standard

P3109 Draft Notes — [IEEE_D1_concepts.md](#). A concept map of the whole IEEE P3109/D1 draft (July 2026): format-defining parameters, value space, operation machinery, and projection, with every node linked back to the draft's line numbers. Status: normative-description.

The full draft transliteration, [IEEE_D1.md](#), is repo-only and is not part of this curated set — it is linked out from the concept map and from the working papers below (see Decision D1). Find it at [IEEE_D1.md](#) in the repository.

Design Rationales

dyadic_rational.md — [dyadic_rational.md](#). The Dyadic $\leftrightarrow$ Rational bridge: how the rung-3 exact carrier converts to and from Julia's Rational exactly, and the three laws that conversion must obey (round-trip on every finite dyadic, $\pm\infty$ agreement with Base's $1/0$, and rejection only at the true NaN slot). Status: design-rationale.

promoterules.md — [promoterules.md](#). Works out why promote_rule between two P3109 formats cannot mean "widen to a bigger format" the way Int/Float64 promotion does, and what the promotion carrier is instead. Status: design-rationale, partially superseded by the promotion-carrier section in the Cheat Sheet and Julia Compatibility guide — those pages reflect the shipped promotecarrier(T) behavior, and this document should be read as the reasoning behind it rather than as the current API description.

Float32inside.md — [Float32inside.md](#). A design study asking whether Float32 can replace Float64 as the internal carrier; verdict is no as a universal replacement, yes as a surface type and as a gated compute carrier for specific kernels. Status: exploratory.

Float32more.md — [Float32more.md](#). Companion to Float32inside.md: a detailed, empirically validated implementation plan for the Float32/BFloat16 surface, the exactness-gated compute path, and k -registered Float32 kernels, with measured pass/fail counts per milestone. Status: exploratory.

Development notes

Four documents plan and re-plan the extension of the format grid from $K \leq 8$ (120 at $K \leq 8$, retained as the default exports) to $3 \leq K \leq 16$ (504 formats). Read in this order; each says explicitly where it overrides its predecessor.

extendingK.md — [extendingK.md](#). The first pass: what the draft permits, the two axes (storage unit and evaluation carrier/rung) that govern the extension, and format counts by starting rung. Status: exploratory/historical.

doingtheextensions.md — [doingtheextensions.md](#). Turns the plan into an architecture — the representation lattice, the carrier dispatch lattice, and generated per-format trait methods — staged so every step compiles and passes independently. Status: exploratory/historical.

implementextensions.md — [implementextensions.md](#). The execution plan: a line-by-line audit of the source tree that found ten findings (errors or gaps) against the inherited plan, followed by a ten-stage, single-commit-per-stage implementation schedule. Status: exploratory/historical.

morejulian.md — [morejulian.md](#). A separate audit of where the package's Julia interface departs from idiomatic Julia — ranked by user impact rather than by extension stage — covering broken `AbstractFloat` contract methods and deliberate departures alike. Status: exploratory/historical.

See also

[Cheat Sheet](#), [Glossary](#), [Architecture and Invariants](#).